



Generative Design

Generative Adversarial Networks

The collage consists of four panels arranged horizontally, each featuring a central image with overlaid text and a smiley face icon.

- Plants:** Shows a close-up of green leaves and small green berries. A yellow flower-shaped smiley face is on the left, and a pink star-shaped smiley face is at the top right. Below the image is a white box containing the word "Plants".
- Machine Learning:** Shows a close-up of a computer motherboard. Below the image is a white box containing the text "Machine Learning".
- How does happiness looks like?**: Shows a person in a white hoodie working on a laptop. A green hexagonal smiley face is at the top right. Below the image is a white box containing the text "How does happiness looks like?".
- Human Designers:** Shows a person in a white hoodie working on a laptop. A green hexagonal smiley face is at the top right. Below the image is a white box containing the text "Human Designers".

Communicate Emotions

Paul Ekman identified six basic emotions that he suggested were universally experienced in all human cultures.

Happiness
Sadness
Disgust
Fear
Surprise
Anger

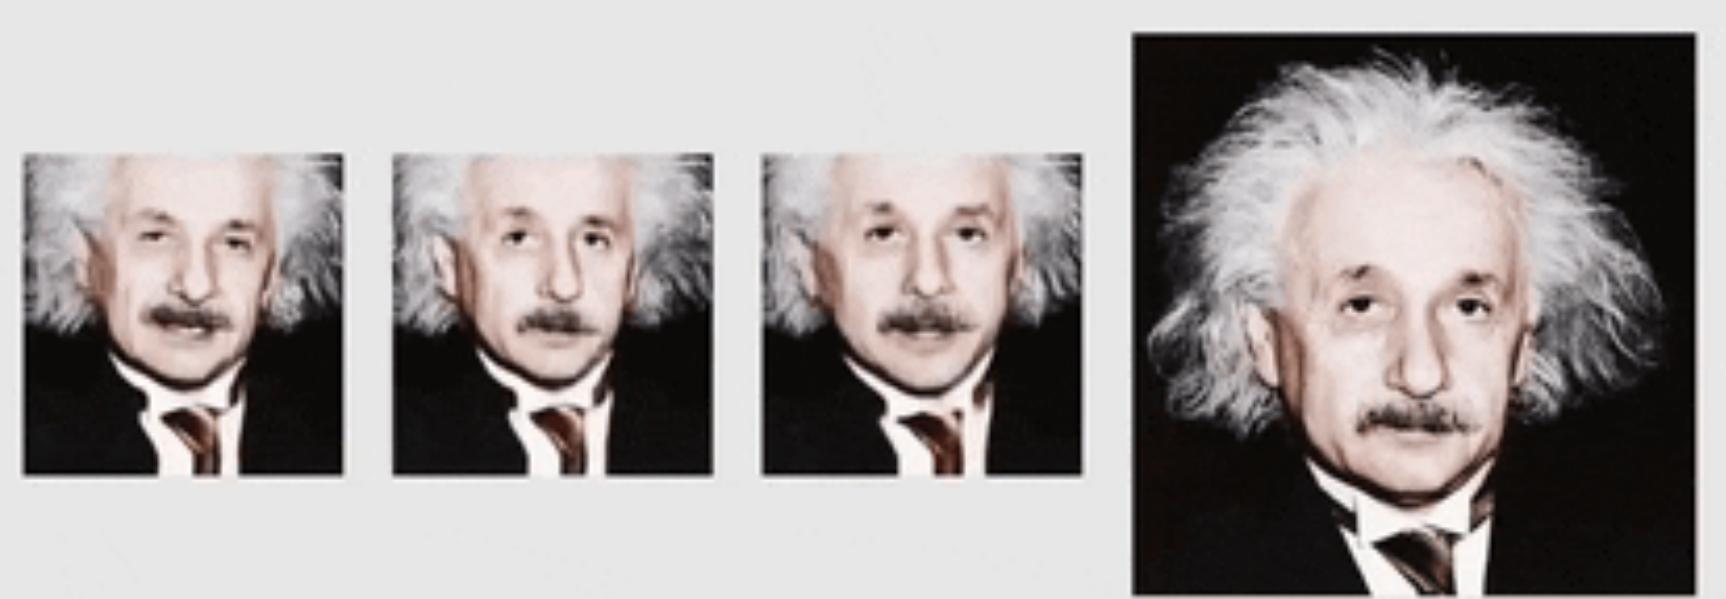
The collage consists of four panels arranged horizontally, each featuring a central image with overlaid text and a large, stylized emoji.

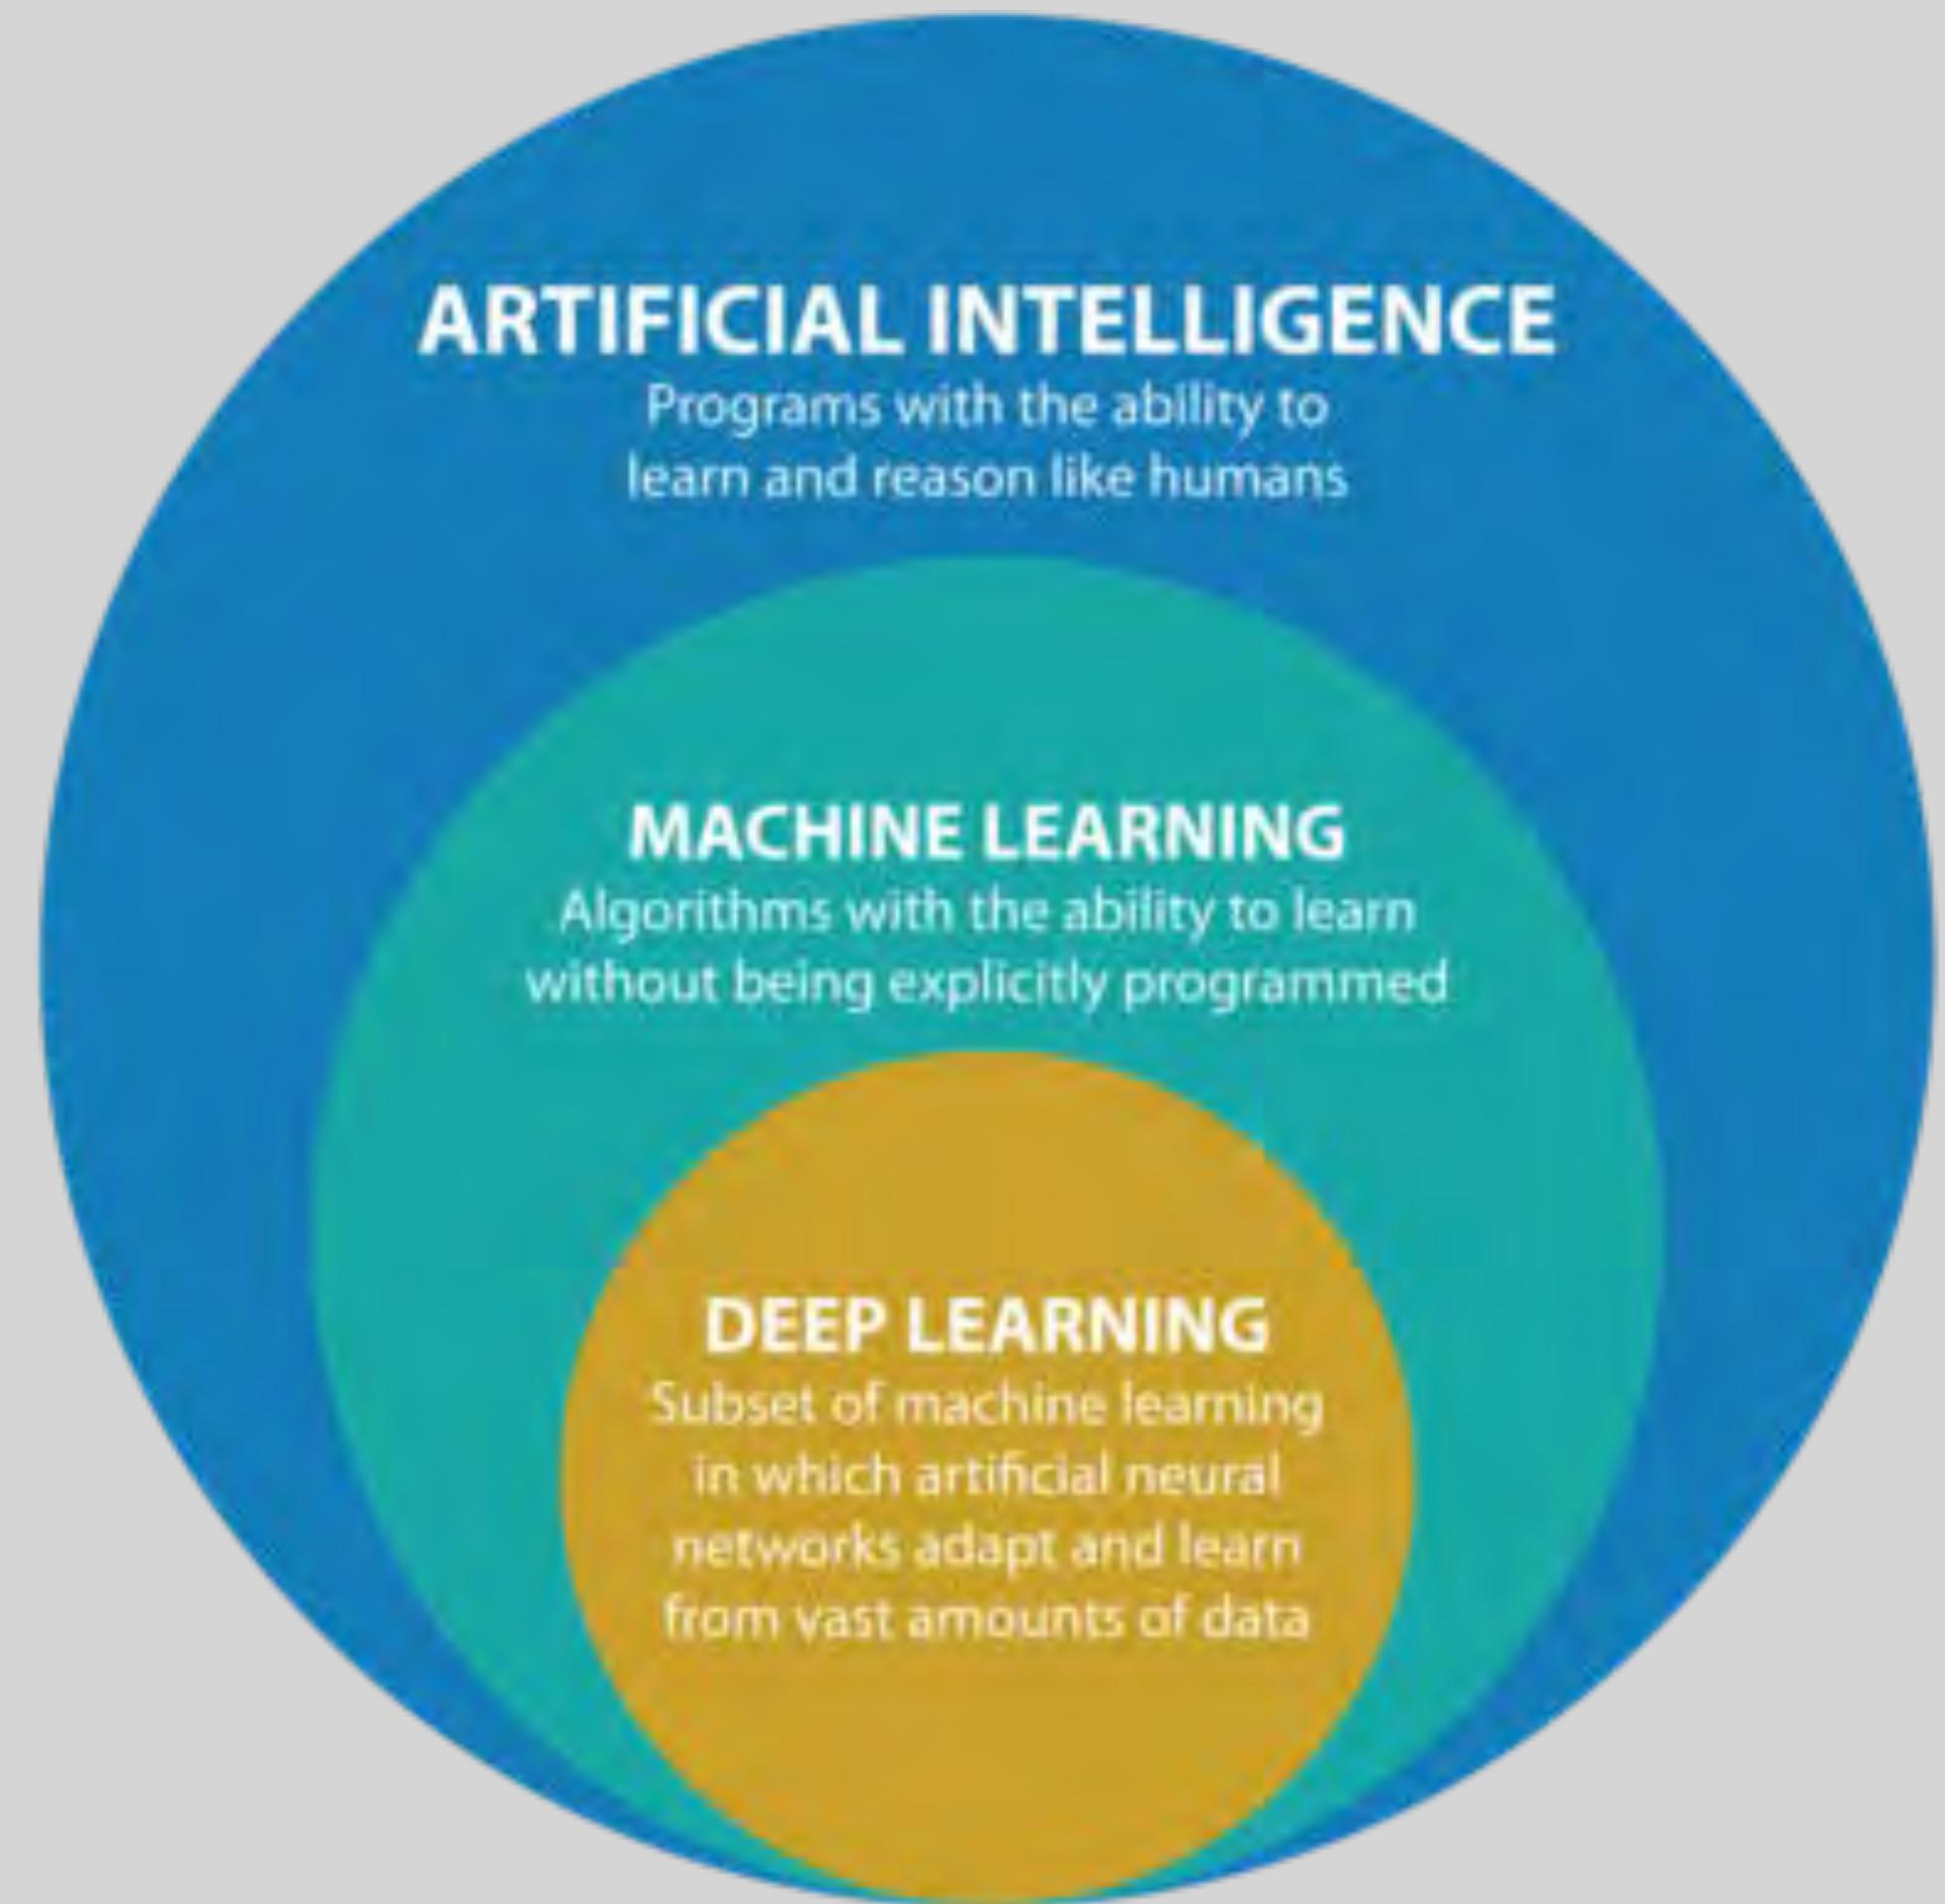
- Plants:** Shows a close-up of green leaves and small green fruits. A yellow flower emoji is in the bottom left, and a pink smiley face emoji is in the top right. Below the image is a white box containing the word "Plants".
- Machine Learning:** Shows a close-up of a computer motherboard. A red star emoji is in the top right. Below the image is a white box containing the text "Machine Learning".
- Communicate Emotions:** Shows a pink gradient background with a yellow smiley face emoji. Below the image is a white box containing the text "Communicate Emotions" and a paragraph about Paul Ekman's work on basic emotions.
- Human Designers:** Shows a person in a white lab coat working on a laptop. A green hexagonal emoji with a sad face is in the top right. Below the image is a white box containing the text "How does happiness looks like?" and a green box containing the text "How does happiness looks like?" followed by two smaller flower and star emojis.

As smartphone or computer users we use

We have all heard on seen what AI can do in terms of **image generation**

AI in a daily basis





During this workshop we want you
use Machine Learning (GAN) to:

1

Understand the
possibilities that AI
presents for the Design
process?

2

Learn not only the theory
but also use and try to
generate something

3

What kind of
opportunities can this
present applied in a
services, product or
experience?

Our Approach

To use AI as a
design tool

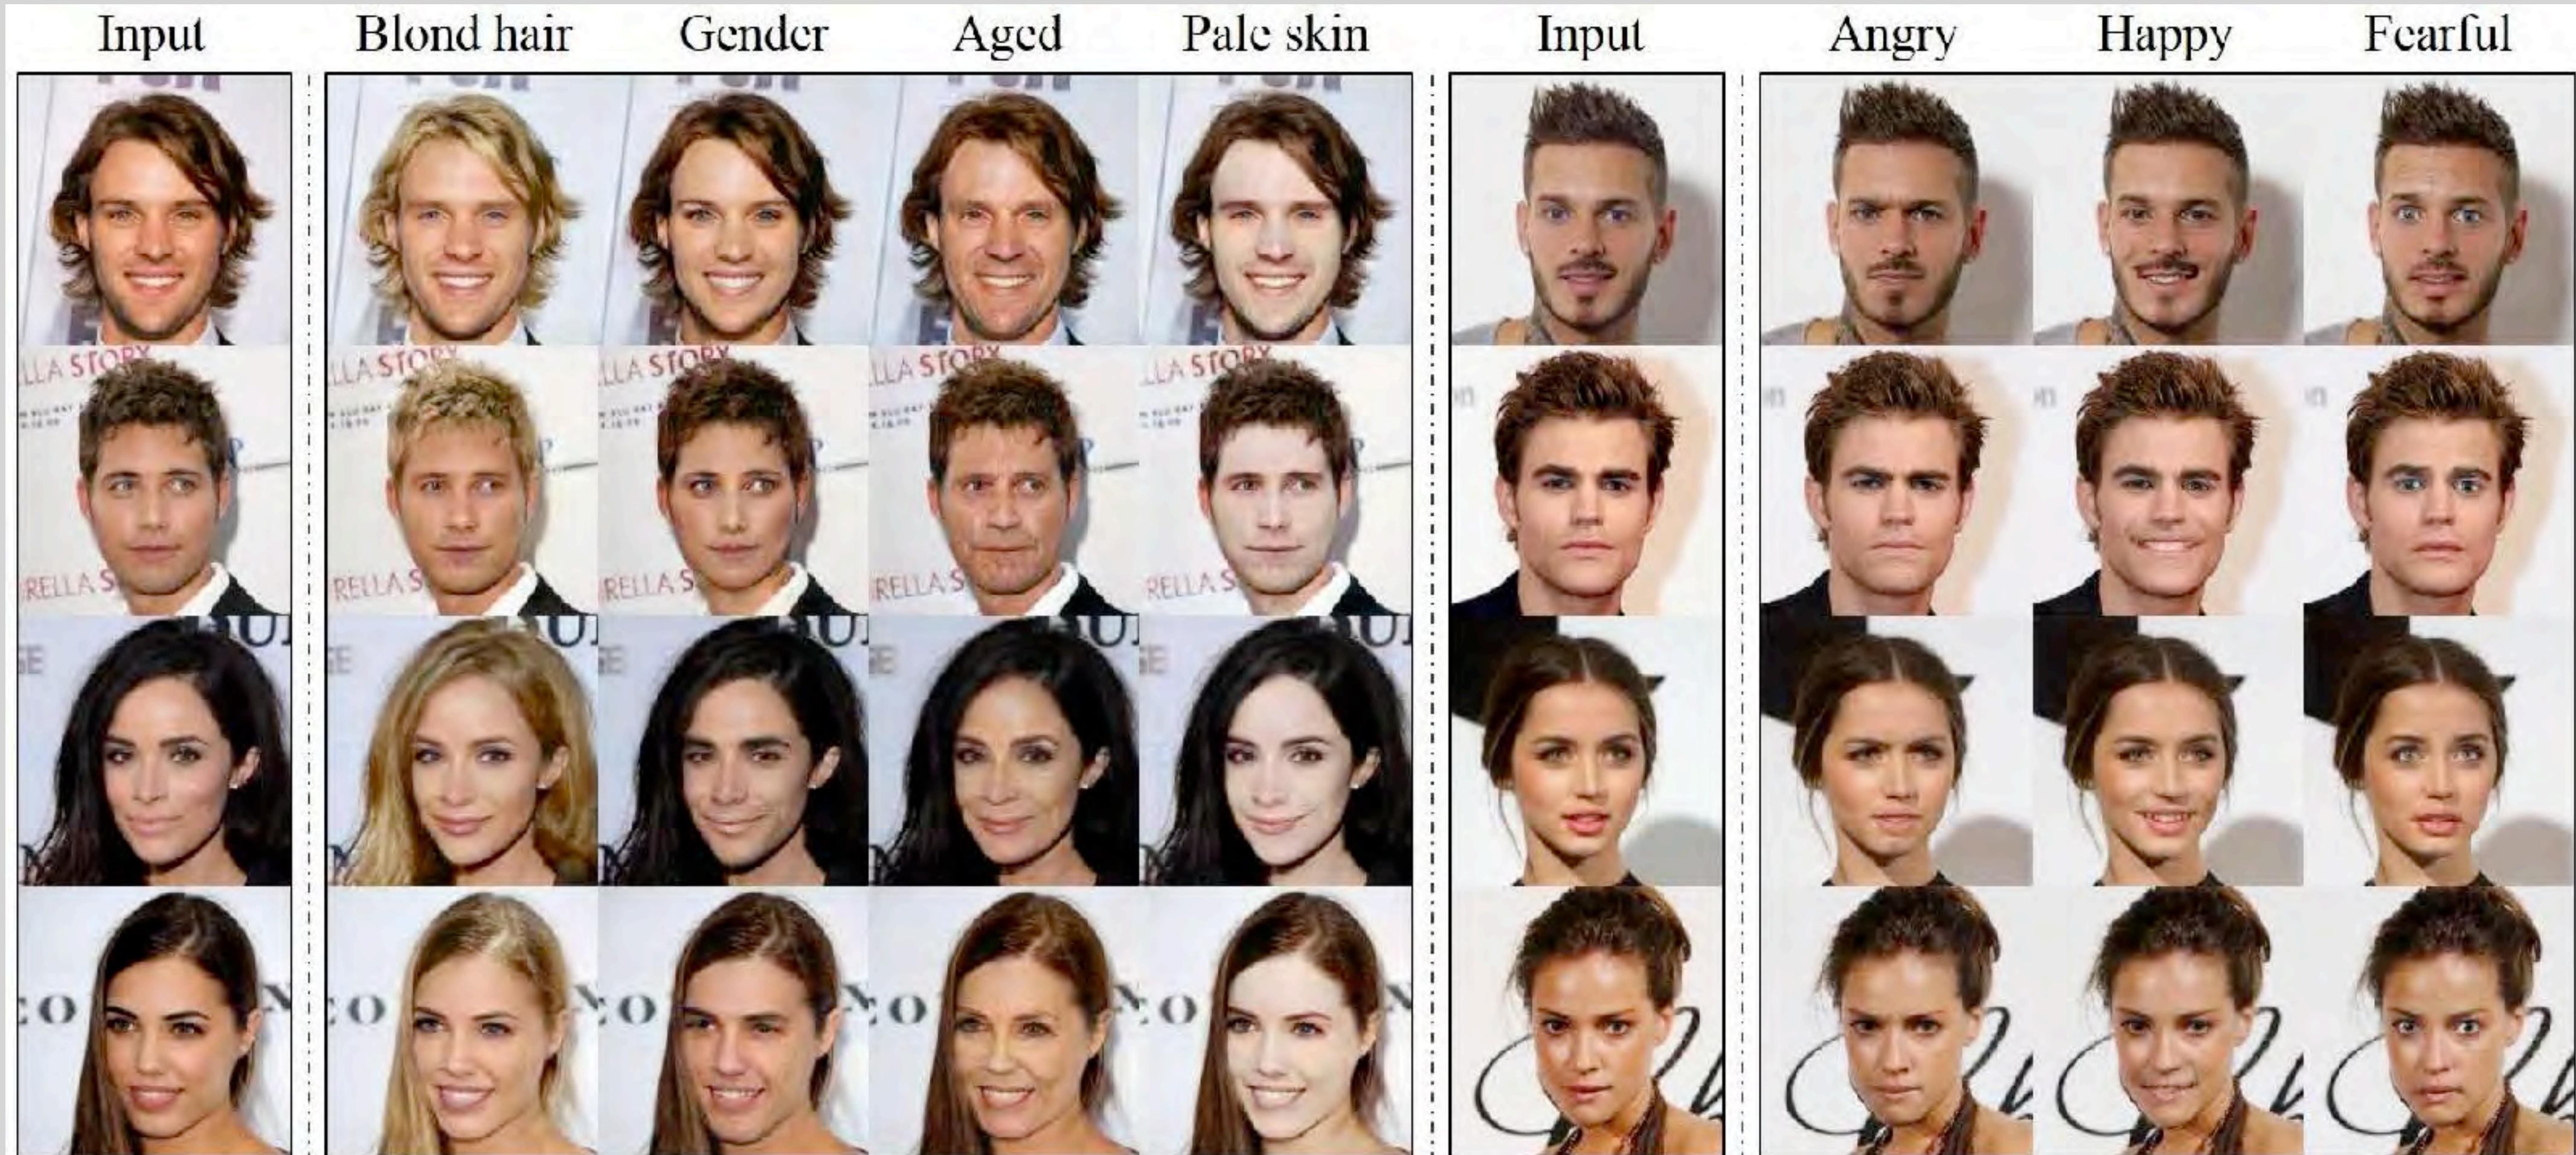
Specifically approaching this from a visual design perspective.

Generative algorithms are part of unsupervised learning techniques.

They underpin one of the **most innovative concepts in machine learning** in the past decade: **Generative Adversarial Networks (GANs)**.

A generative model can learn to mimic any distribution of it. Their potential is huge as they can be taught to recreate similar models in any domain. Some of these domains include, but are not limited, to the following:

Images
Music
Speech
Text
Videos



Example of a Deepfake (Source: StarGAN GitHub)



Combining Markov Random Fields and Convolutional Neural Networks for Image Synthesis ([Chuan Li, Michael Wand 2016](#))



Colorful Image Colorization (Zhang, Richard and Isola, Phillip and Efros, Alexei A 2016)



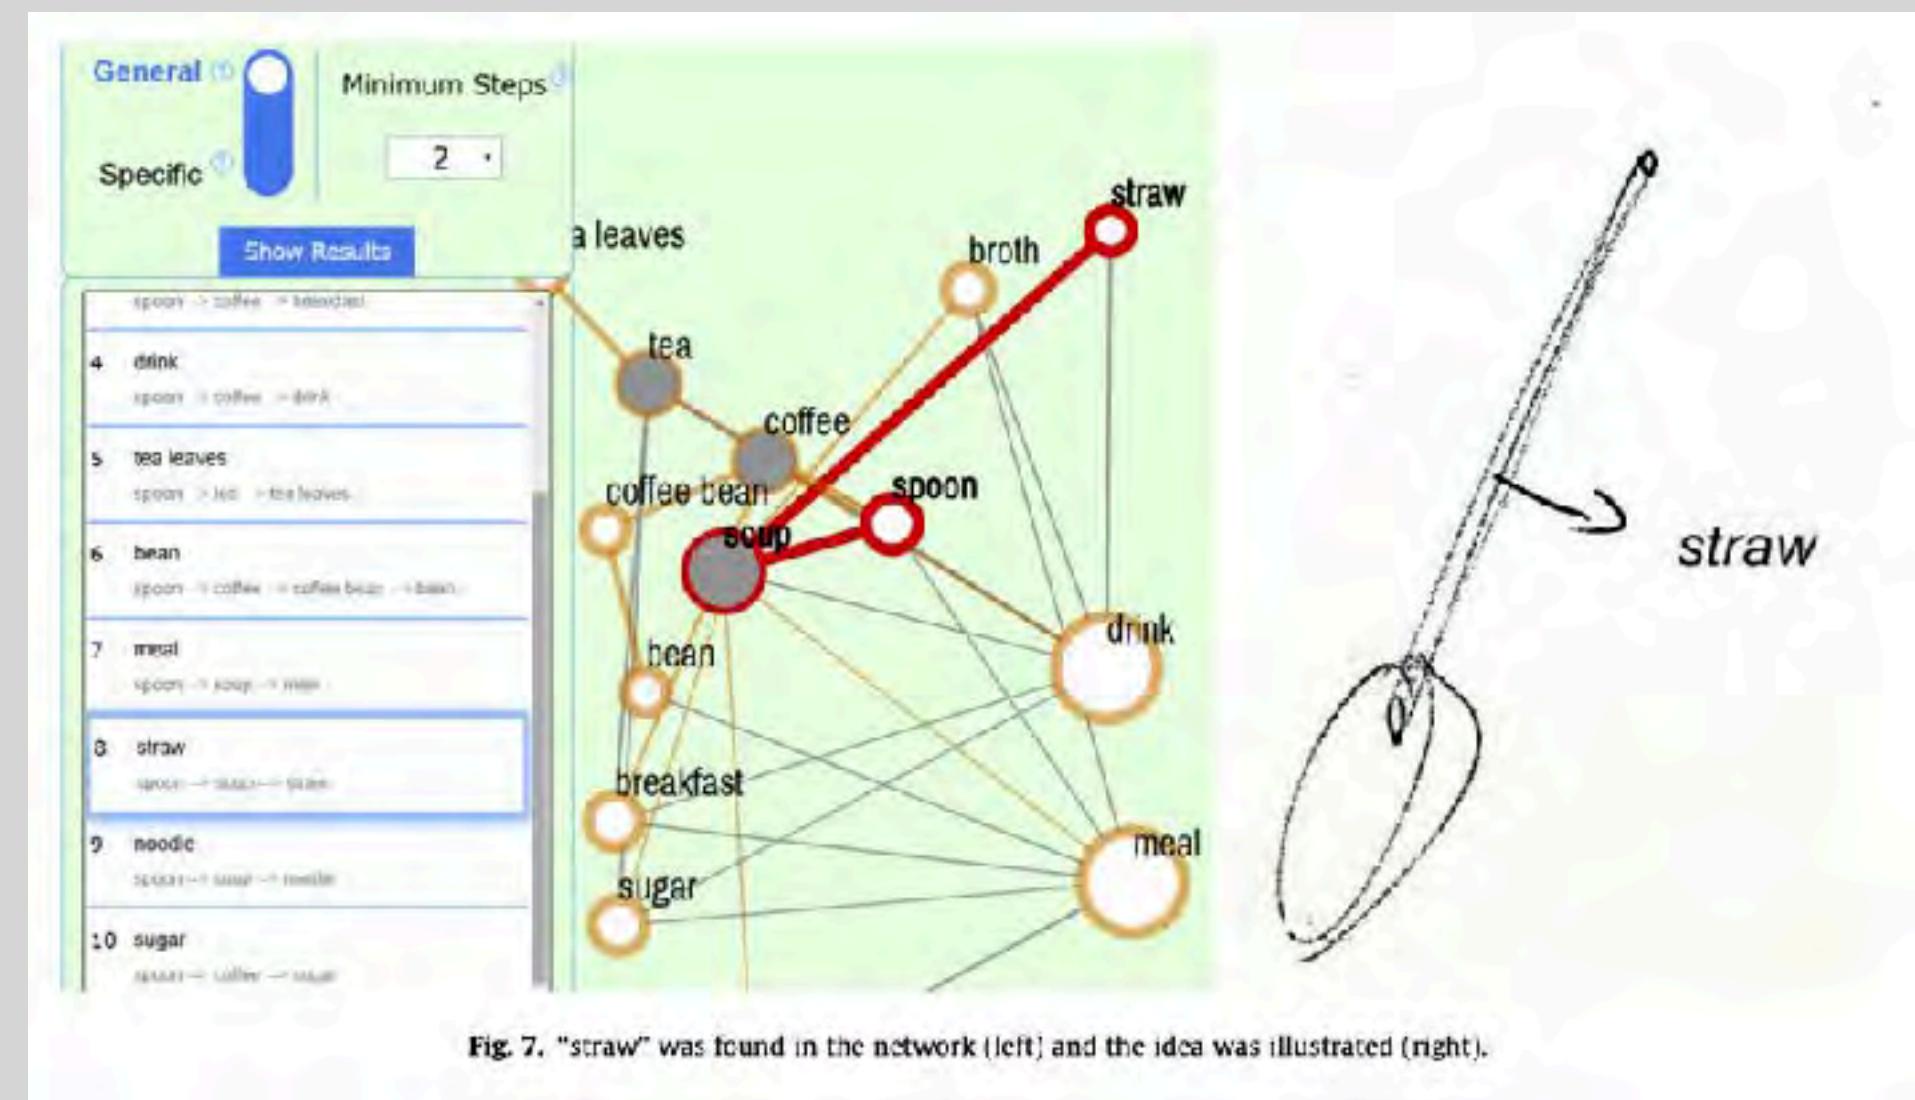
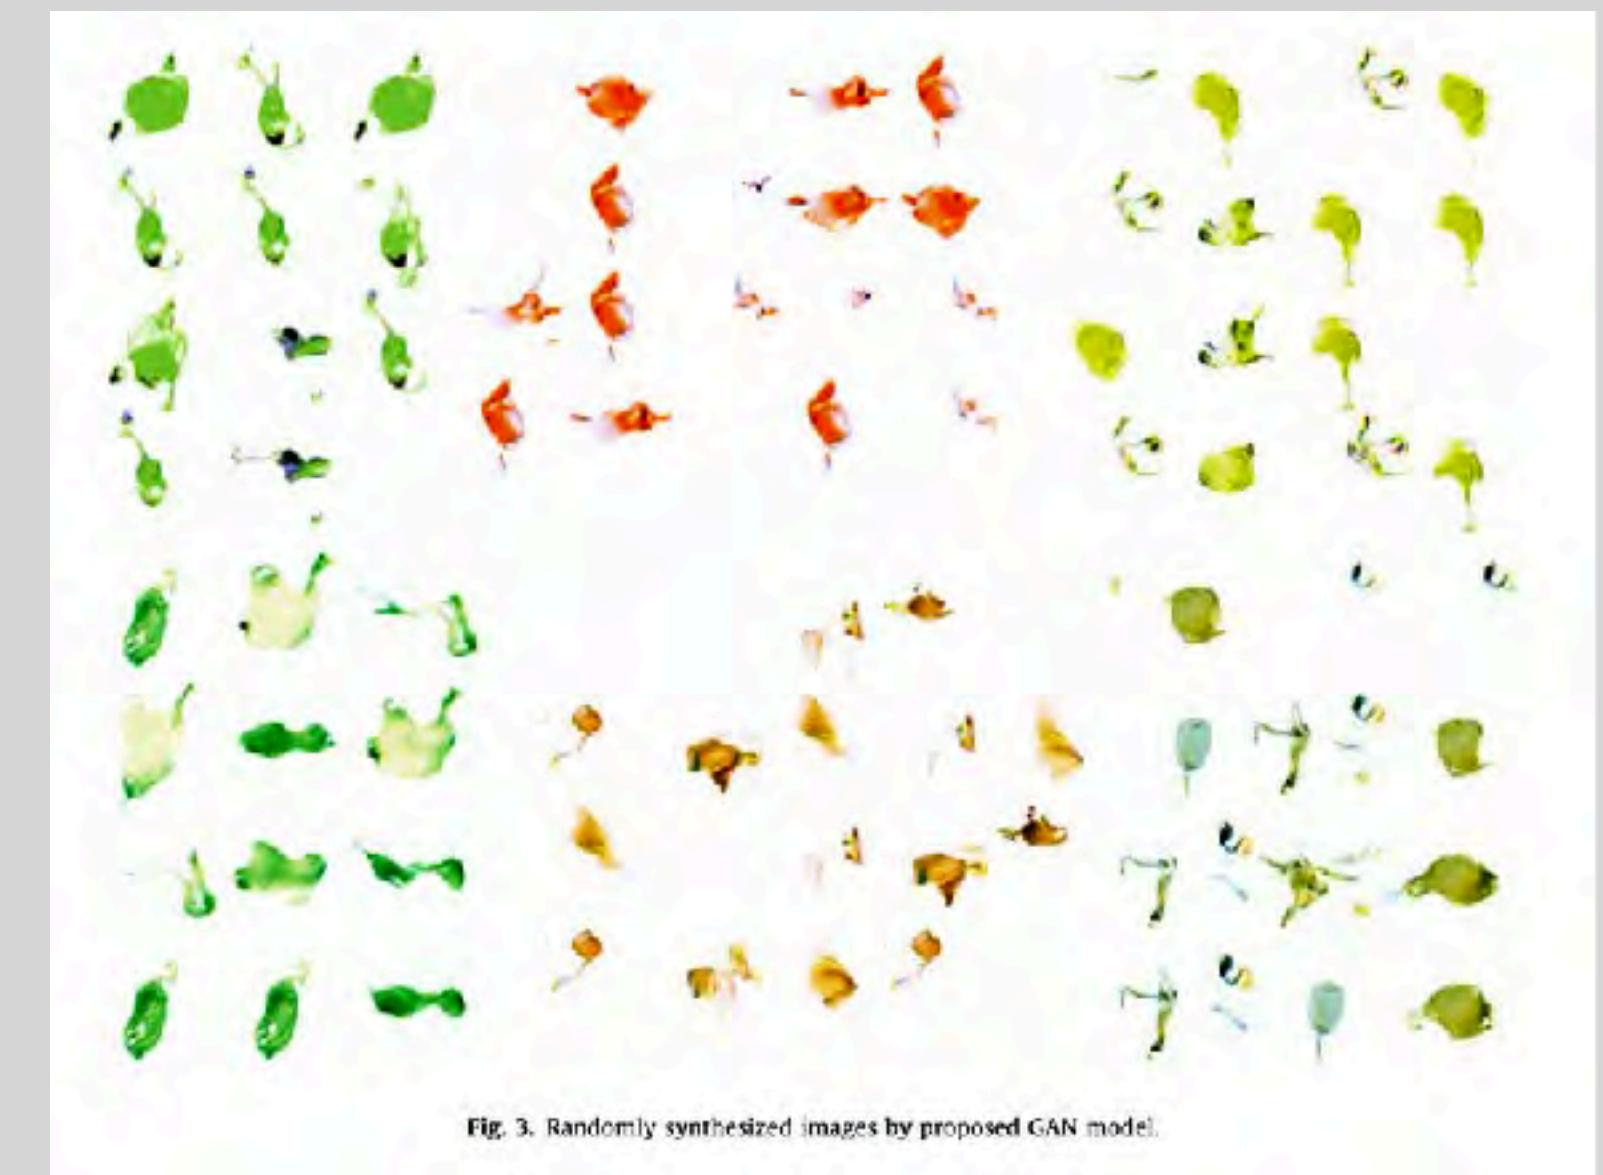
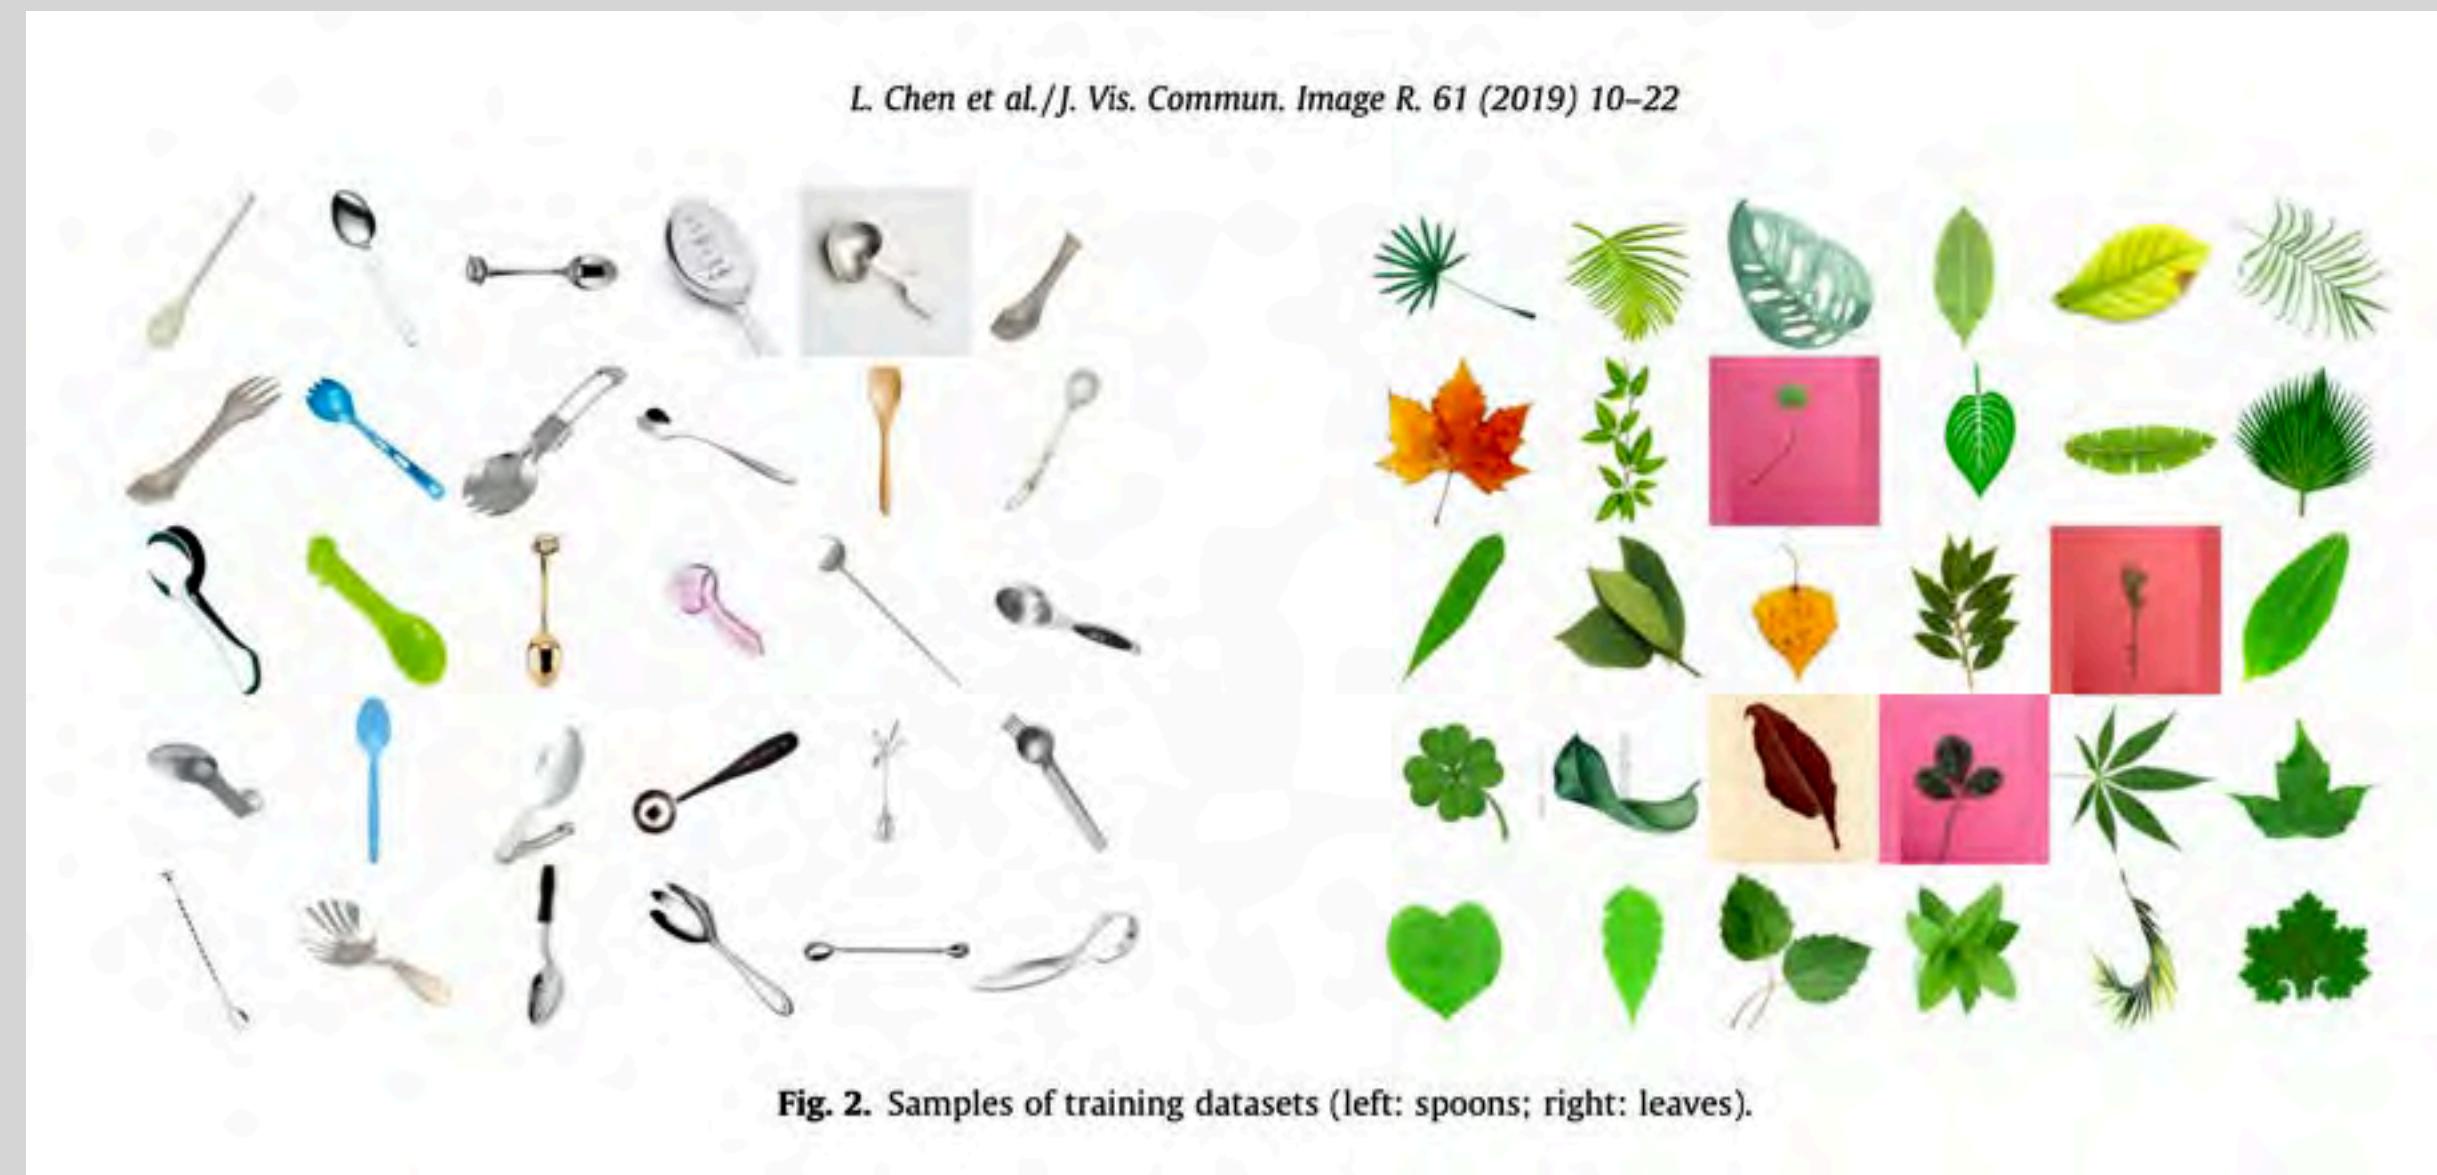
DeblurGAN: Blind Motion Deblurring Using Conditional
Adversarial Networks (Kupyn, Orest and Budzan, Volodymyr and
Mykhailych, Mykola and Mishkin, 2017)

Ai as an Assistant

The **designer** and the **machine** making **decisions together** to expand and improve the results.



Data-driven approach for design ideation, where Liuqing Chen and Pan Wan propose the use of different tools, including generative adversarial networks, to create design concepts by implementing Koesther's biosociation (Chen et al., 2019).





The DesING project, as a final example, explores the use of algorithms for creative inspiration in the area of fashion design (Sbai, 2019).

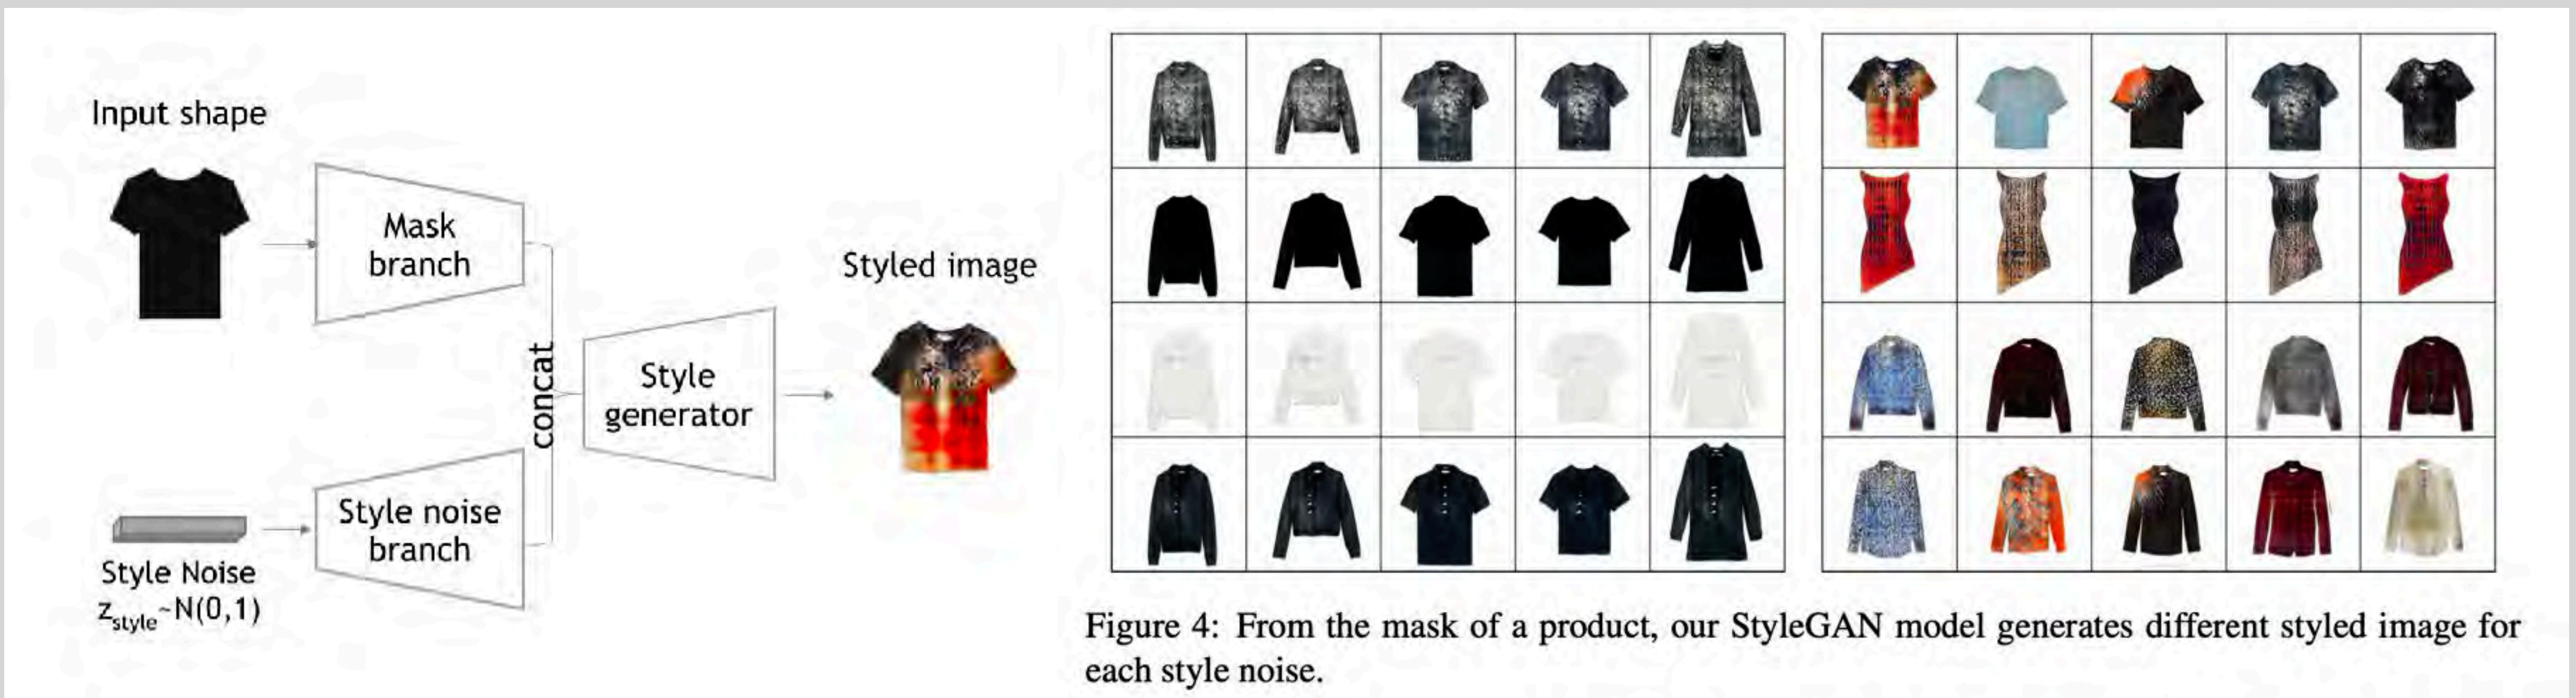
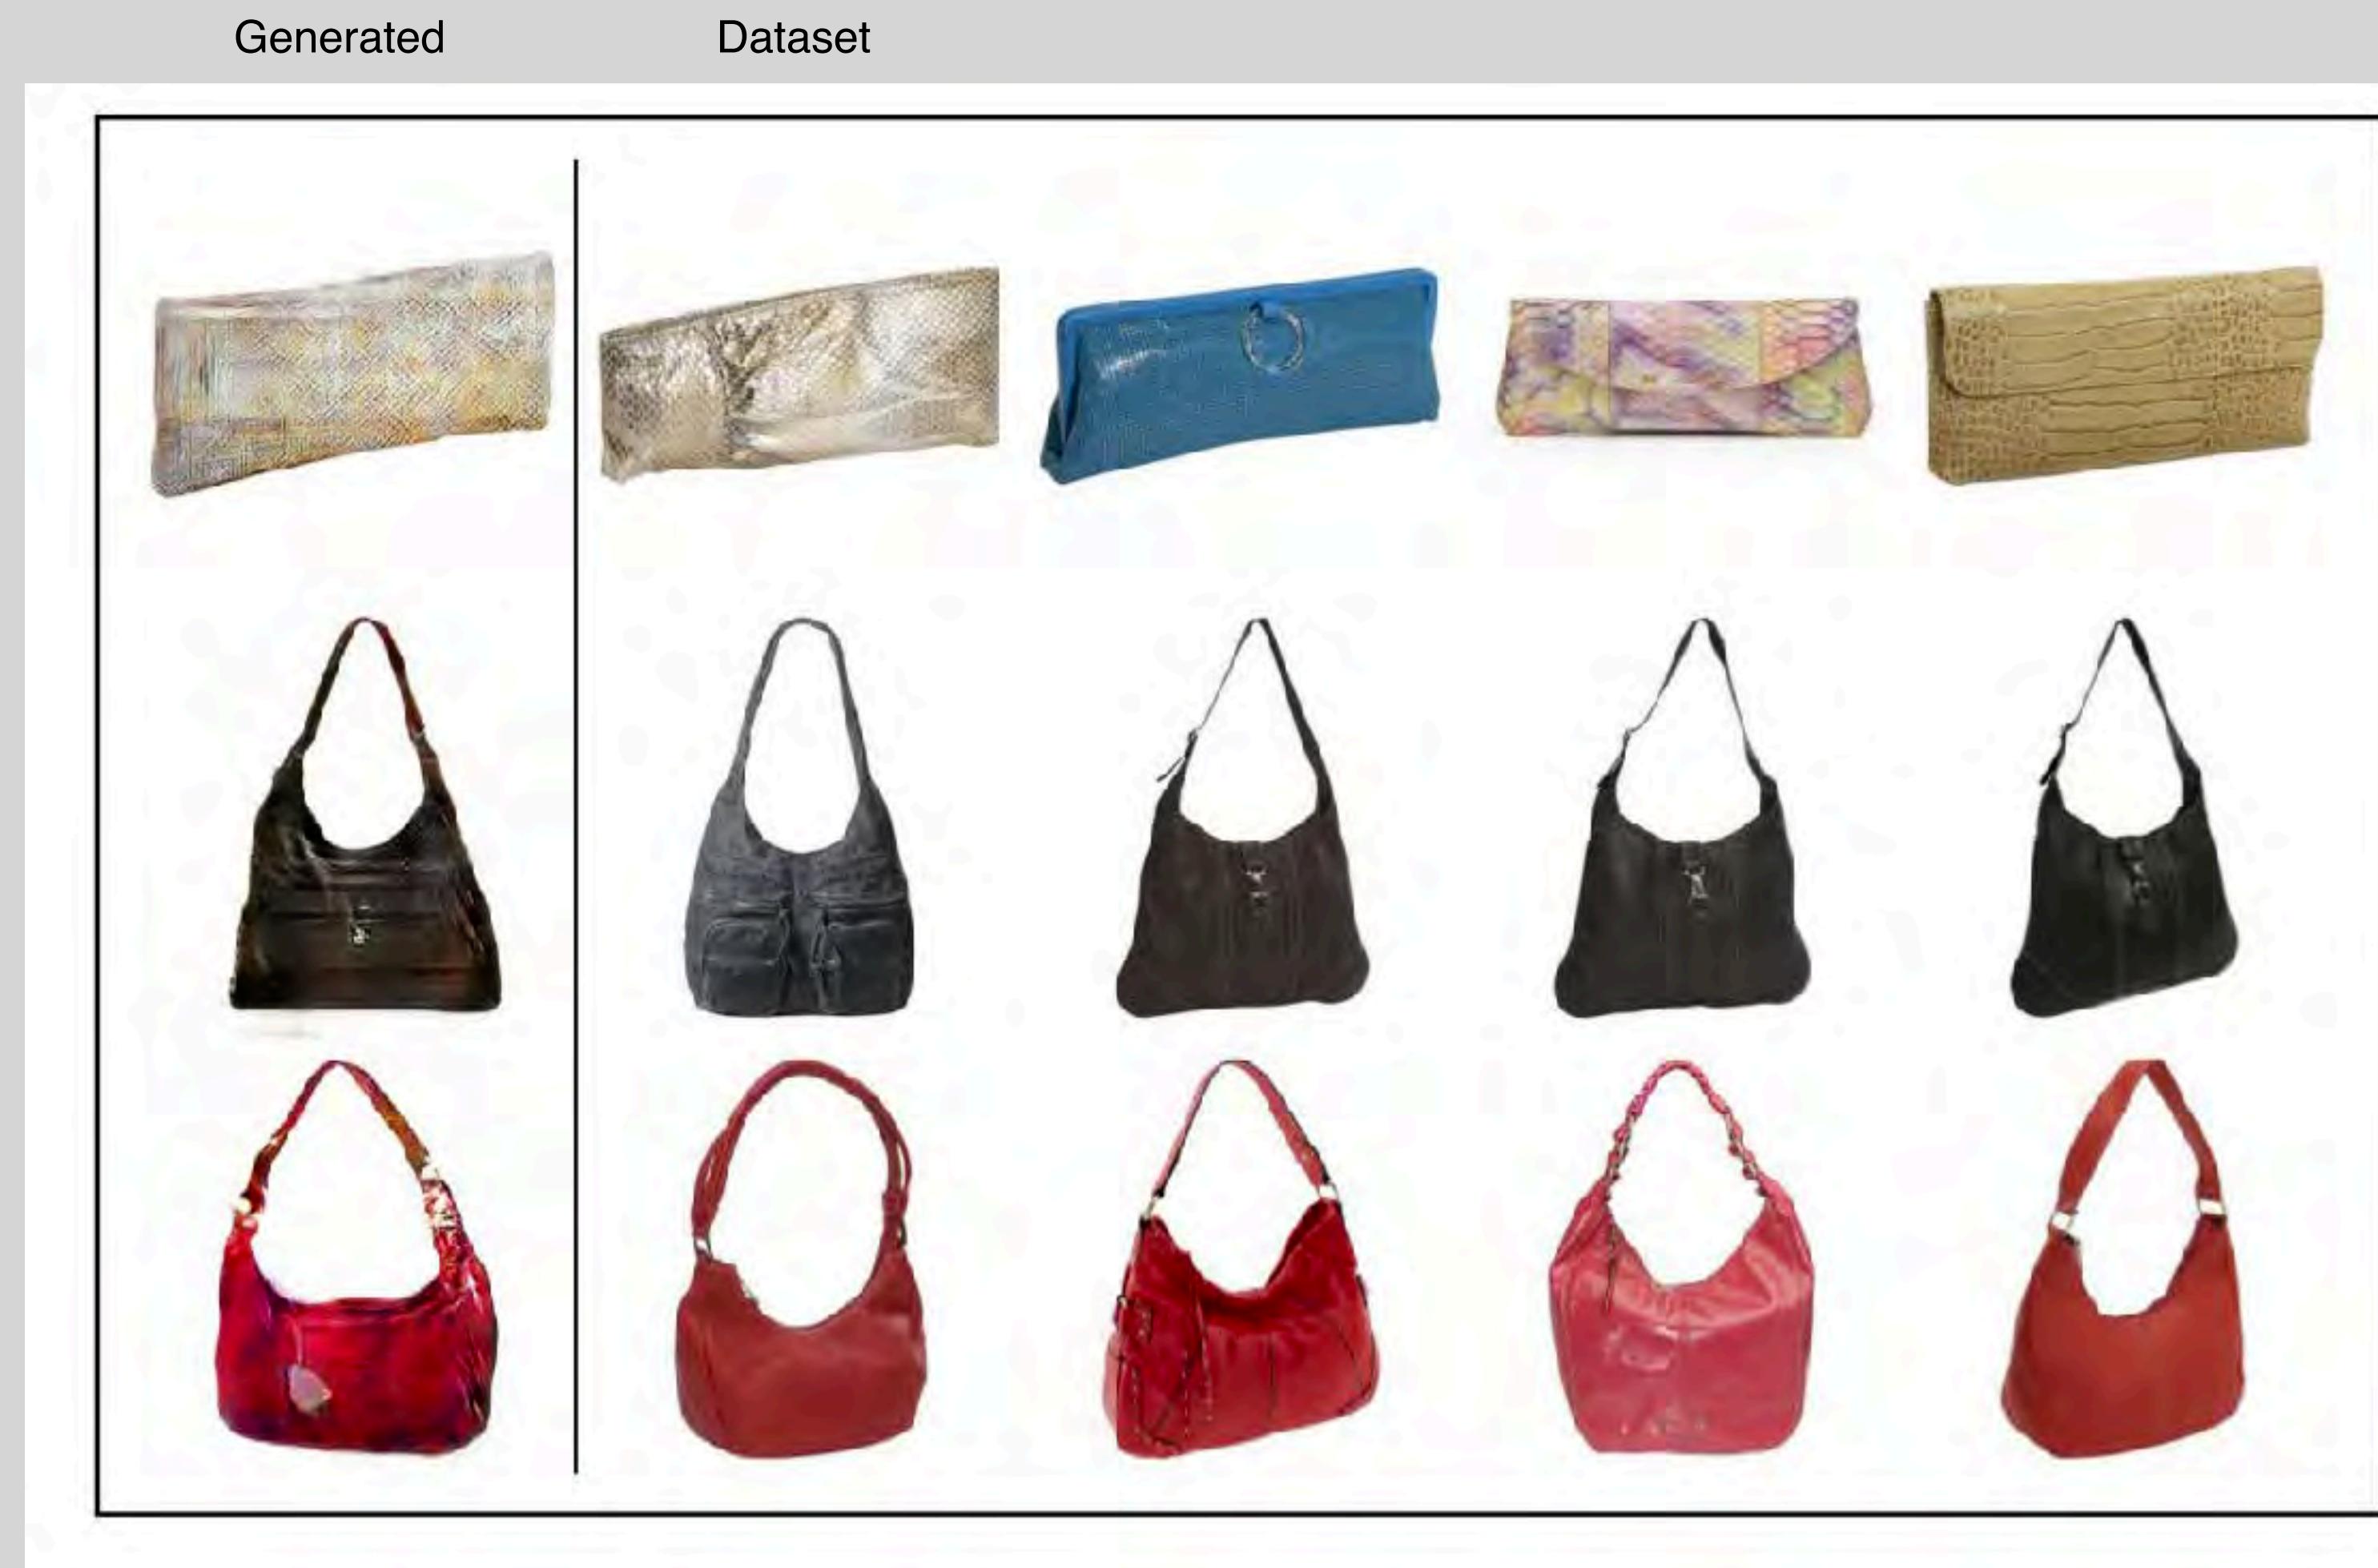


Figure 4: From the mask of a product, our StyleGAN model generates different styled image for each style noise.

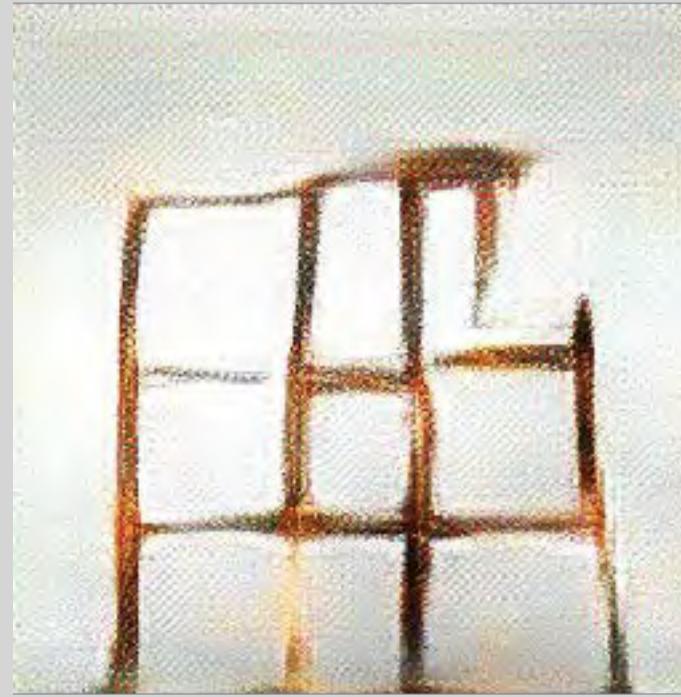
The DesING project, as a final example, explores the use of algorithms for creative inspiration in the area of fashion design (Sbai, 2019).



The DesING project, as a final example, explores the use of algorithms for creative inspiration in the area of fashion design (Sbai, 2019).



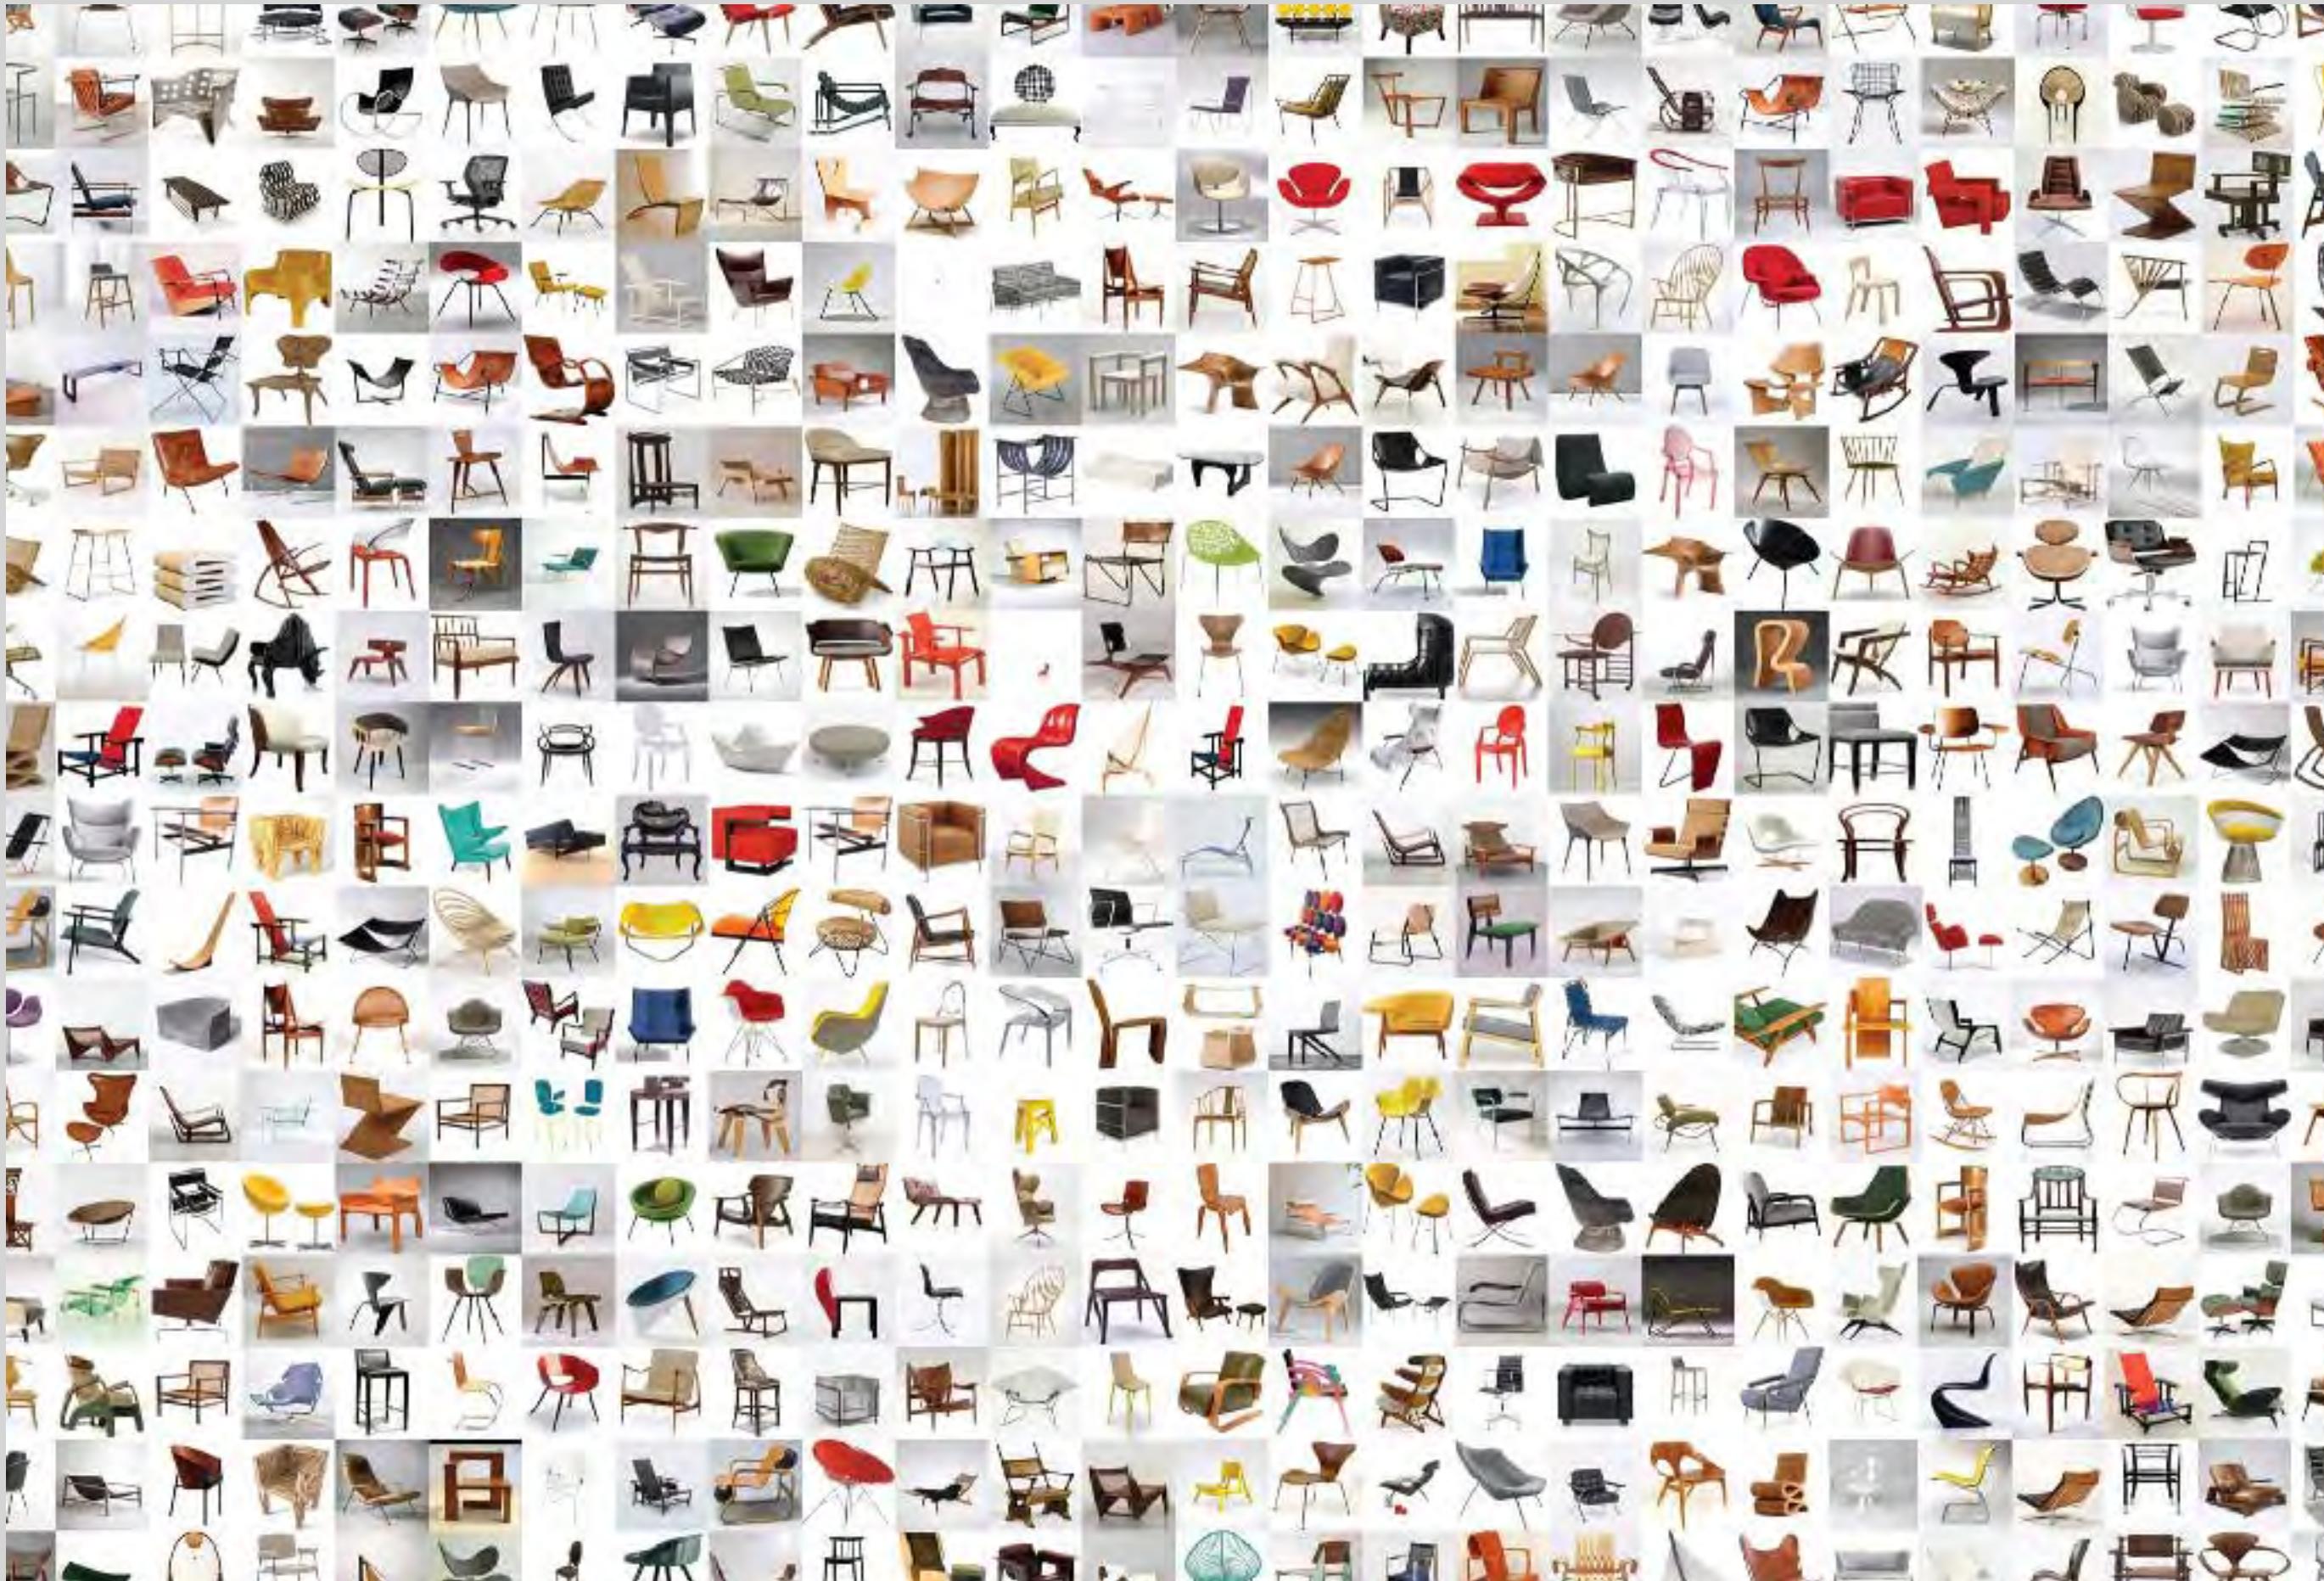
ChAIr Project (Schmitt, 2018), which uses AI to generate chair shapes that designers can reinterpret and then build.



ChAIr Project (Schmitt, 2018), which uses AI to generate chair shapes that designers can reinterpret and then build.



ChAIr Project (Schmitt, 2018), which uses AI to generate chair shapes that designers can reinterpret and then build.



ChAIr Project (Schmitt, 2018), which uses AI to generate chair shapes that designers can reinterpret and then build.



ChAIr Project (Schmitt, 2018), which uses AI to generate chair shapes that designers can reinterpret and then build.



ChAIr Project (Schmitt, 2018), which uses AI to generate chair shapes that designers can reinterpret and then build.



ChAIr Project (Schmitt, 2018), which uses AI to generate chair shapes that designers can reinterpret and then build.

Generative Adversarial Networks / GAN

Is a deep learning framework in which two models, a generative model G and a discriminative model D, are trained simultaneously.

GAN STORY

“Catch me if you can”

Leonardo Dicaprio
As a the “Generator”

Tom Hanks
As a the “Discriminator”

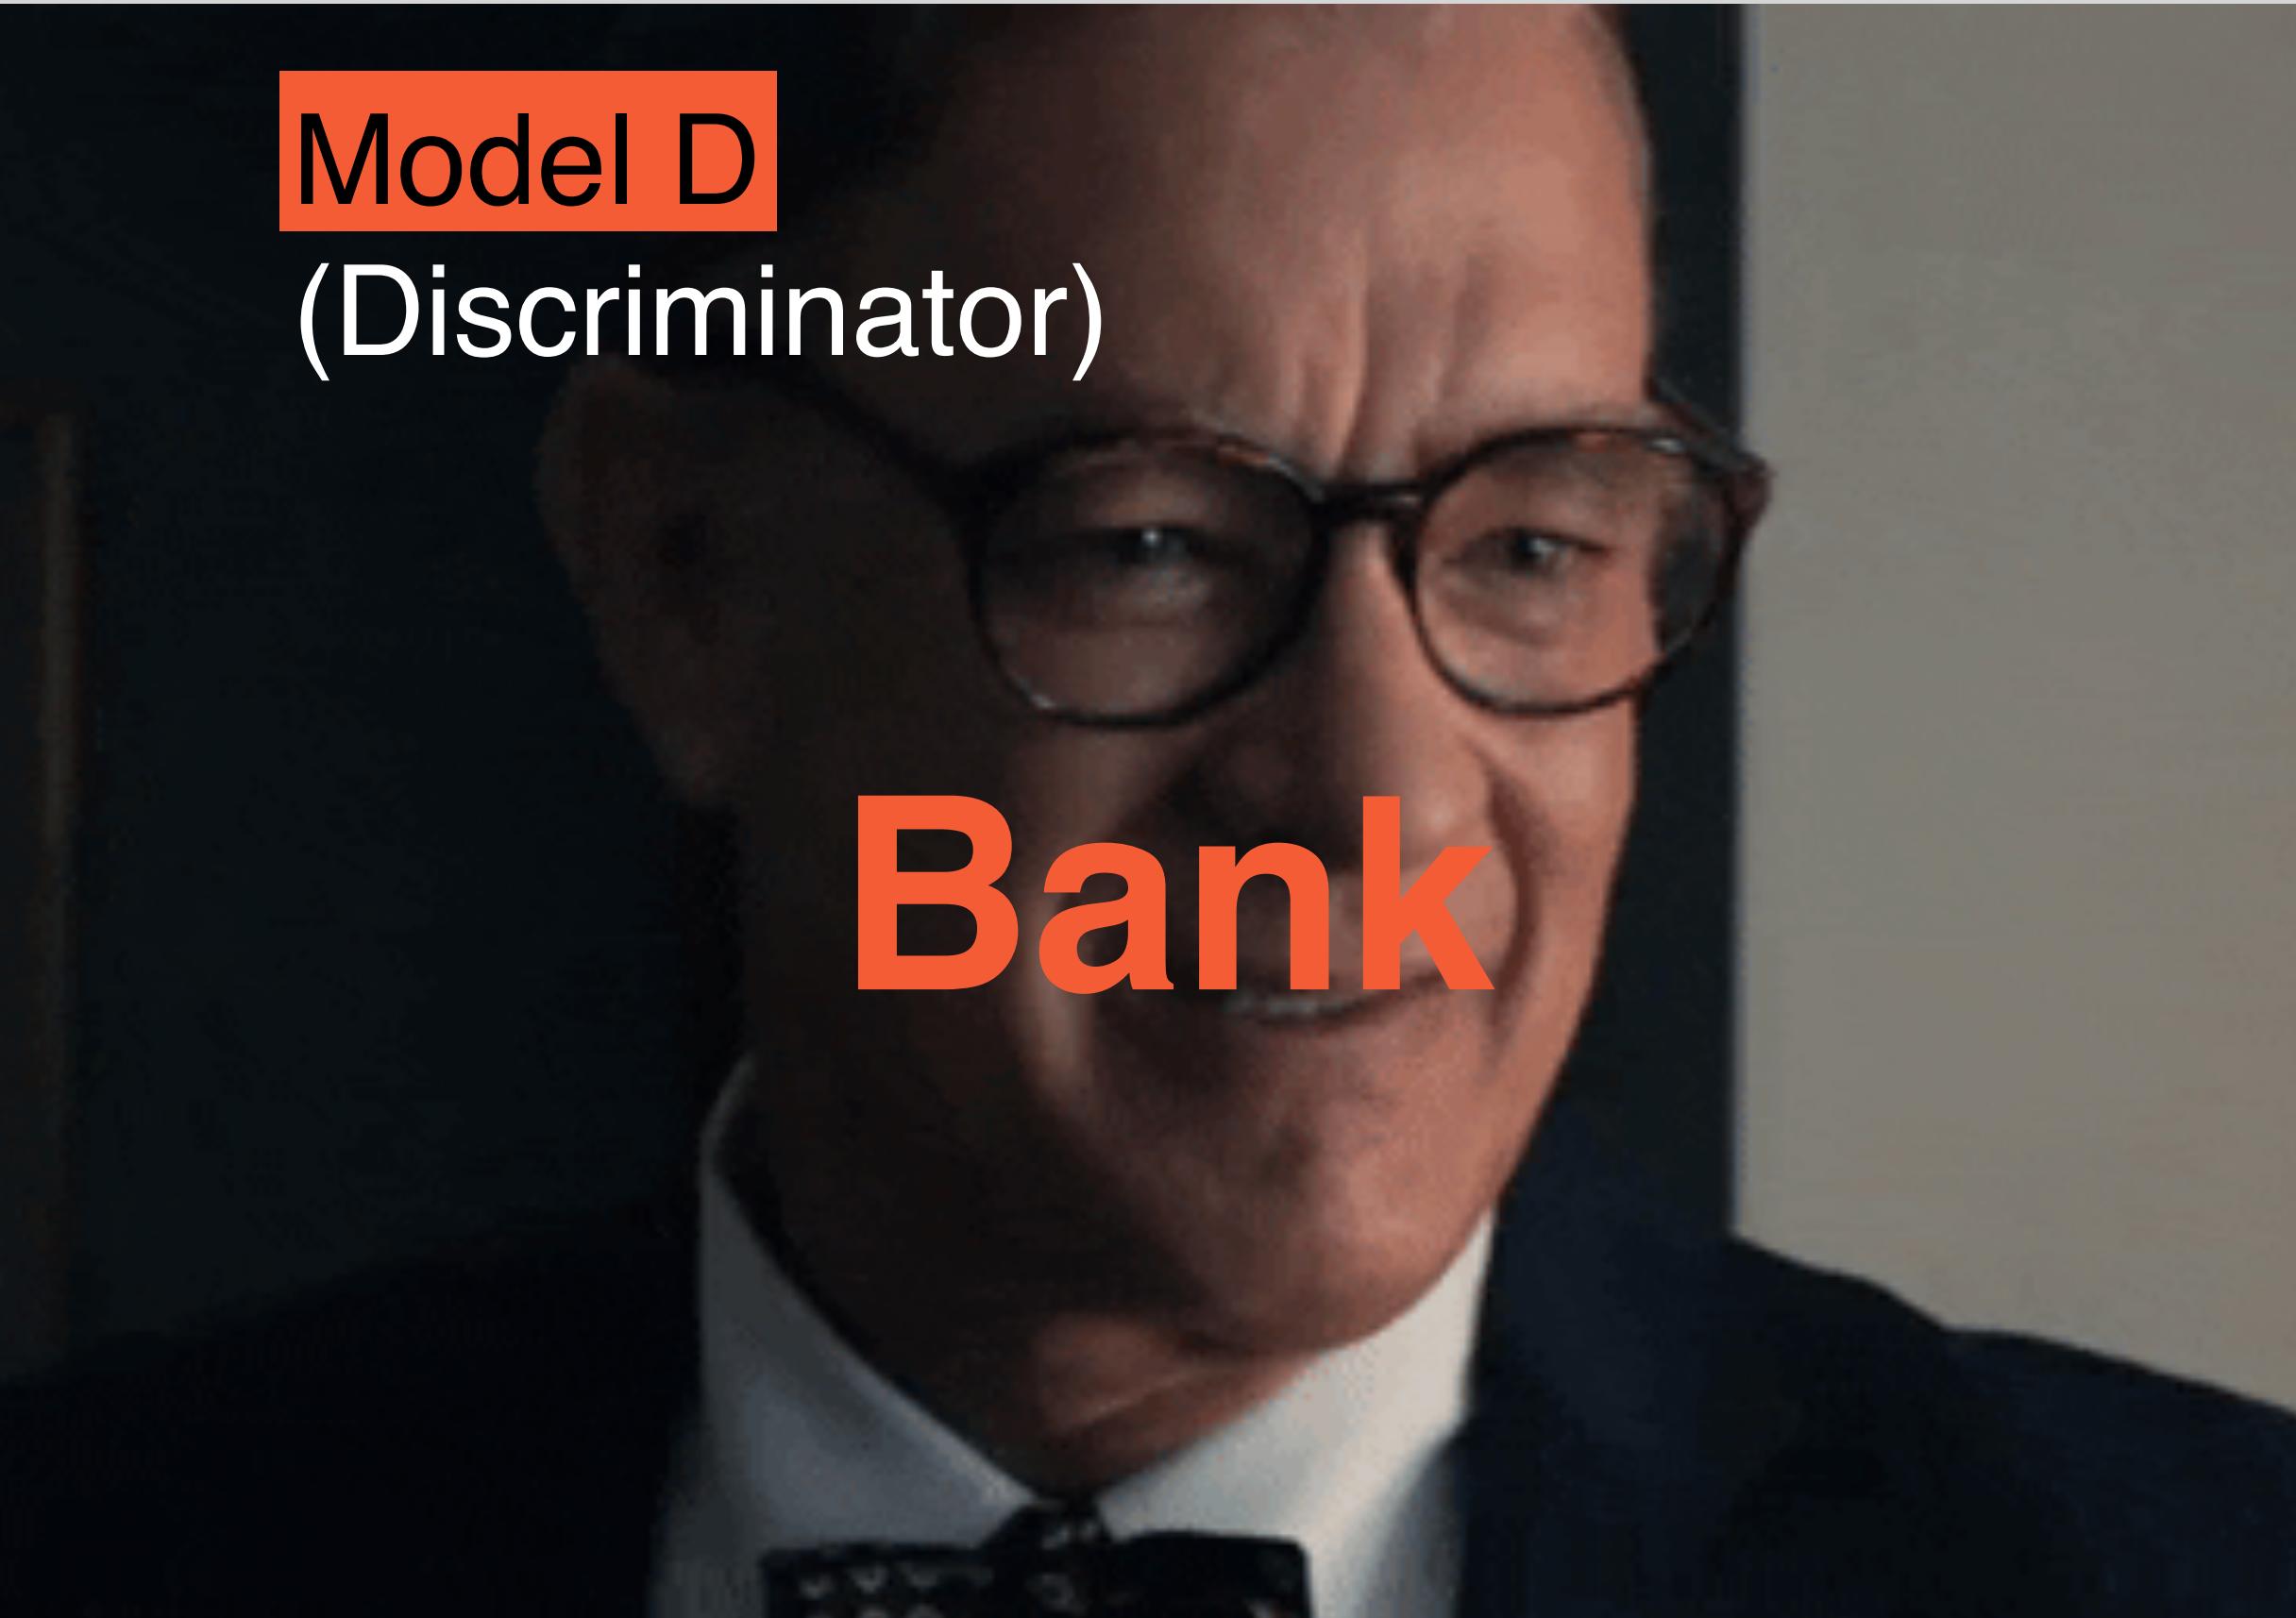
The characters

Model D
(Discriminator)

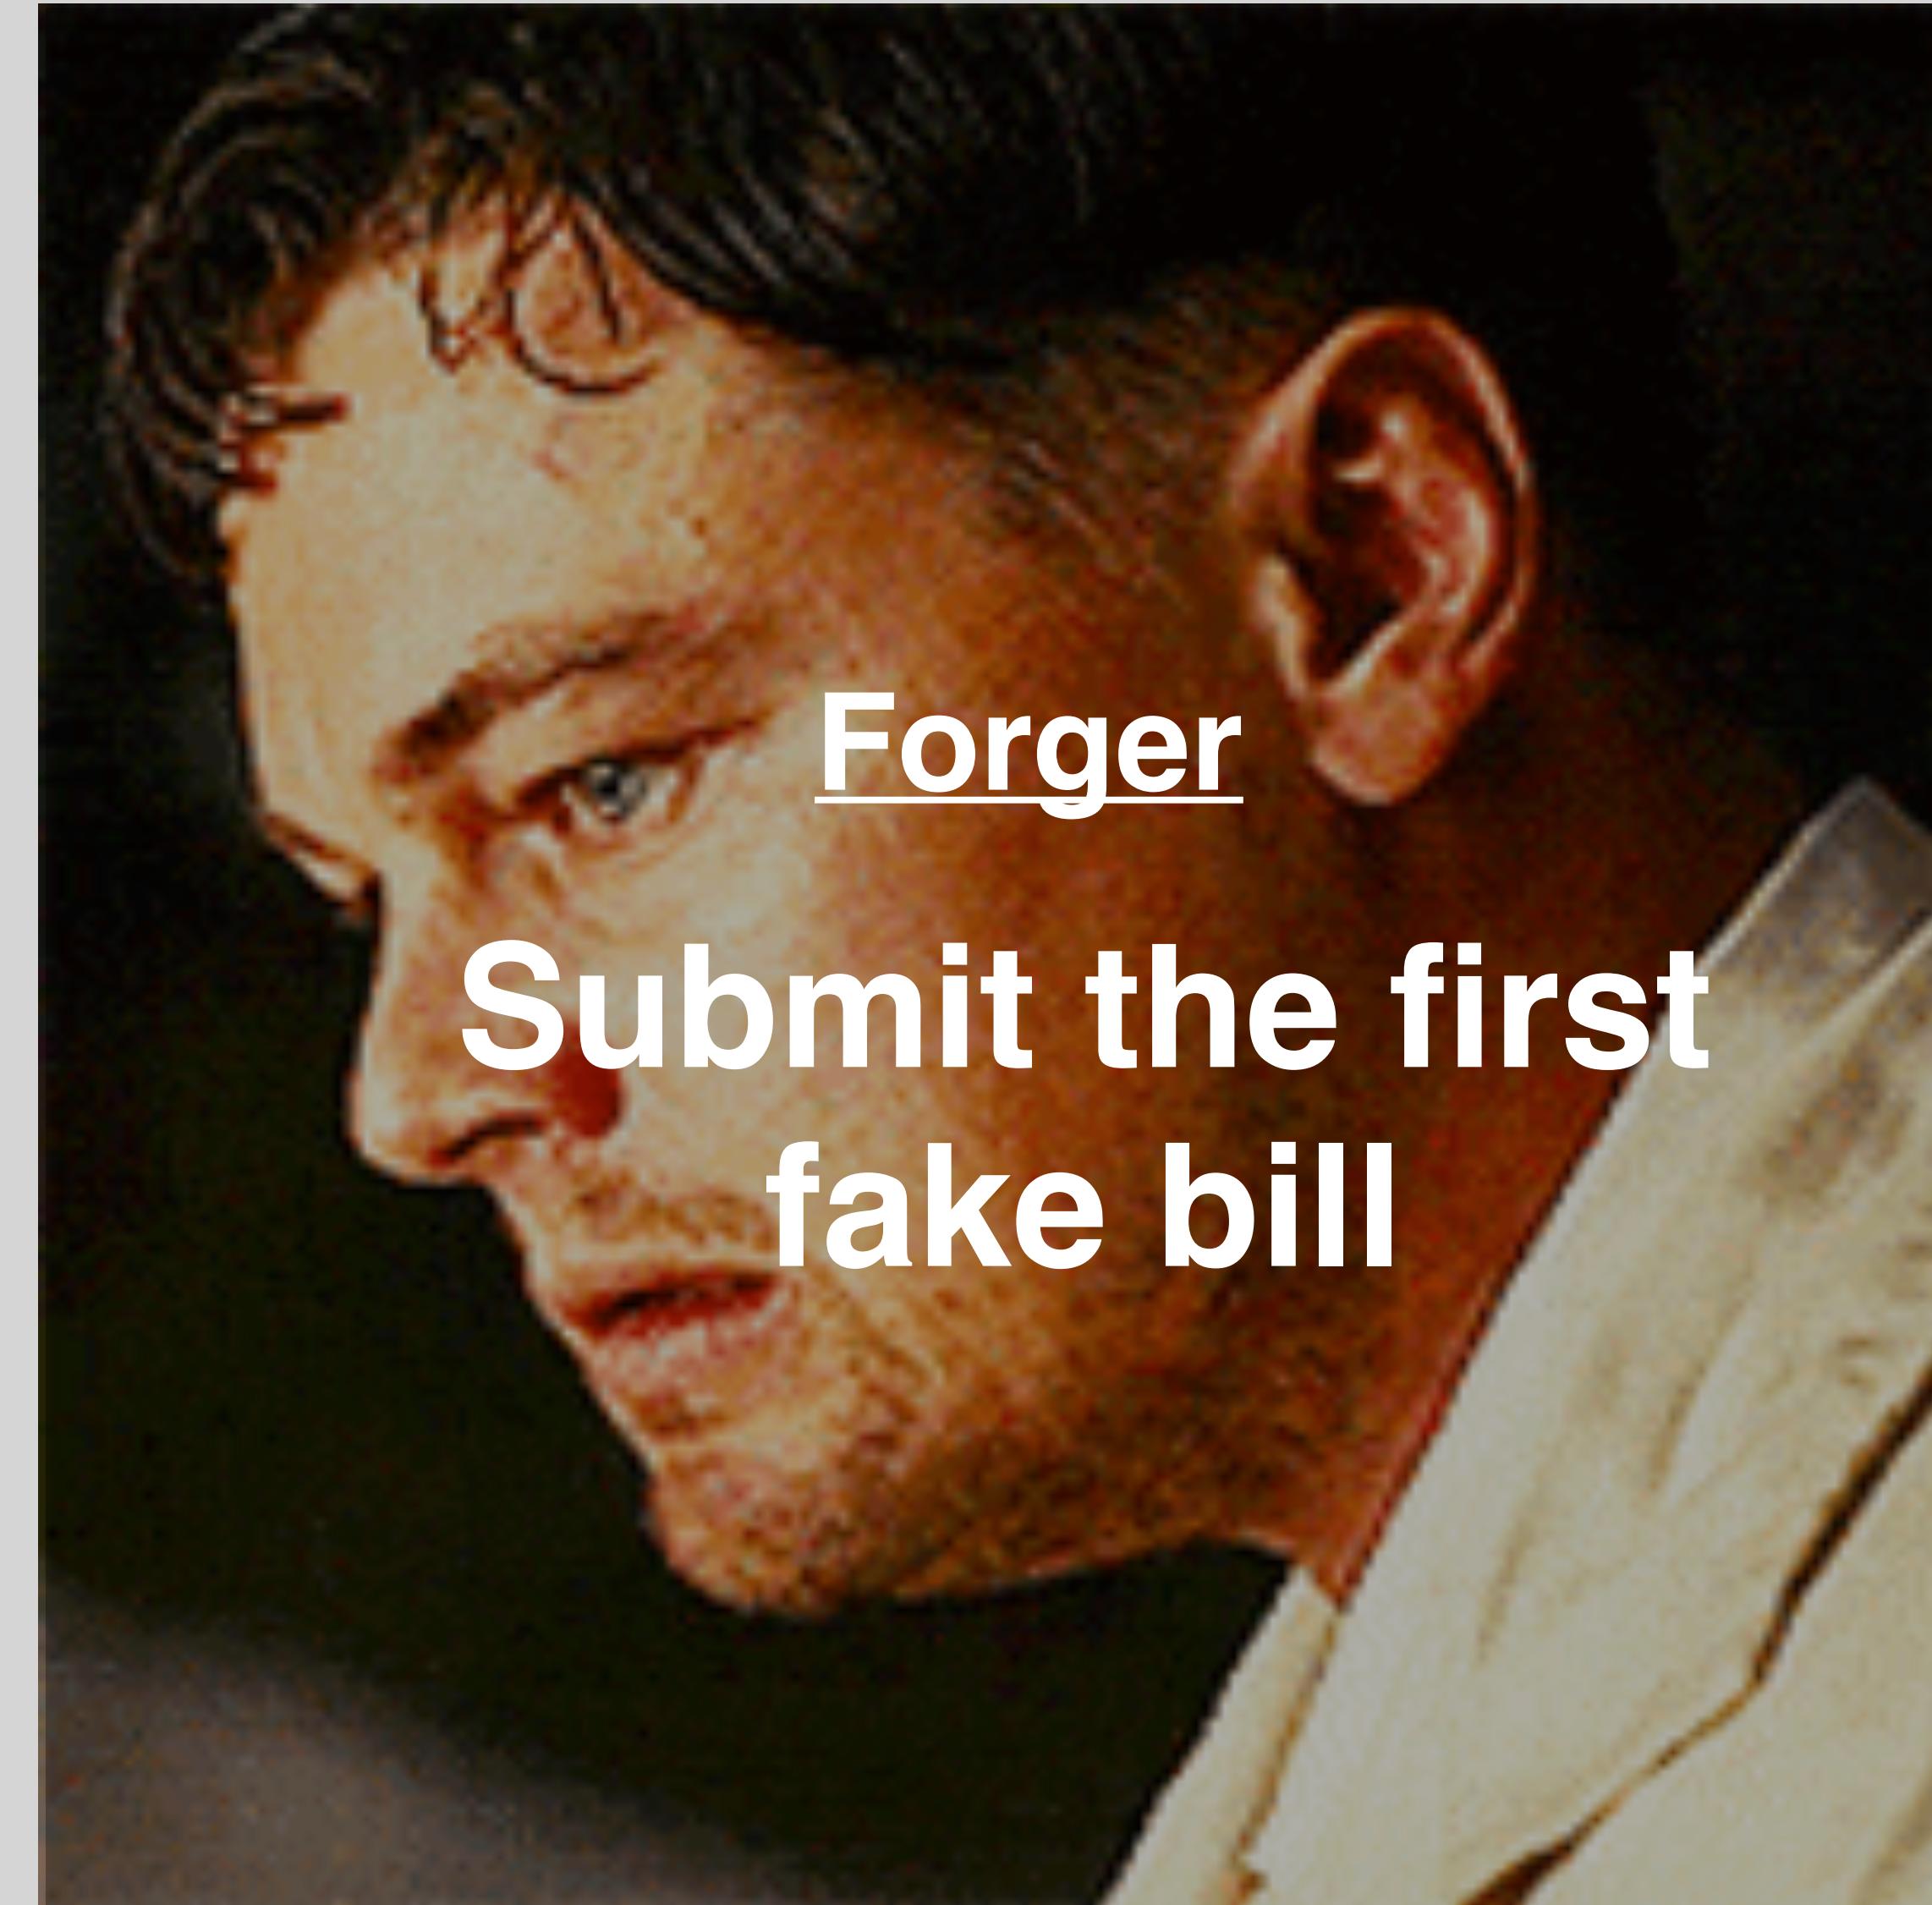
Bank

Model G
(Generator)

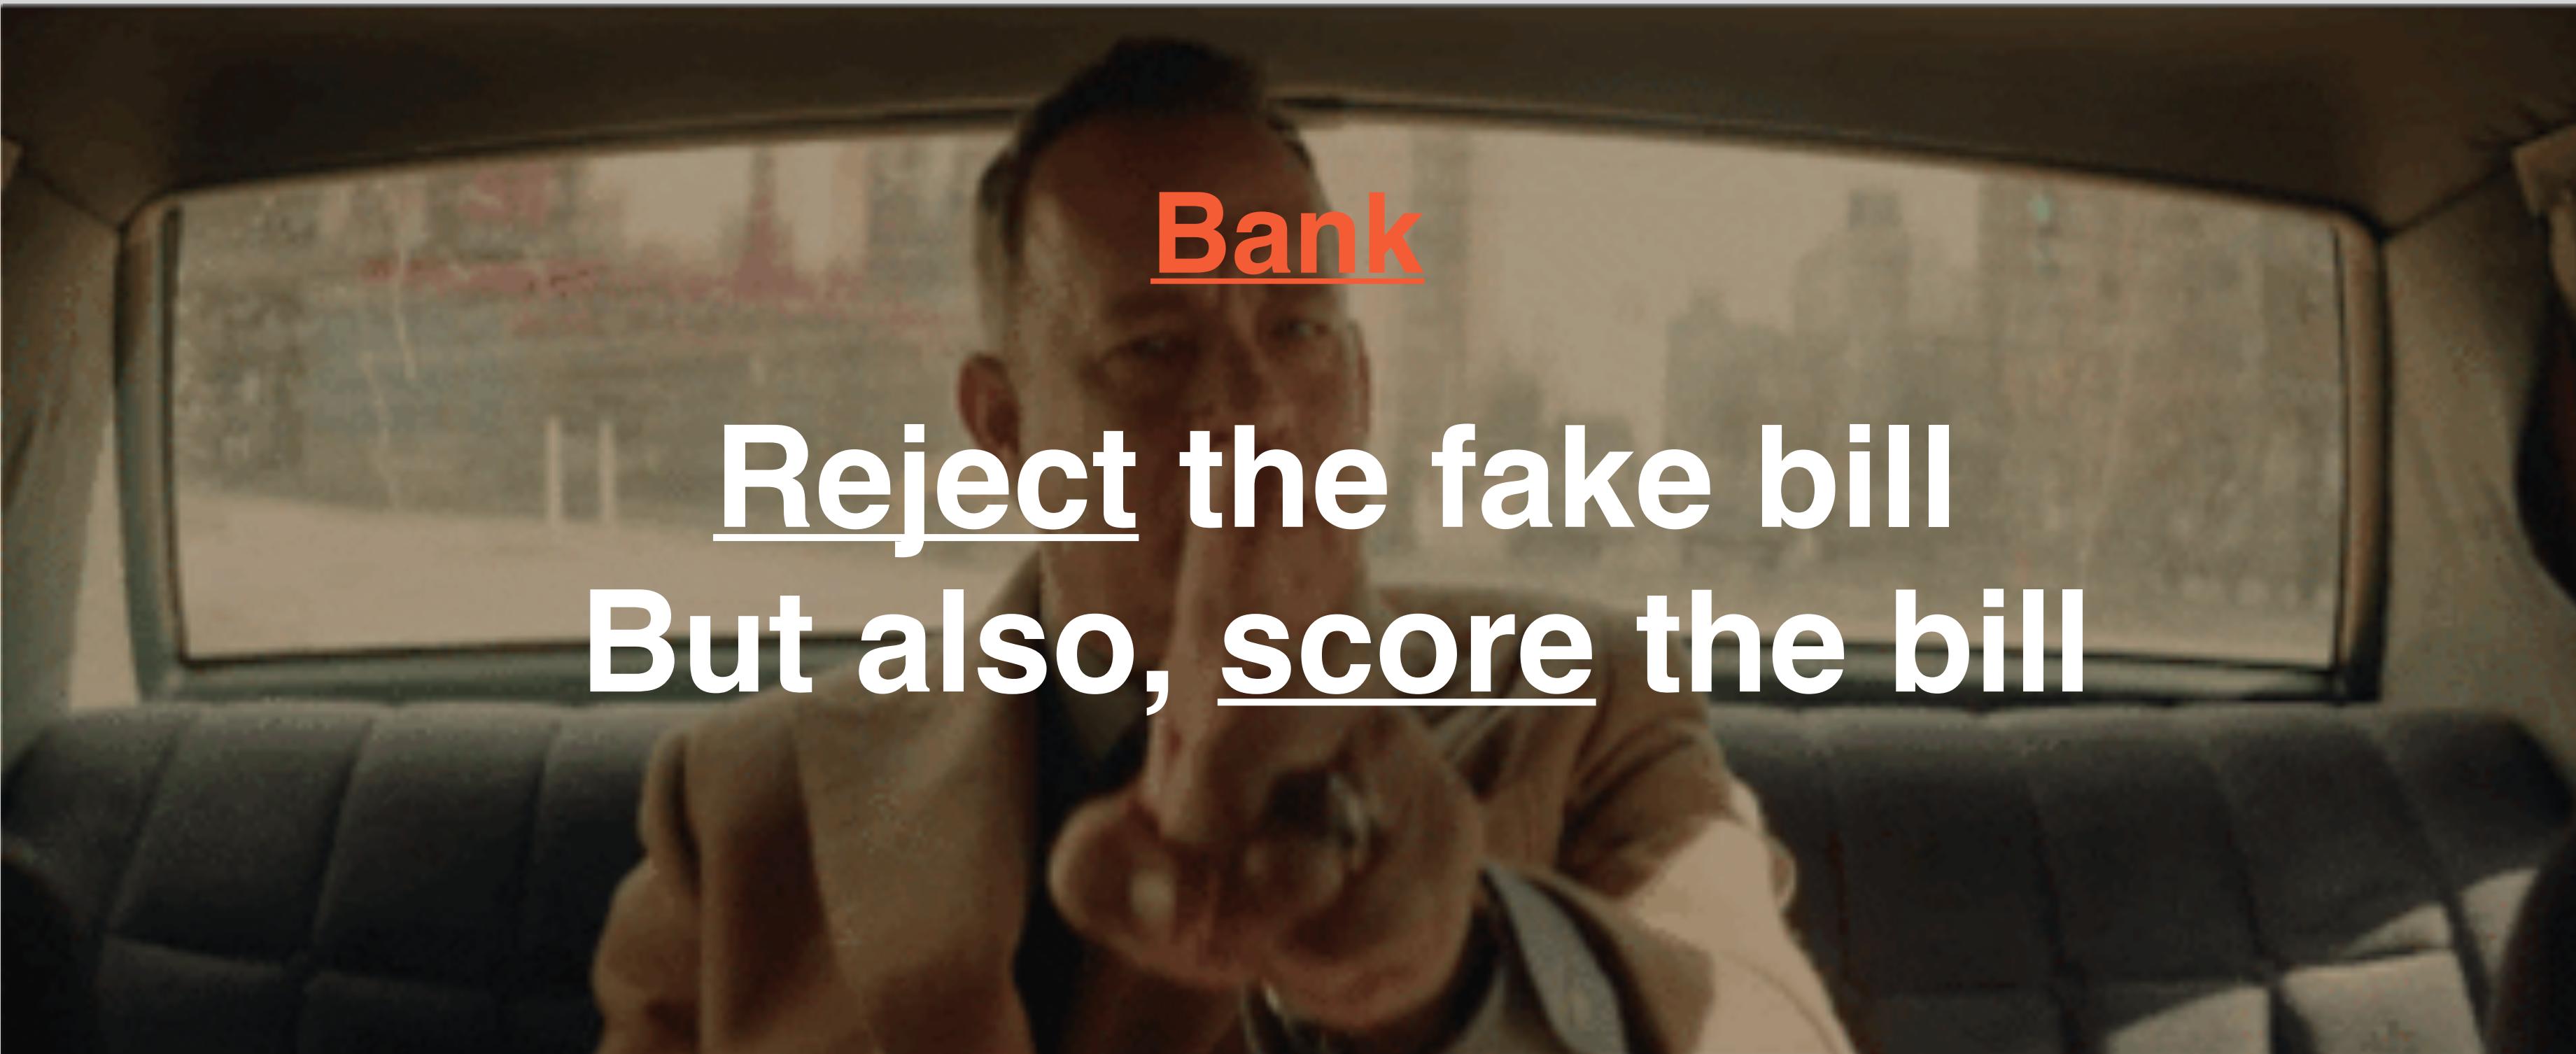
Forger



Model G (Generator)



Model D (Discriminator)



Model G (Generator)



Model G (Generator)



Model D (Discriminator)







Horse/zebra image translation using a pre-trained DC-GAN

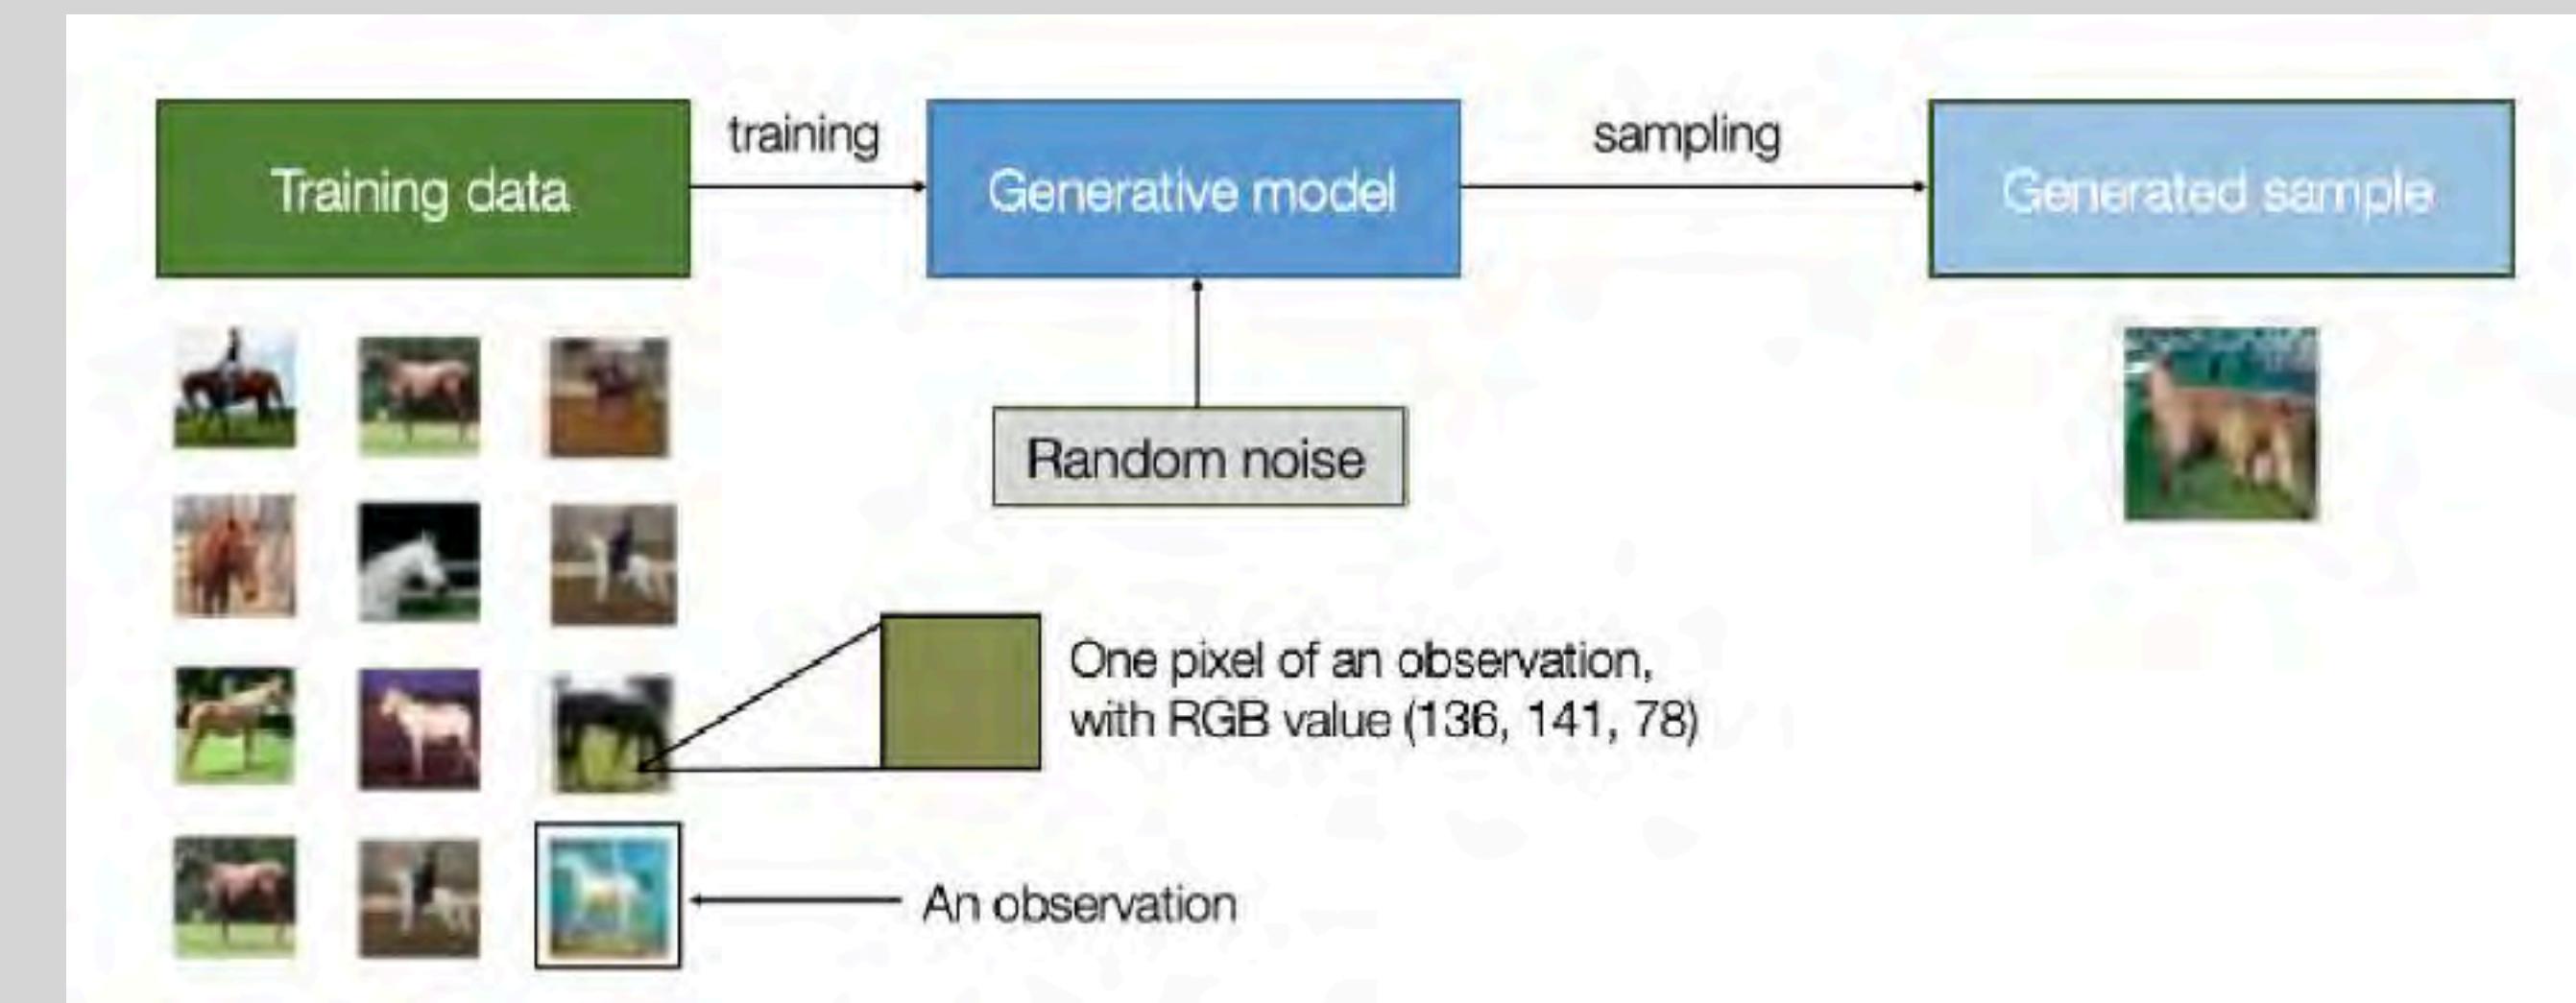
We shall first look at what it means to say that a model is **generative** and learn how it differs from the more widely studied **discriminative modeling**.

Generative Modeling

By sampling from this model, we are able to generate new data.

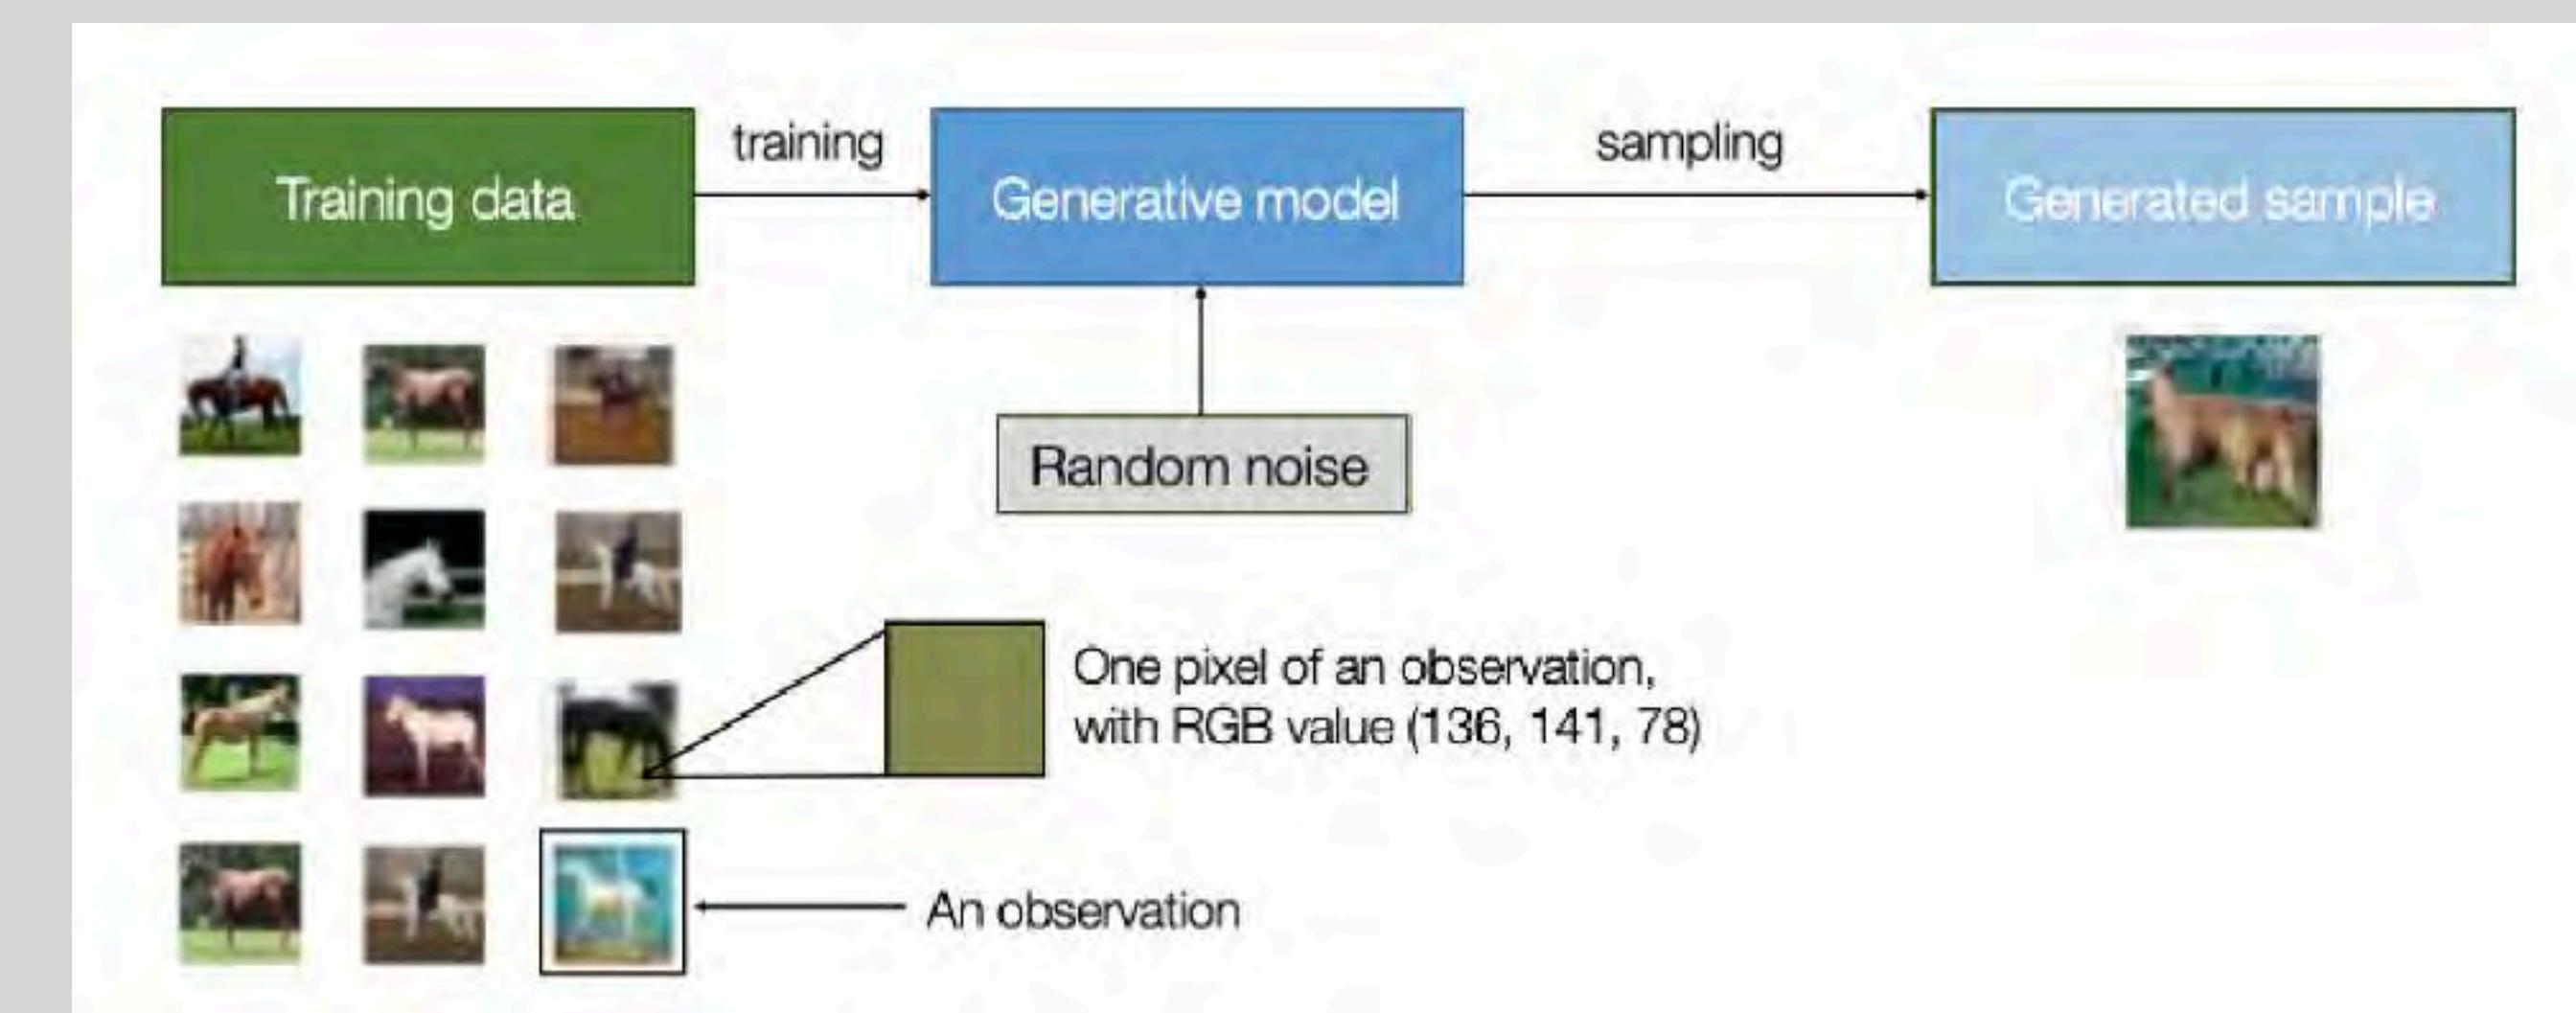
What is Generative Modeling?

A generative model can be broadly defined as follows: A generative model describes how a dataset is generated, in terms of a probabilistic model.

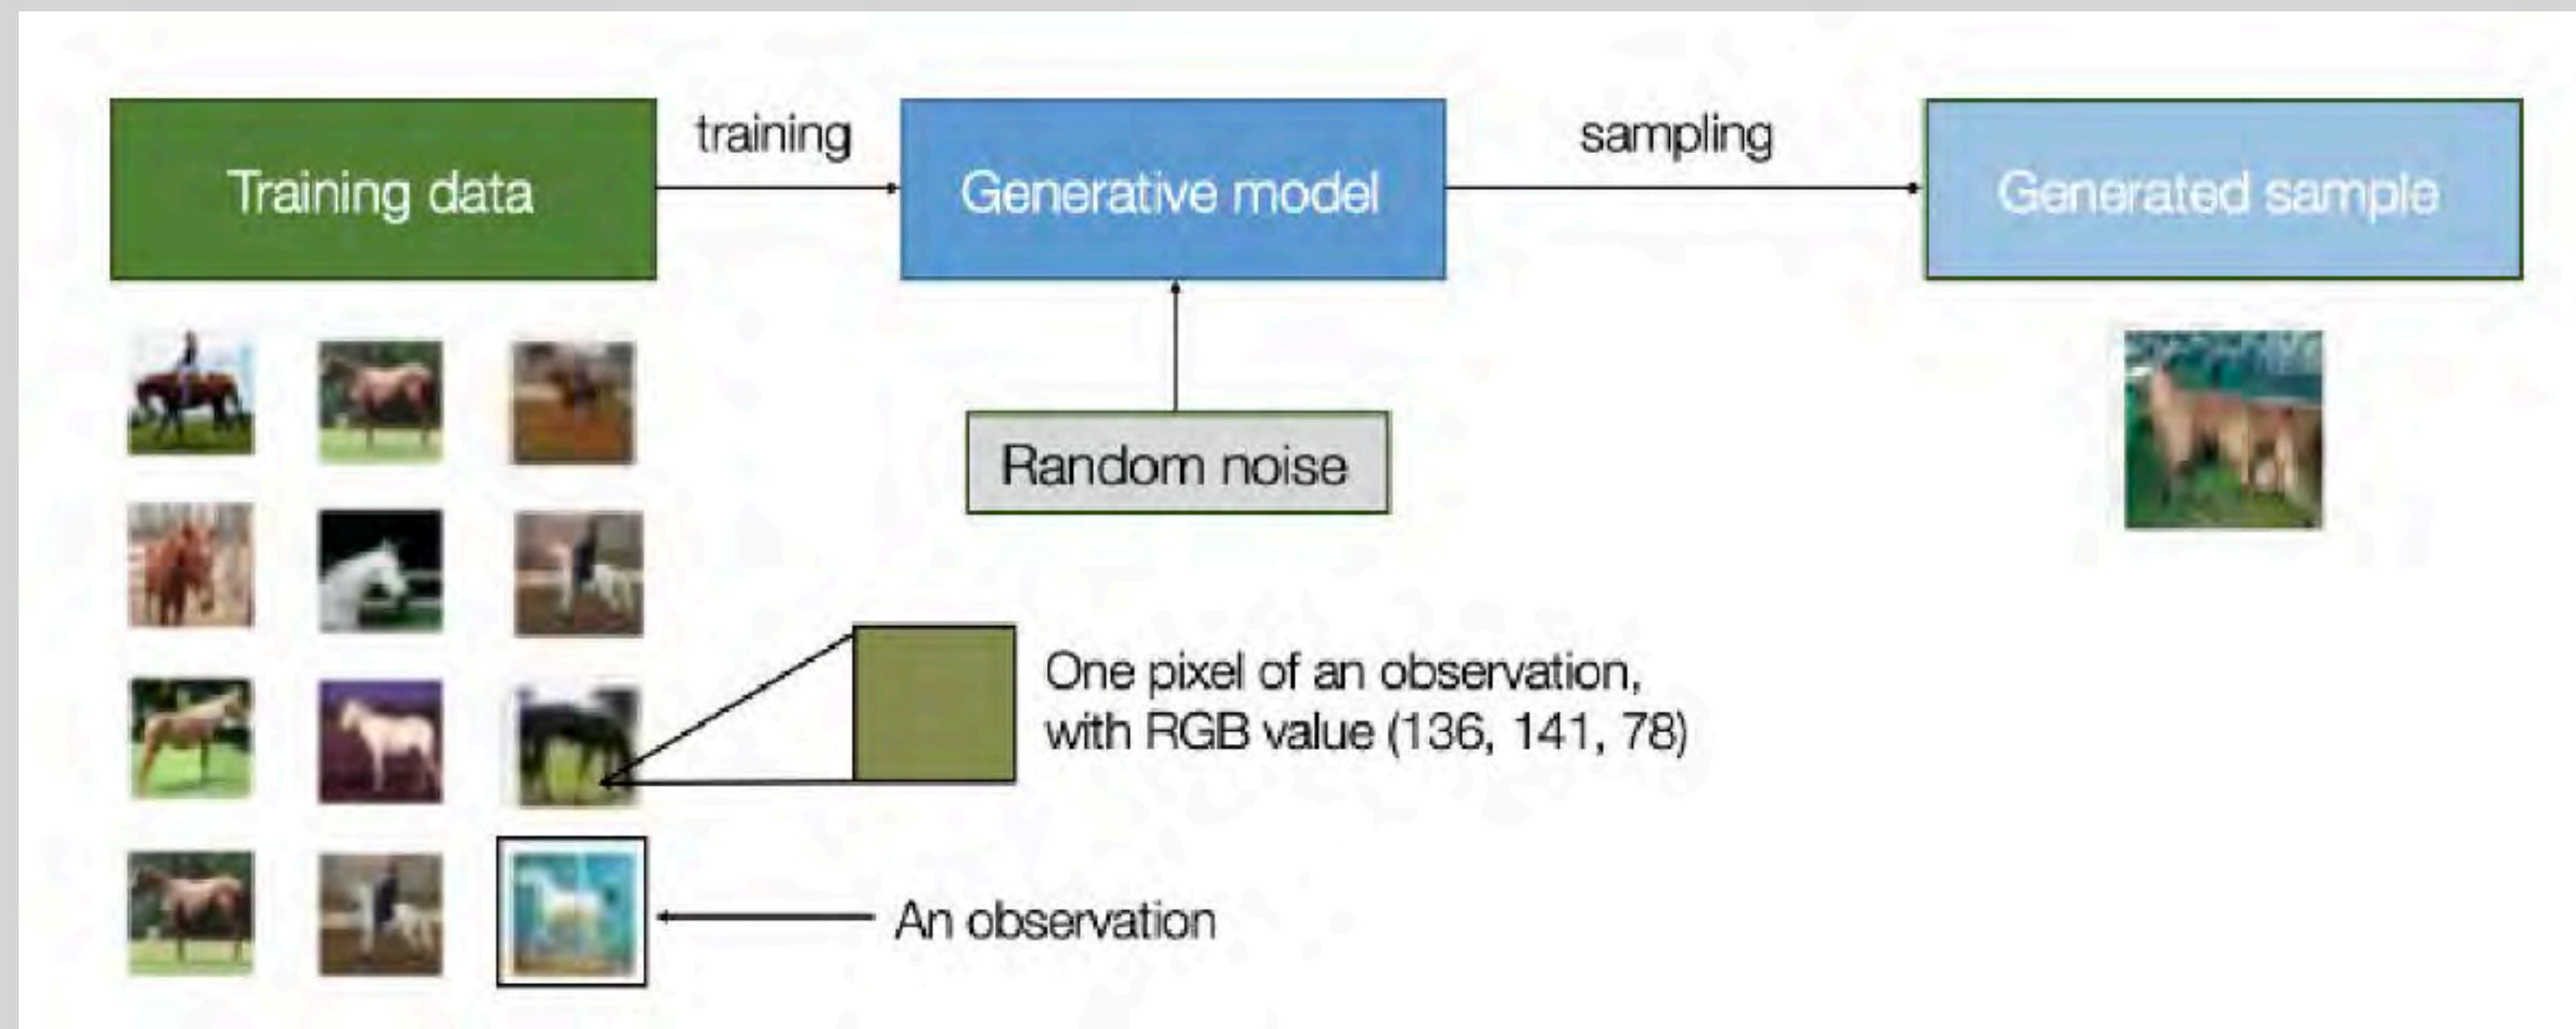


Suppose we have a dataset containing images of horses.

We may wish to build a model that can generate a new image of a horse that has never existed but still looks real because the model has learned the general rules that govern the appearance of a horse.

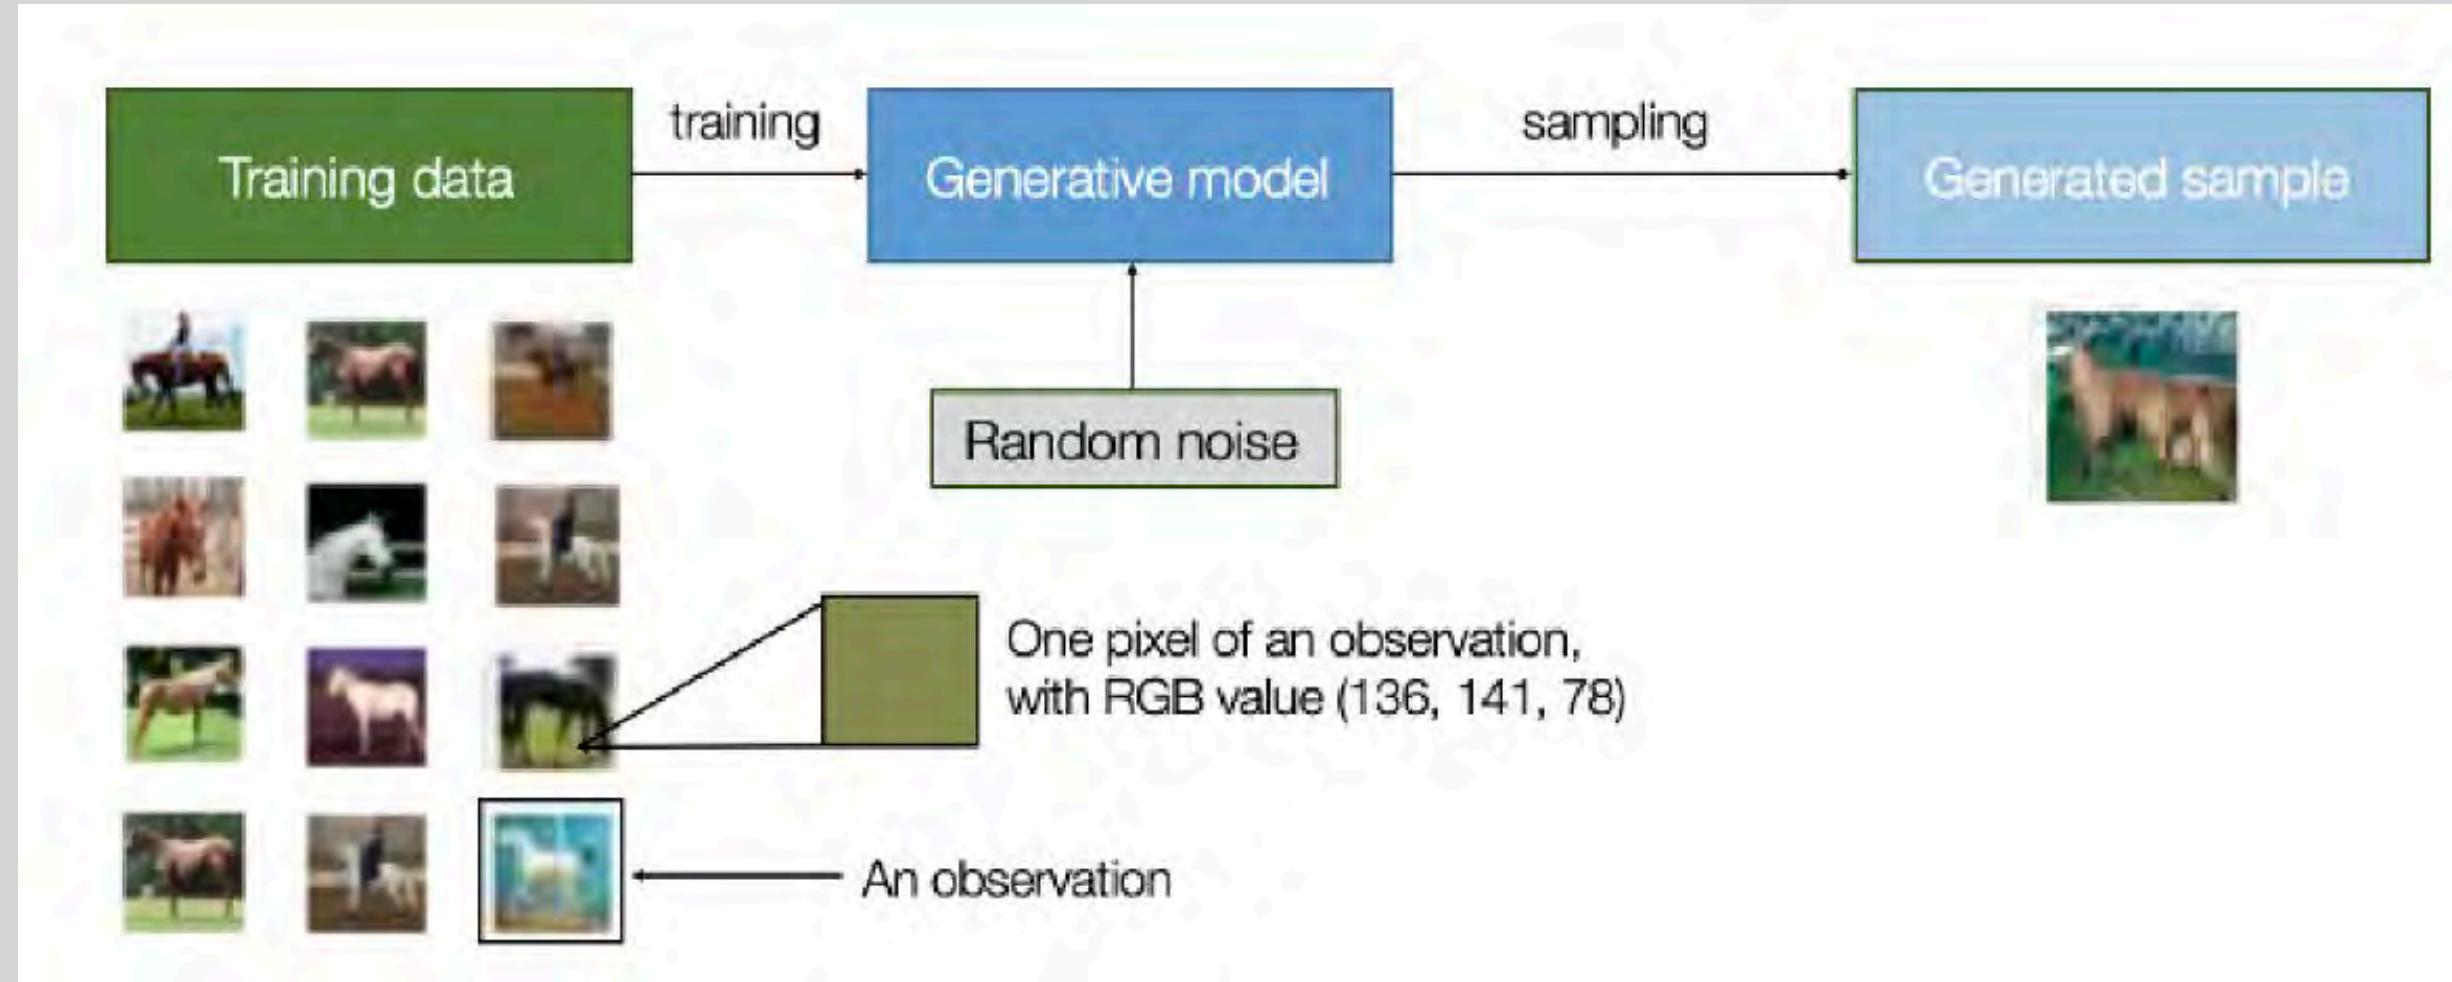


This is the kind of problem that can be solved using generative modeling.



First, we require a **dataset** consisting of many examples of the entity we are trying to generate.

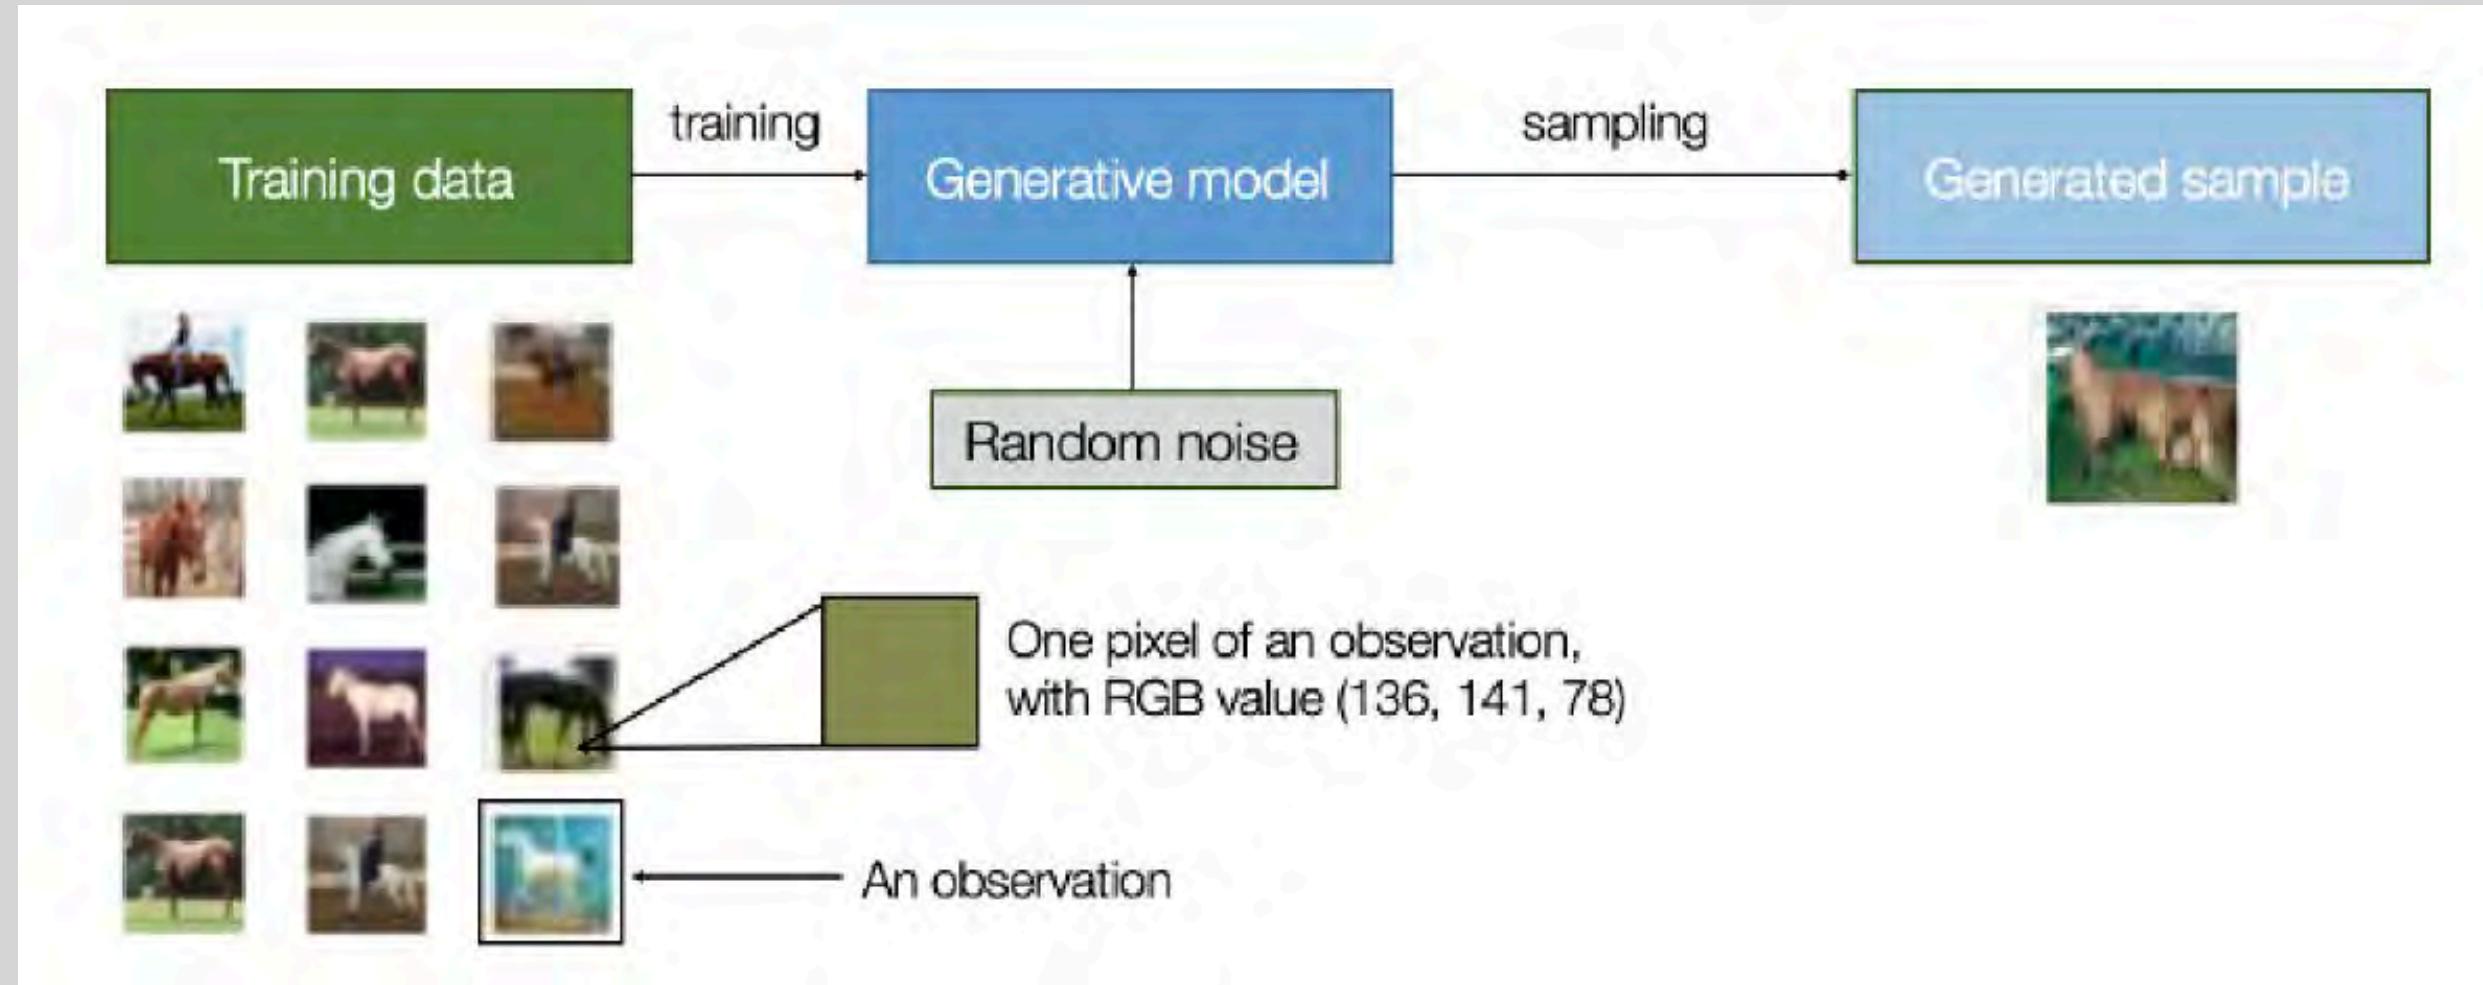
This is known as the **training data**, and one such data point is called an **observation**.



A generative model must also be probabilistic rather than deterministic.

If our model is merely a fixed calculation, such as taking the average value of each pixel in the dataset, **it is not generative** because the model produces the same output every time.

The model must include a stochastic (random) element that influences the individual samples generated by the model.



In other words, we can imagine that there is some unknown probabilistic distribution that explains why some images are likely to be found in the training dataset and other images are not.

It is our job to build a model that mimics this distribution as closely as possible and then sample from it to generate new images.

Generative vs Discriminative Modeling

Generative vs Discriminative Modeling

In order to comprehend generative algorithms, it can be helpful to contrast them with discriminative algorithms.

Generative algorithms, on the other hand, do the opposite; they aim to predict features given a certain label.

When input data is fed into a discriminative algorithm, it aims to predict the label to which the data belongs. The algorithm aims to map features to labels.

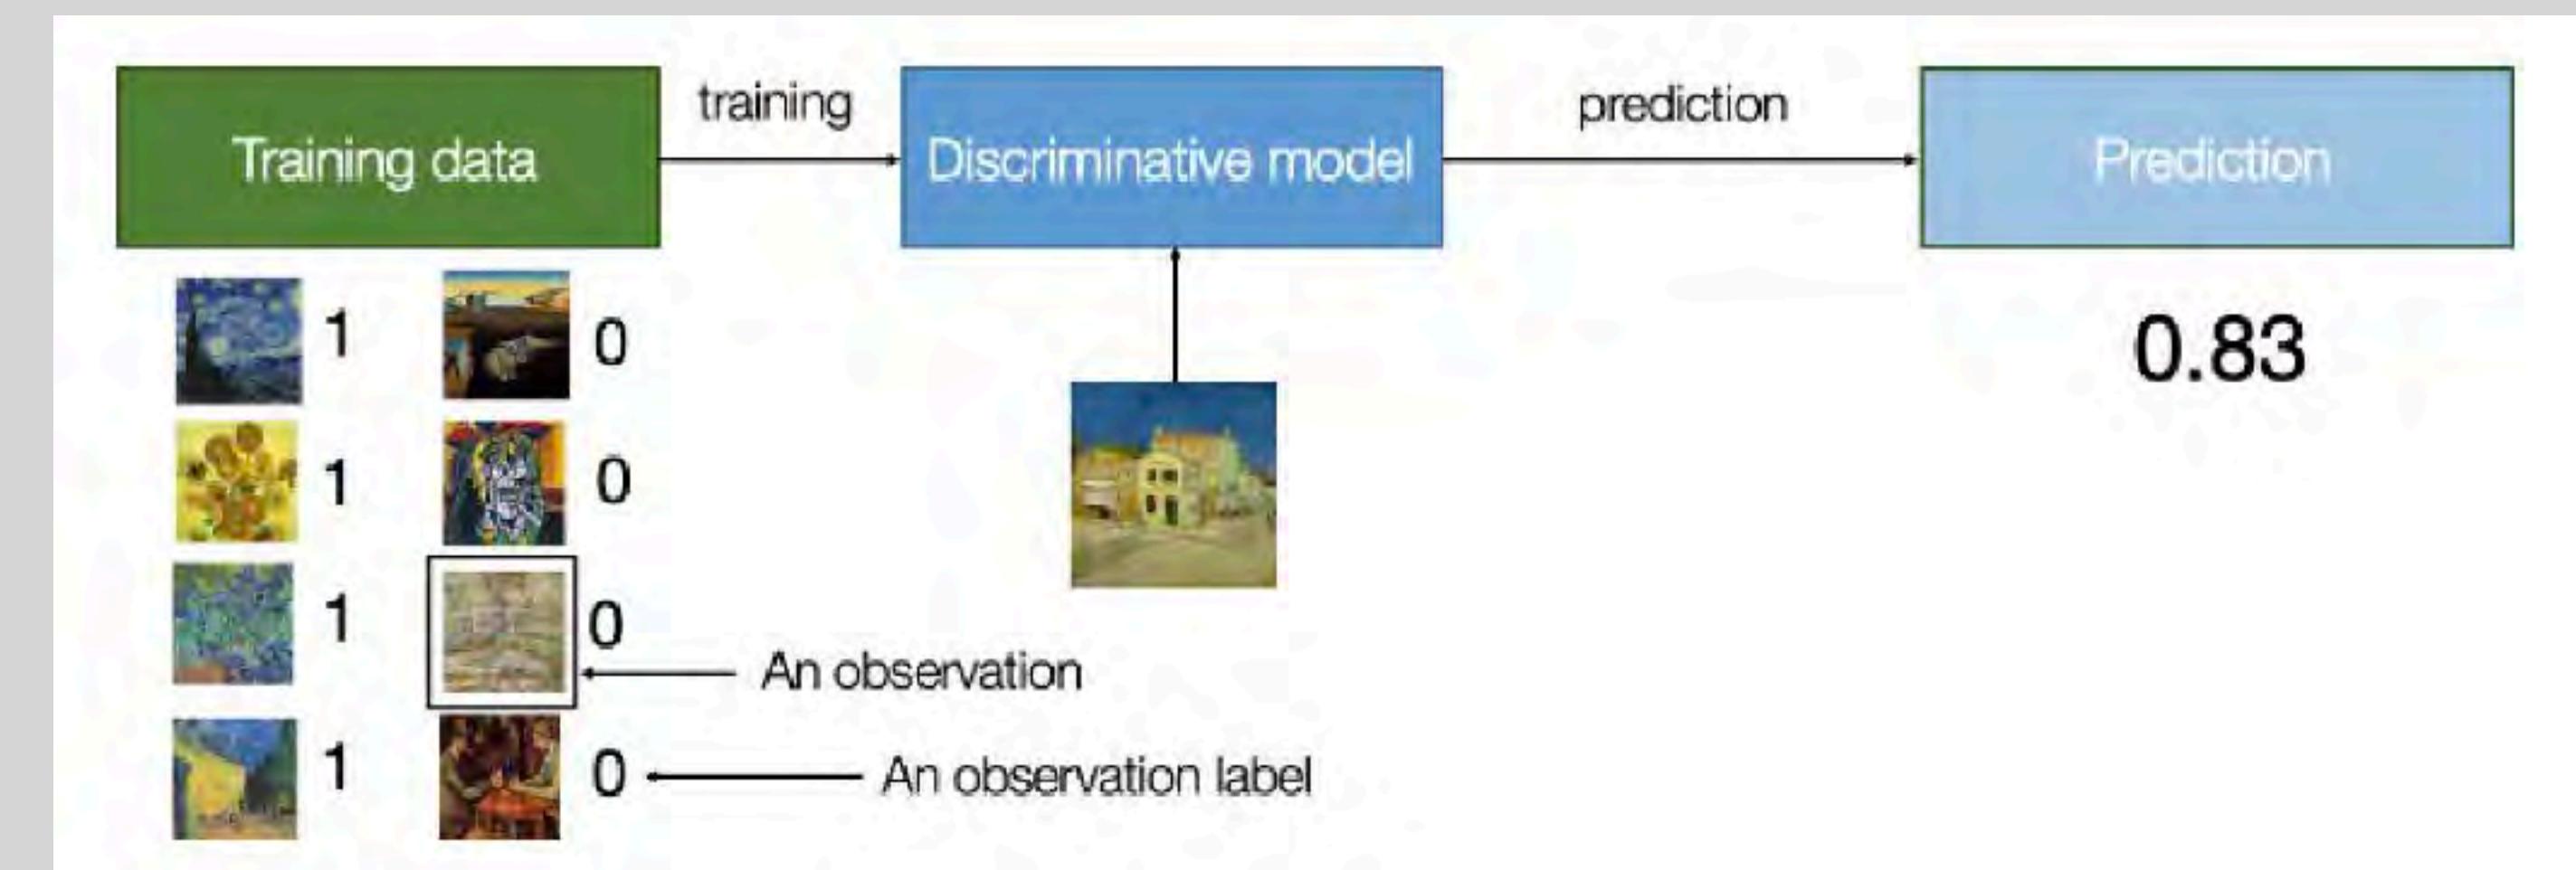
To truly understand what **generative modeling** aims to achieve and why this is important, it is useful to **compare** it to its counterpart, **discriminative modeling**.

Discriminative modeling

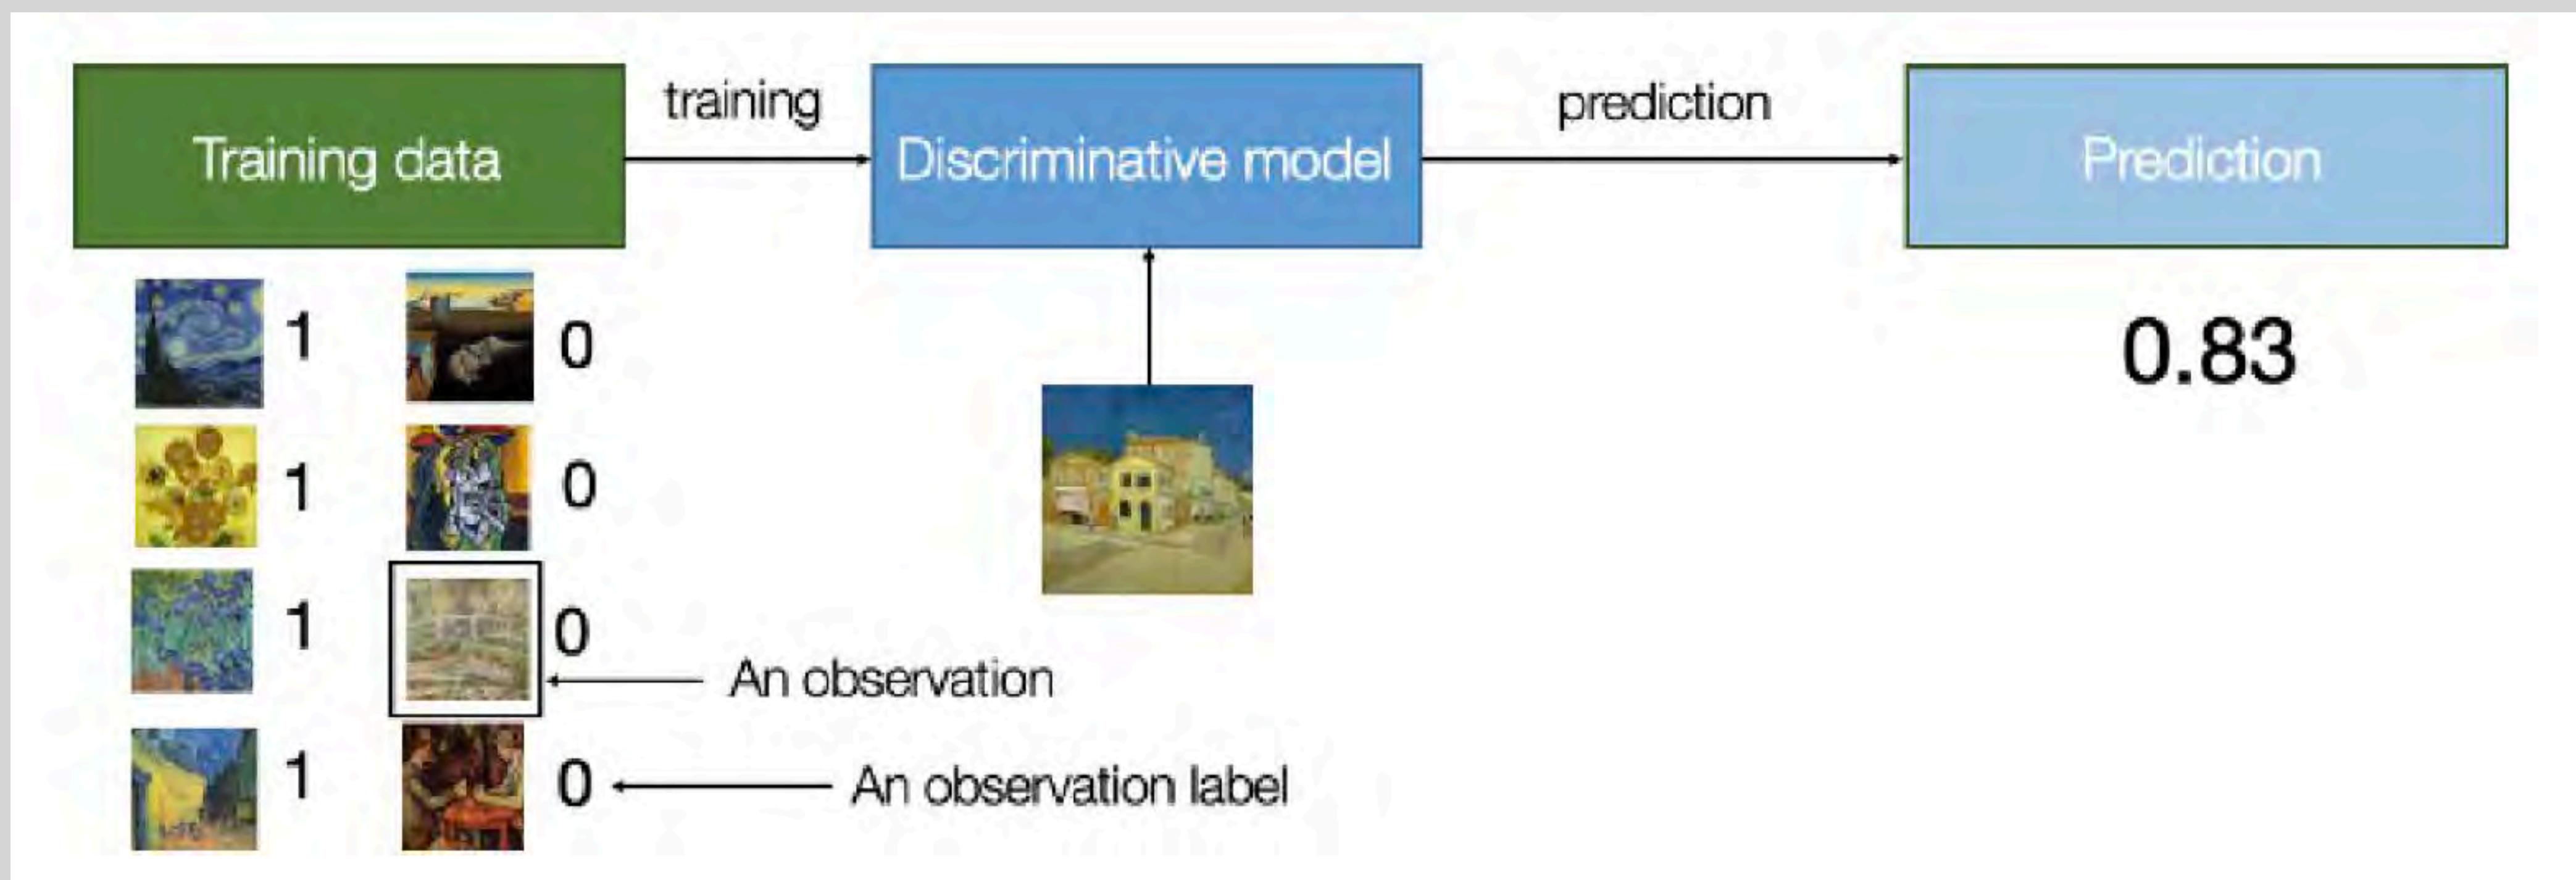
Example:

Suppose we have a dataset of paintings, some painted by Van Gogh and some by other artists.

With enough data, we could train a **discriminative** model to predict if a given painting was painted by Van Gogh.

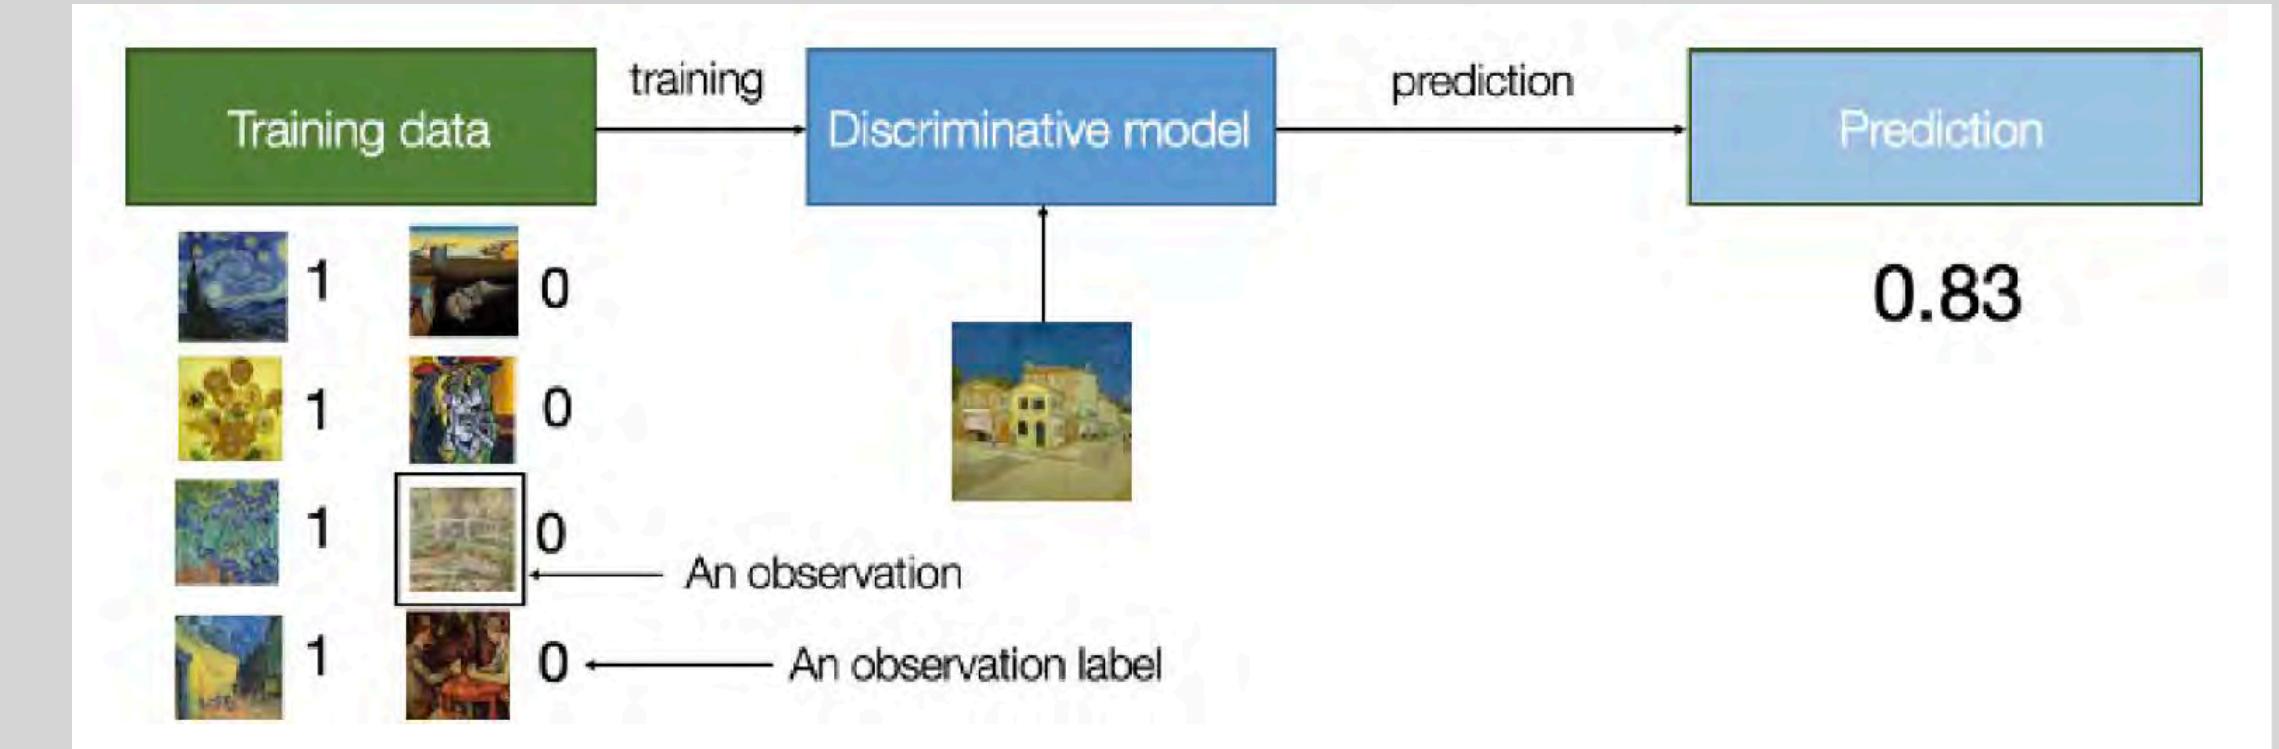


Our model would learn that certain colors, shapes, and textures are more likely to indicate that a painting is by Van Gogh, and for paintings with these features, the model would upweight its prediction accordingly.



Our model then learns how to discriminate between these two groups and outputs the probability that a new observation has label 1 that it was painted by Van Gogh.

For this reason, discriminative modeling is synonymous with **supervised learning**, or learning a function that maps an input to an output using a labeled dataset.



Generative modeling is usually performed with an **unlabeled dataset** (that is, as a form of unsupervised learning), though it can also be applied to a labeled dataset to learn how to generate observations from each distinct class.

Let's take a look at some mathematical notation to describe the difference between generative and discriminative modeling.

Discriminative modeling estimates $p(y|x)$ —the probability of a label y given observation x .

Generative modeling estimates $p(x)$ —the probability of observing observation x .

If the dataset is labeled, we can also build a generative model that estimates the distribution $p(x|y)$.

In other words, **Discriminative Modeling** attempts to estimate the probability that an observation x belongs to category y .

Generative modeling doesn't care about labeling observations. Instead, it attempts to estimate the probability of seeing the observation at all.

Discriminative Modeling

The key point is that even if we were able to build a perfect discriminative model to identify Van Gogh paintings, it would still have no idea how to create a painting that looks like a Van Gogh.

Generative Modeling

We would instead need to train a generative model, which can output sets of pixels that have a high chance of belonging to the original training dataset.

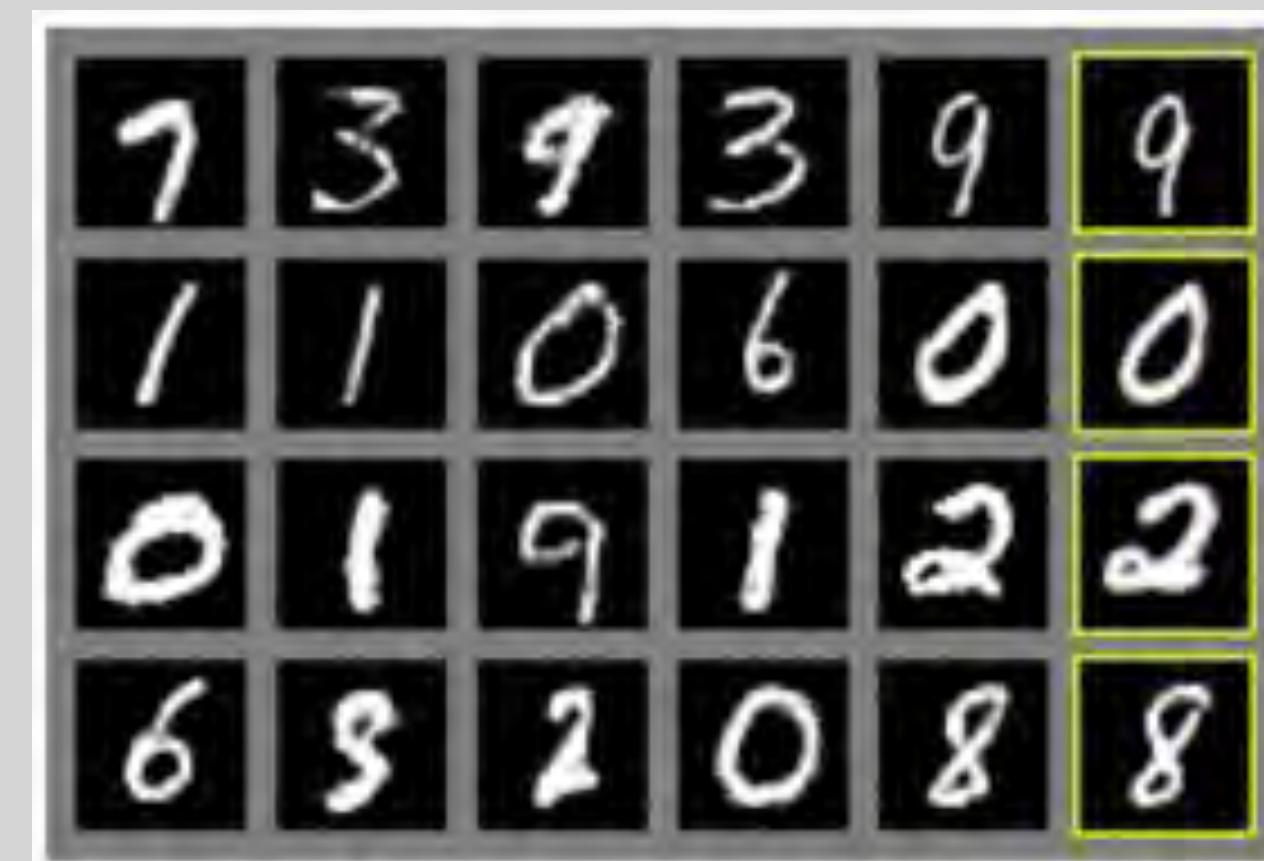
Generative Adversarial Networks (GAN)

GANs are comprised of two neural networks: a **generator** and a **discriminator**.

They are able to generate new, synthetic data.

The generator outputs new instances of the data, while the discriminator determines whether each instance of the data belongs to the training dataset.

The following image gives an illustration of the output from a GAN datasets. The images on the far-right side of the grid are the true values and the others are generated by the model:

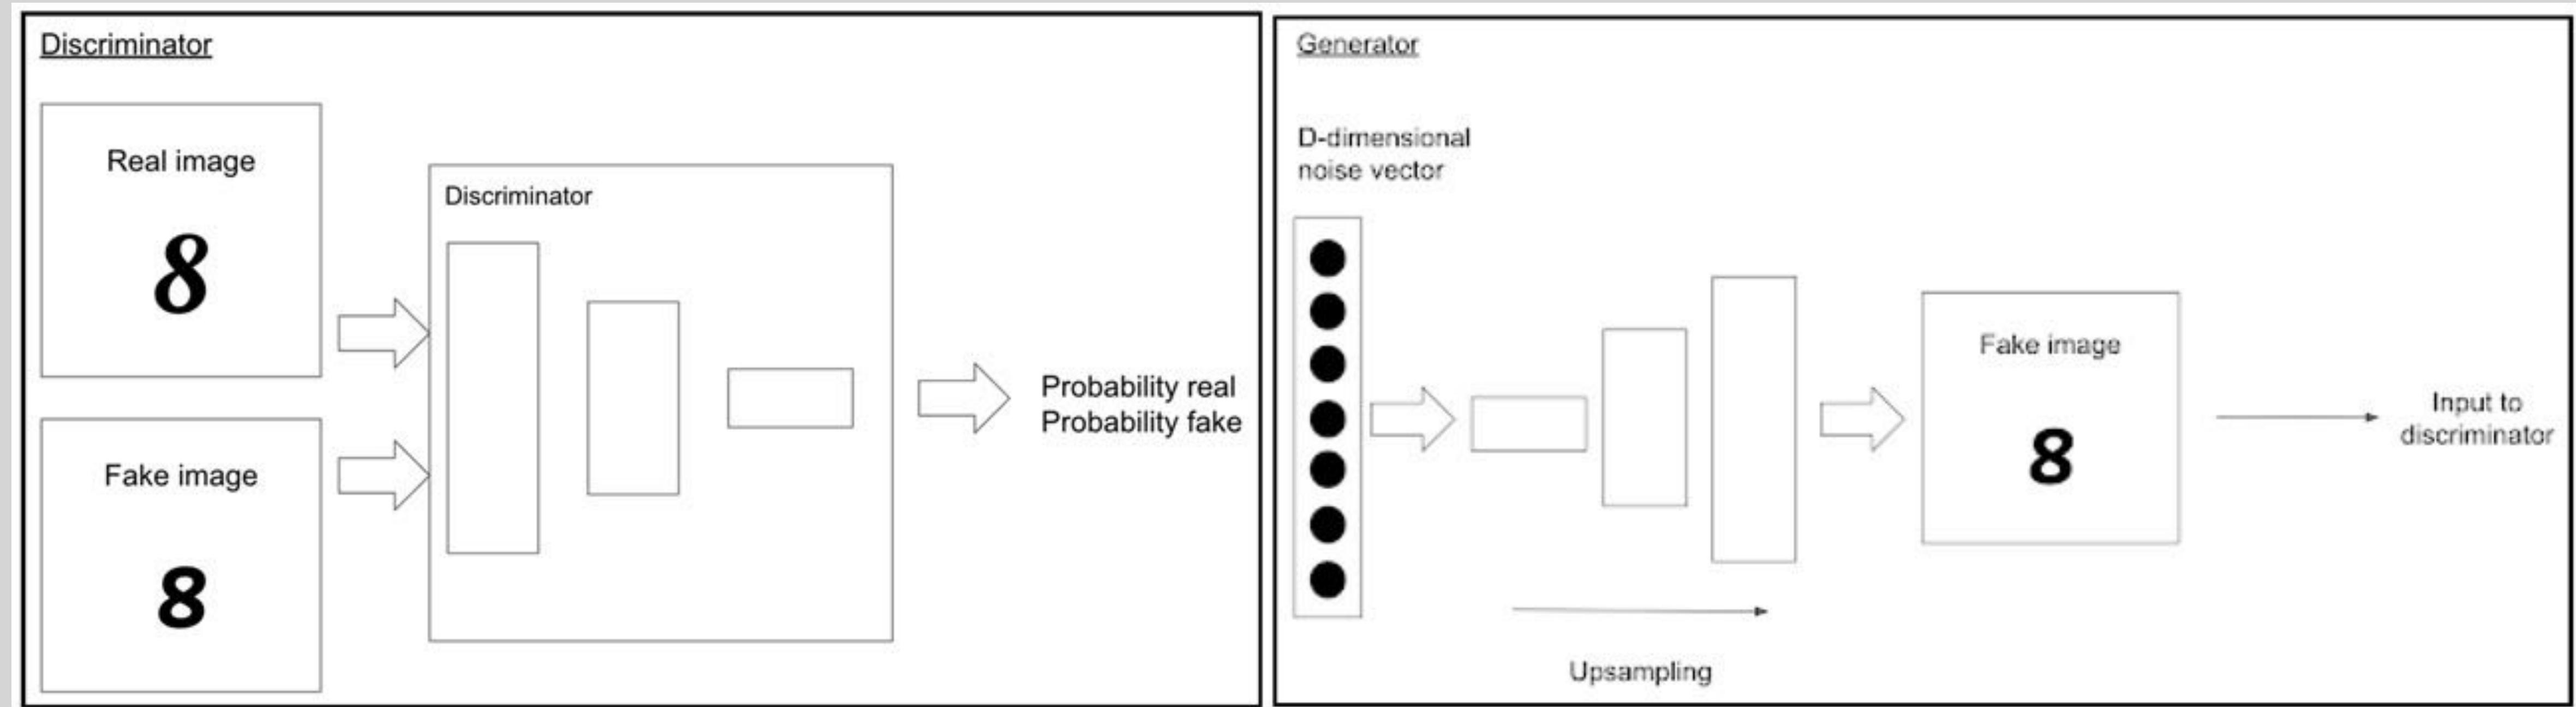


The following are the steps that a GAN takes:

1. Random numbers are fed into the generator and an image is generated

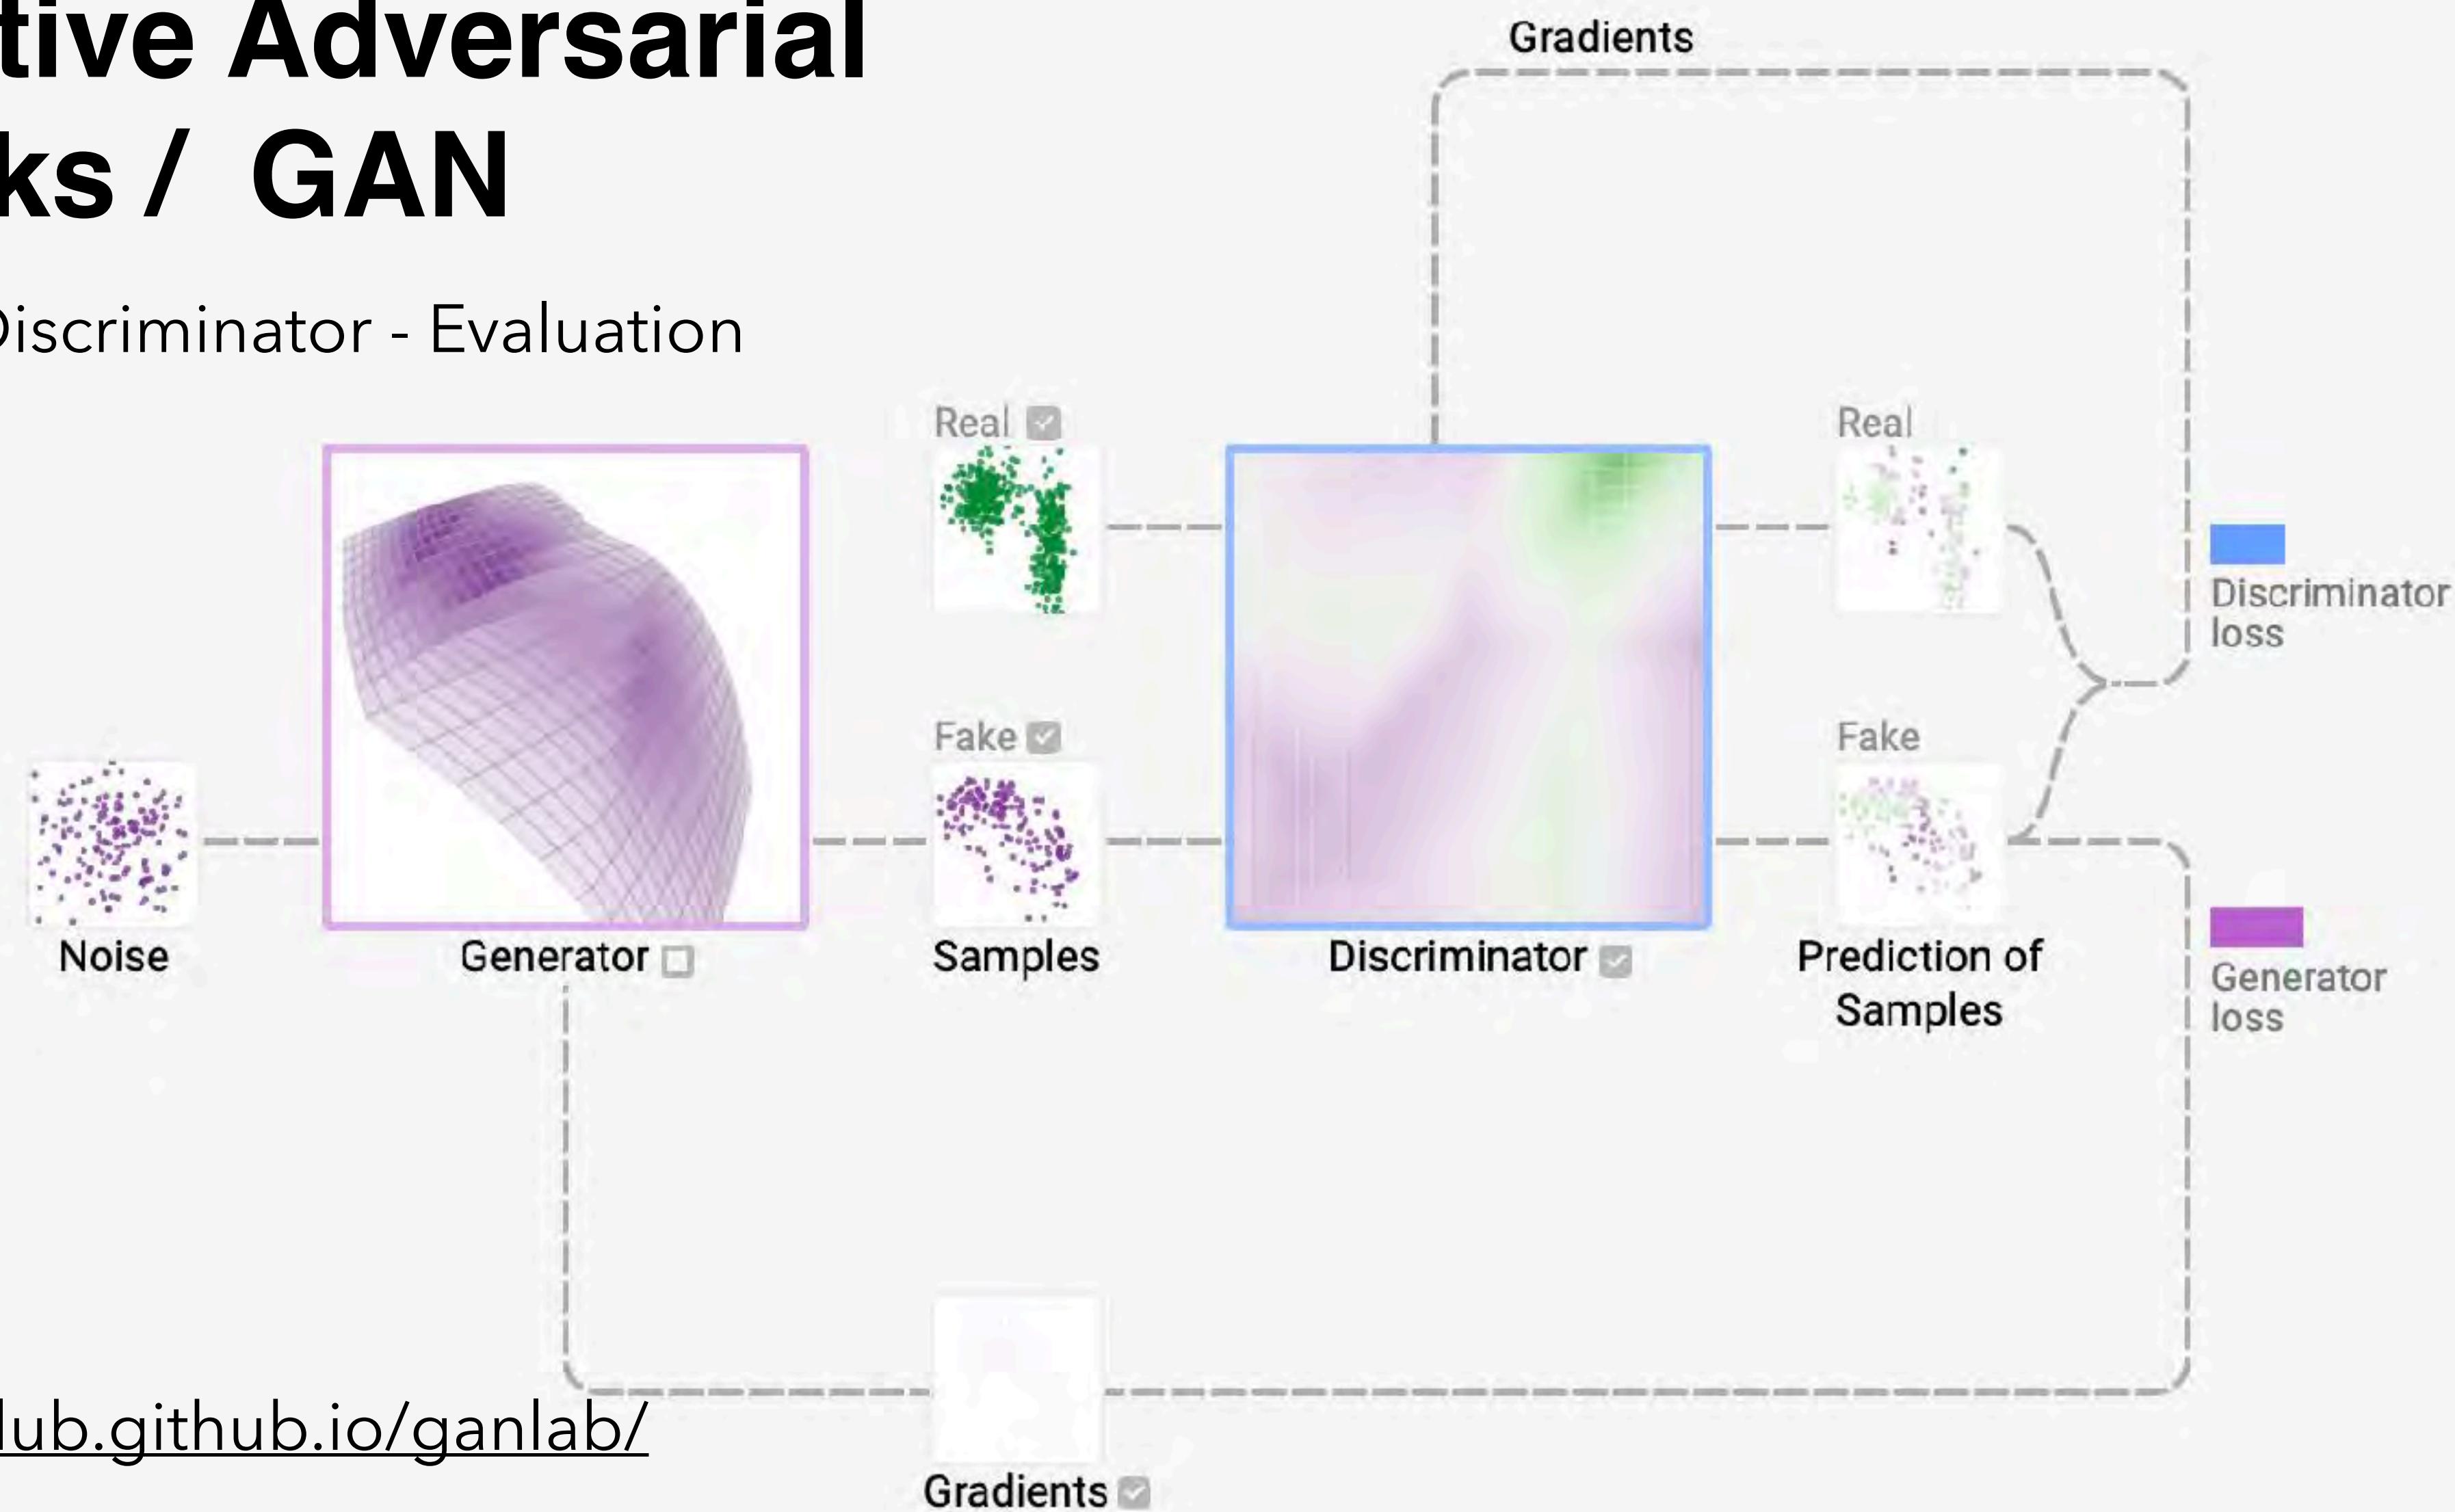
2. The generated image is fed into the discriminator along with other images taken from the real dataset

3. The discriminator considers all of the images fed into it and returns a probability as to whether it thinks the image is real or fake



Generative Adversarial Networks / GAN

Generator - Discriminator - Evaluation



Training Generative Adversarial Networks (GANs)

The training process of a GAN is as follows:

Training the discriminator: Generator values should be held constant

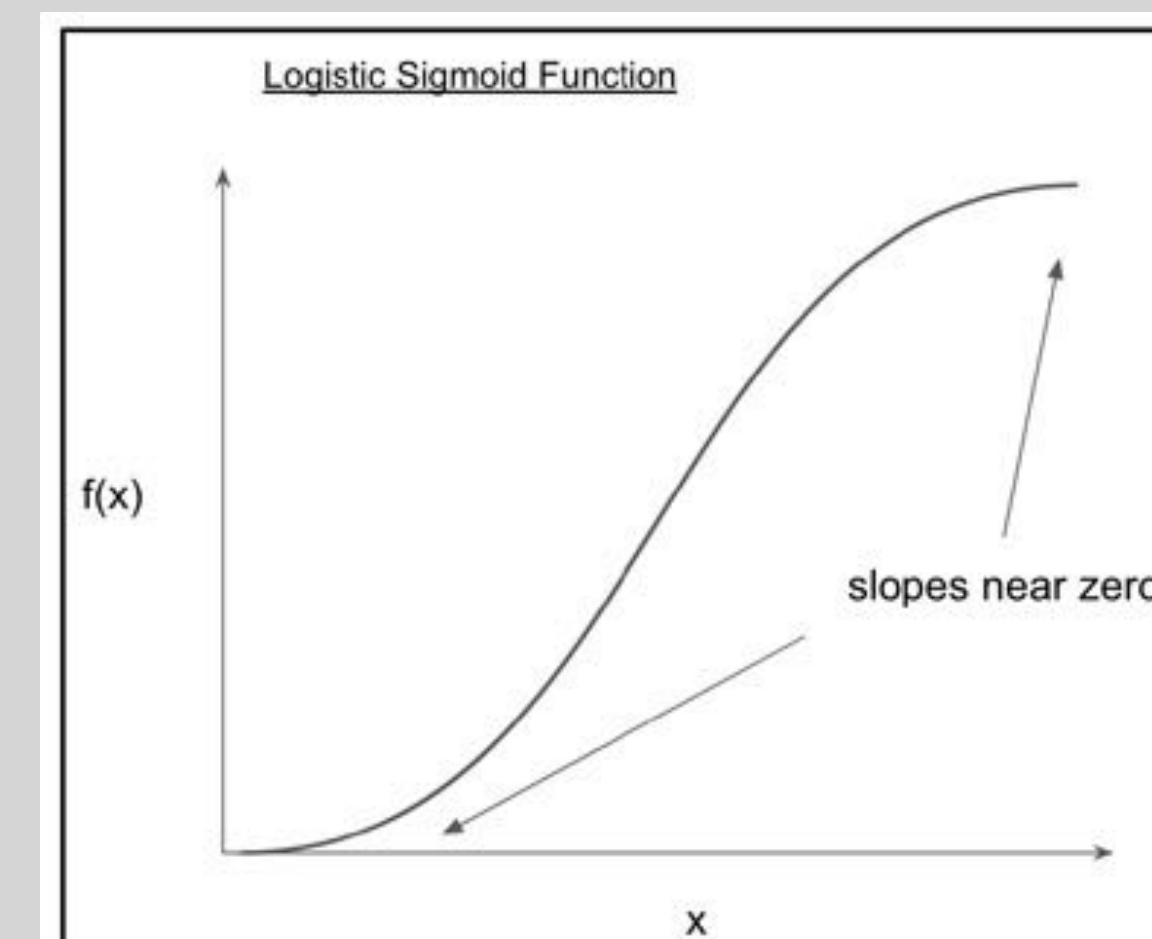
Training the generator: The discriminator should be pretrained against the original dataset

Either side of the GAN can overpower the other if one is significantly higher-performing than the other.

For example, if the discriminator is too high-performing, it will return a value that is very close to either 0 or 1, and the generator will have difficulties reading the gradient.

On the other hand, if the performance of the generator is too high, then it will continually exploit weaknesses in the discriminator that lead to false negatives.

Such a dilemma during training is illustrated in the following graph for the case of a binary classification problem. We want to stop one side of the GAN from winning so that both sides can continue to learn together for an extended period of time:



GAN challenges

GANs often face the following major challenges:

Mode collapse: When the generator collapses and produces limited varieties of samples

Diminished gradient: When the discriminator performs too well that the generator gradient vanishes and learns nothing

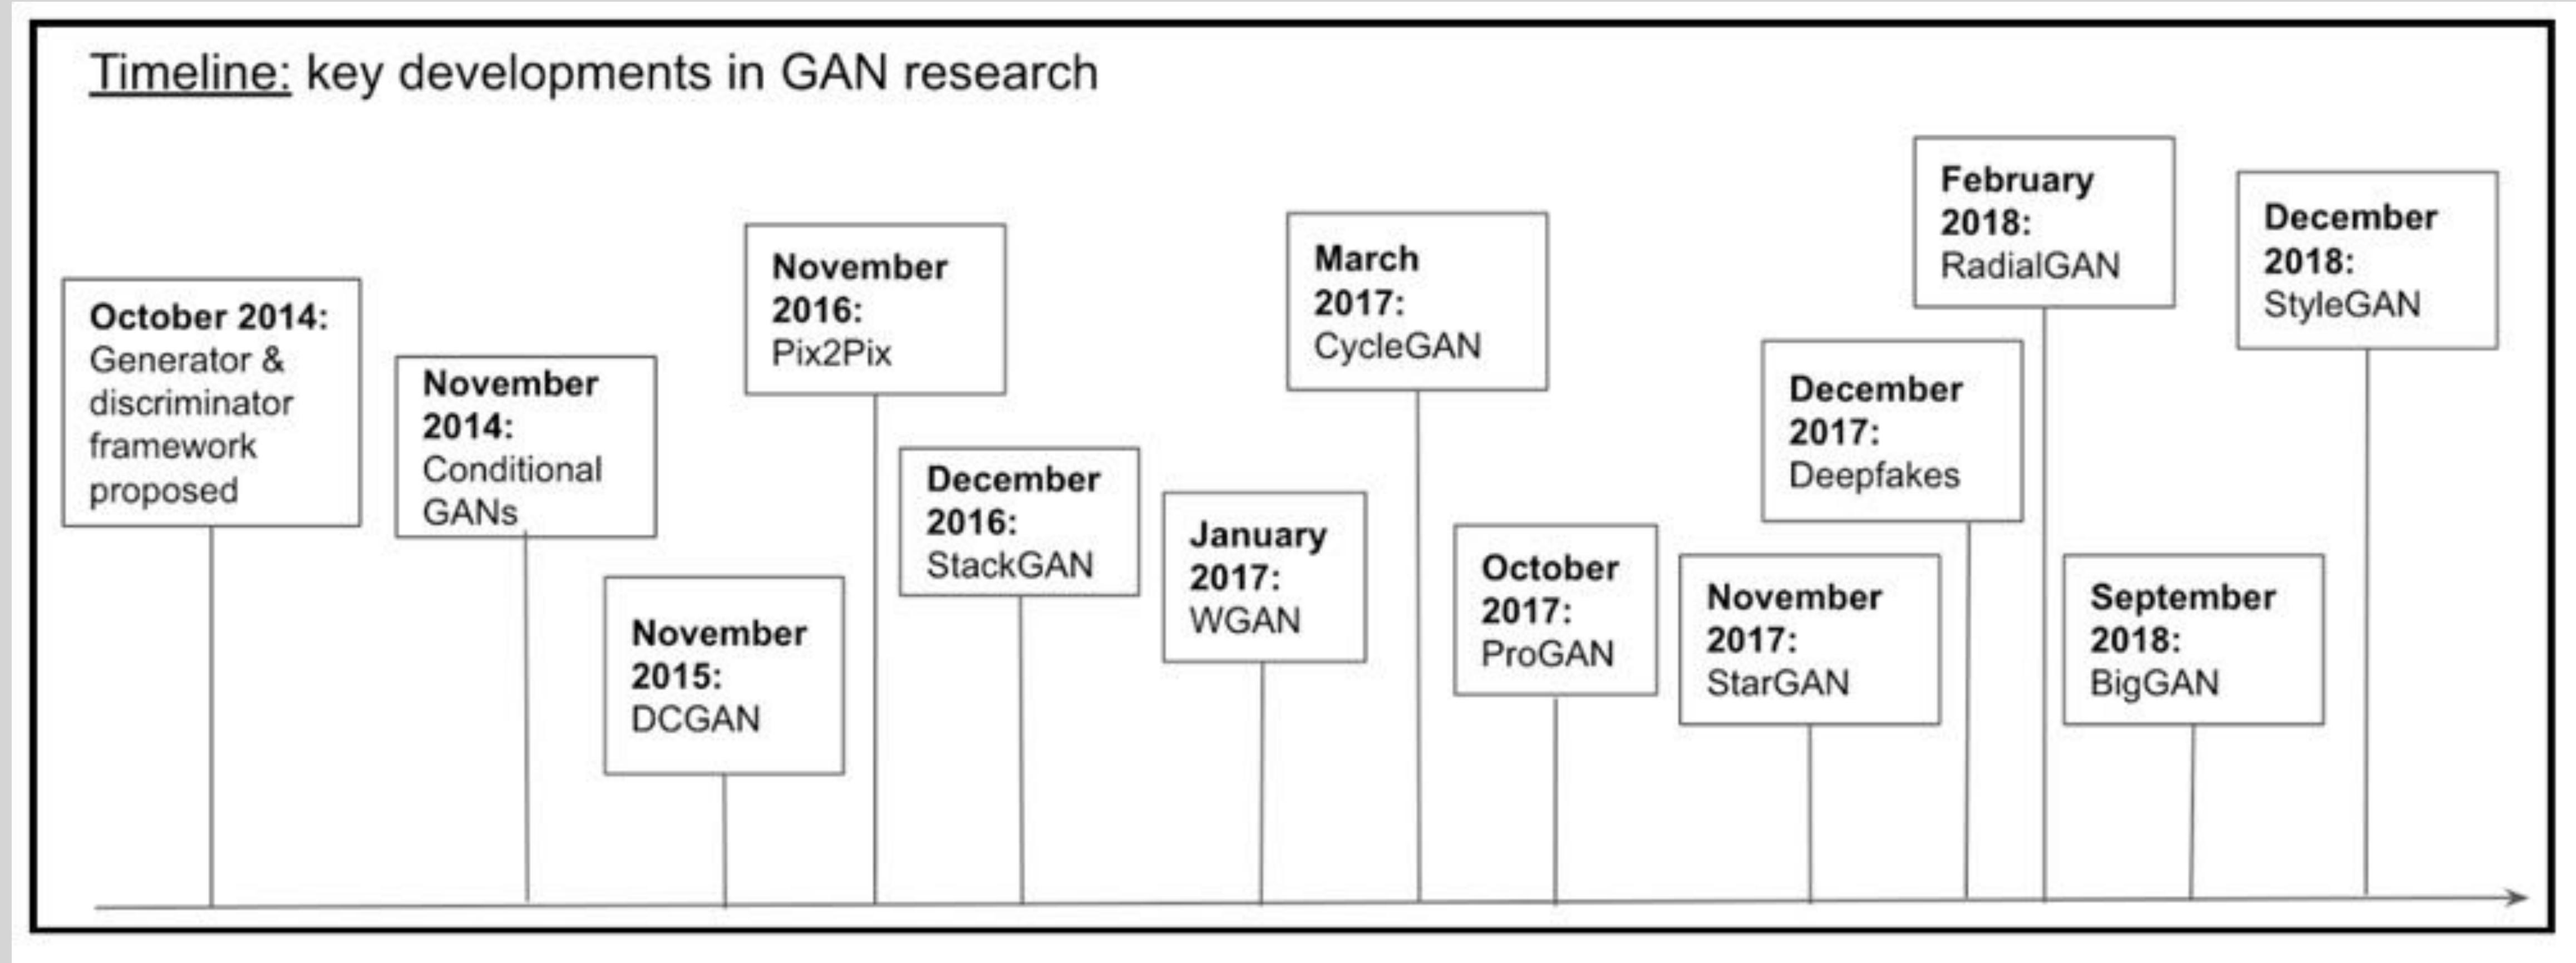
High sensitivity: They are highly sensitive to the hyperparameter selections

Overfitting: An imbalance between the generator and discriminator can cause overfitting

Long training time: If the model is trained on a GPU, it can take hours

Non-convergence: The models may never converge, as the parameters can oscillate and destabilize

GAN variations and timelines

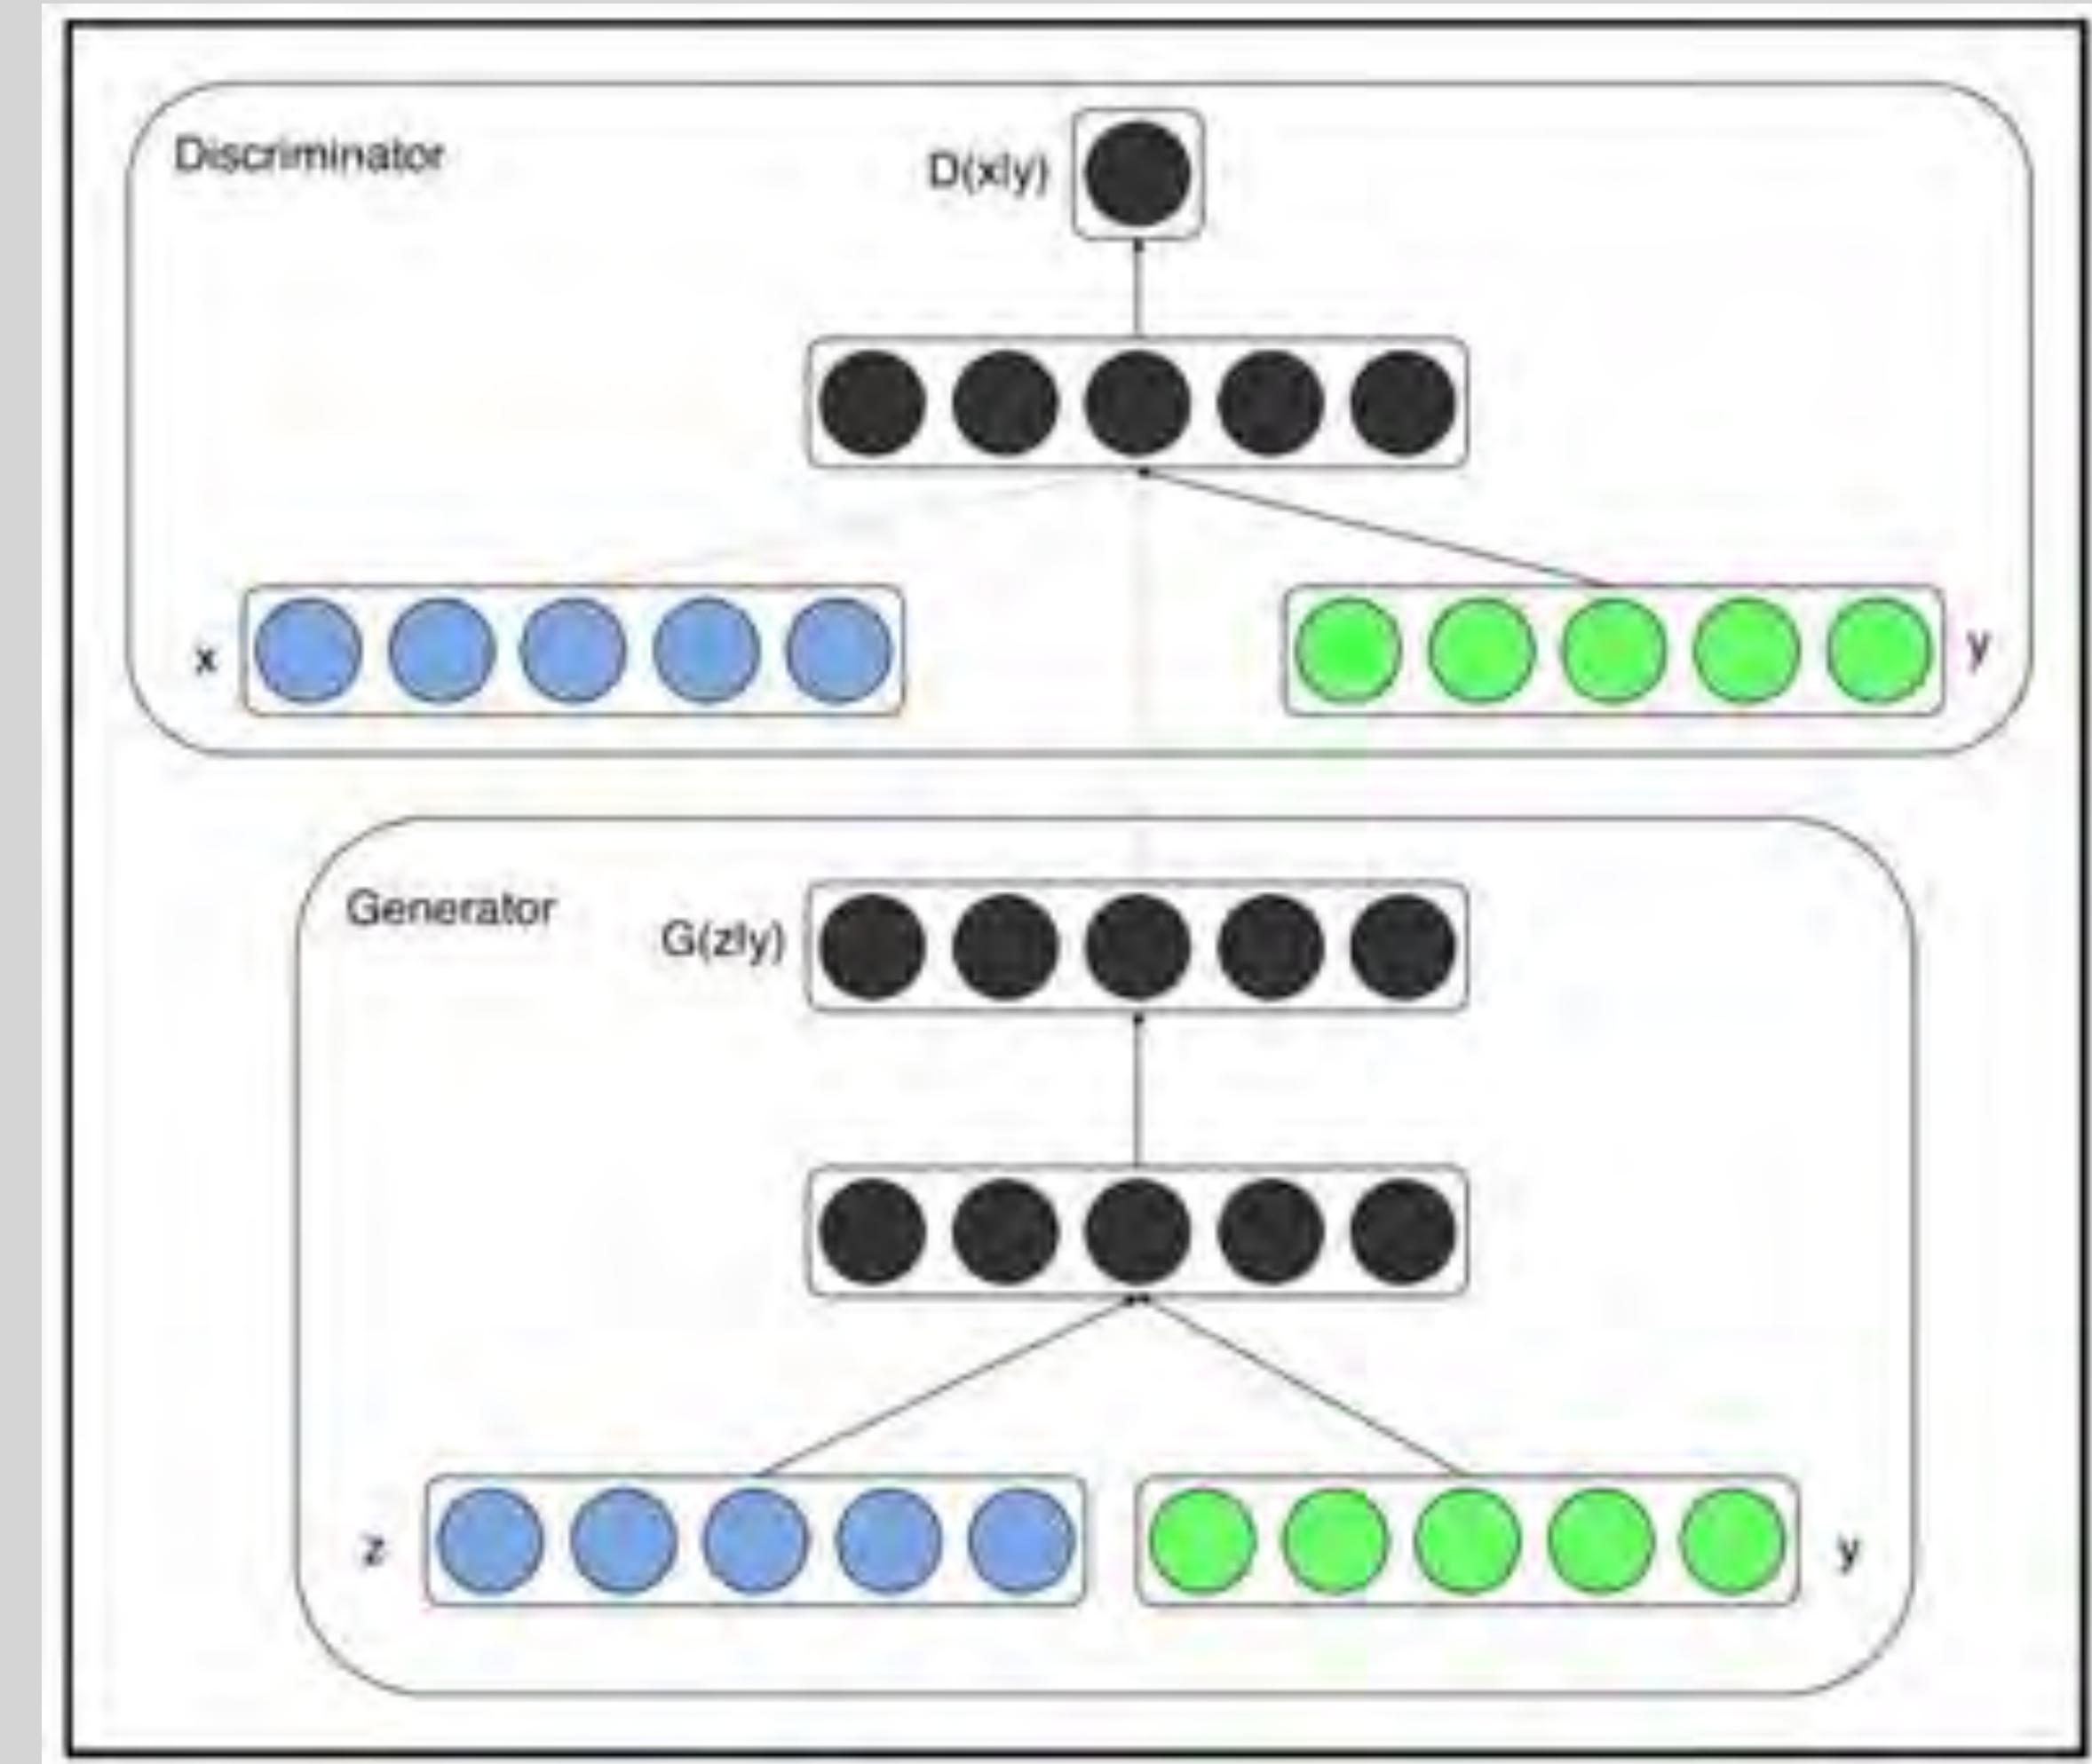


Conditional GANs

Conditional GANs are a central theme that form the building blocks of many state-of-theart GANs.

The paper submitted by Mirza and Osindero in 2014 shows how integrating the class labels of data yields greater stability in GAN train.

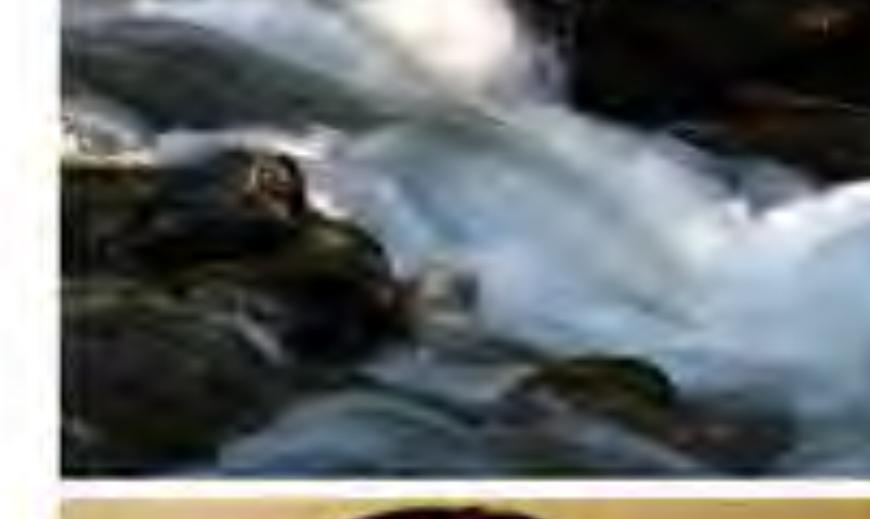
This idea of conditioning GANs with prior information is a common approach in future GAN research. It is particularly important for papers whose main focus is on image-to-image or text-to-image applications:



Conditional GANs Sample

In this sample we can create automated tagging of images, with multi-label predictions, using conditional adversarial nets to generate a (possibly multi-modal) distribution of tag-vectors conditional on image features.

For the experiments they use MIR Flickr 25,000 dataset , and extract the image and tags features using the convolutional model and language model we described above. The first 150,000 examples were used as training set.

User tags + annotations	Generated tags
   	taxi, passenger, line, transportation, railway station, passengers, railways, signals, rail, rails
montanha, trem, inverno, frio, people, male, plant life, tree, structures, transport, car	chicken, fattening, cooked, peanut, cream, cookie, house made, bread, biscuit, bakes
food, raspberry, delicious, homemade	creek, lake, along, near, river, rocky, treeline, valley, woods, waters
water, river	love, people, posing, girl, young, strangers, pretty, women, happy, life
people, portrait, female, baby, indoor	

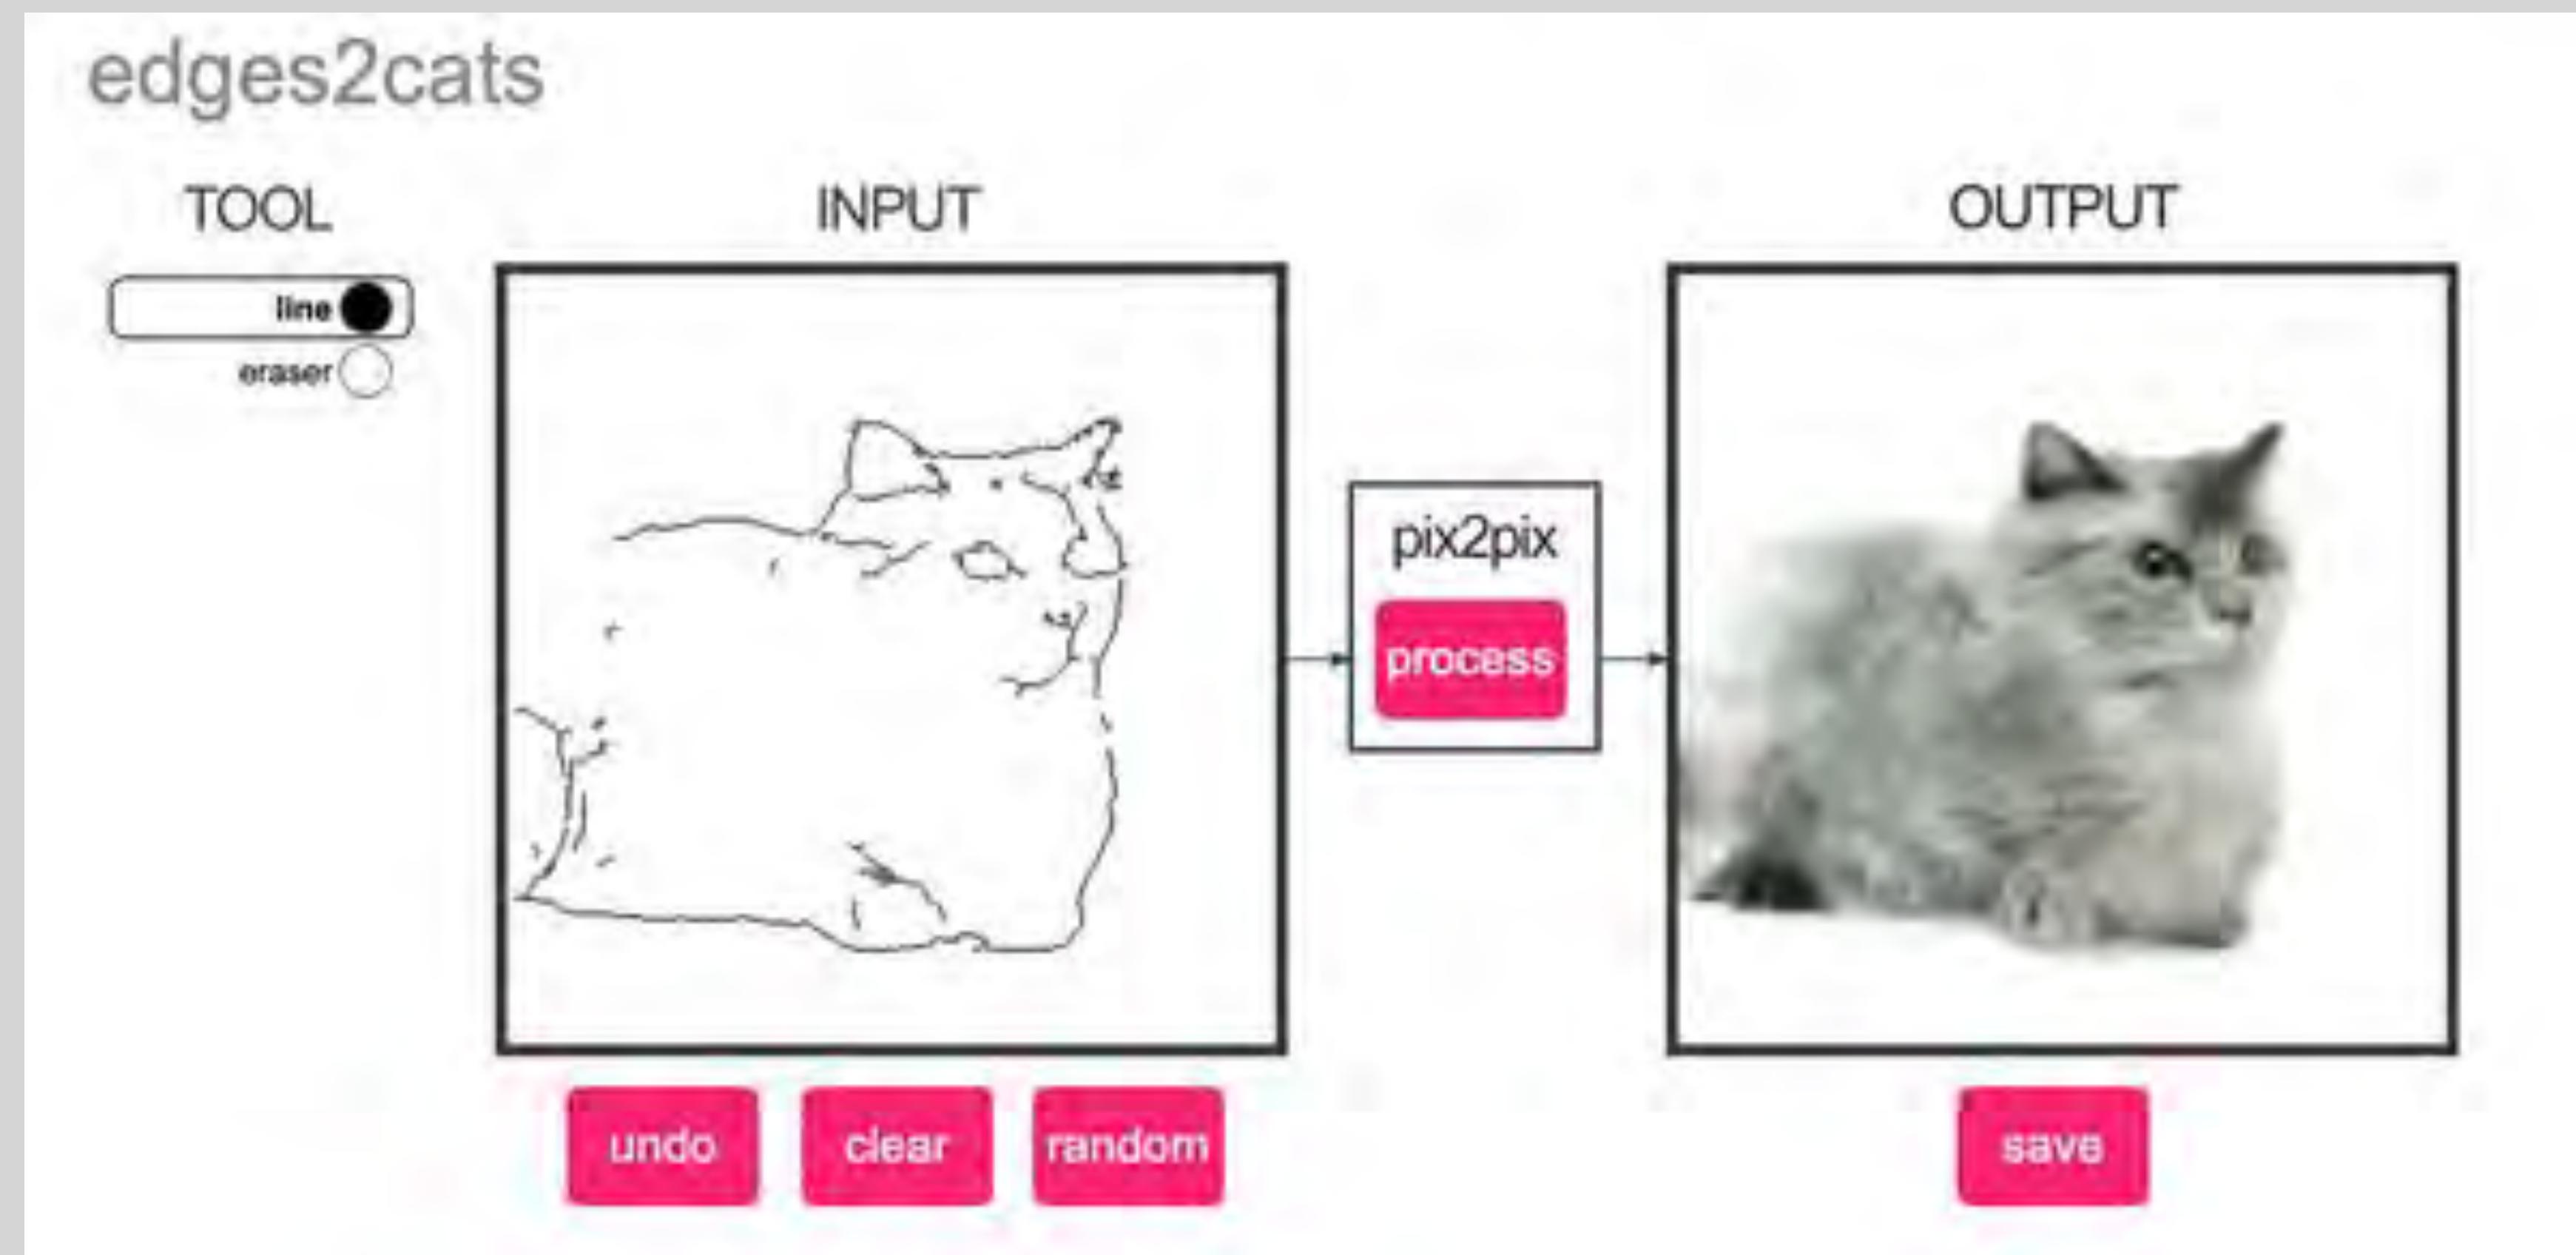
Pix2Pix translation GAN model

Another image-to-image translation GAN model is Pix2Pix. It has many applications such as edge maps to photo-realistic images, and black and white to colour. The architecture uses paired training samples and incorporates a PatchGAN. The PatchGAN looks at individual regions of the image in order to determine whether they are real or fake, versus when considering the whole image.

The following are some examples of images generated by implementing the Pix2Pix GAN:

Mapping the edges of a handbag to a full realistic image is as follows:



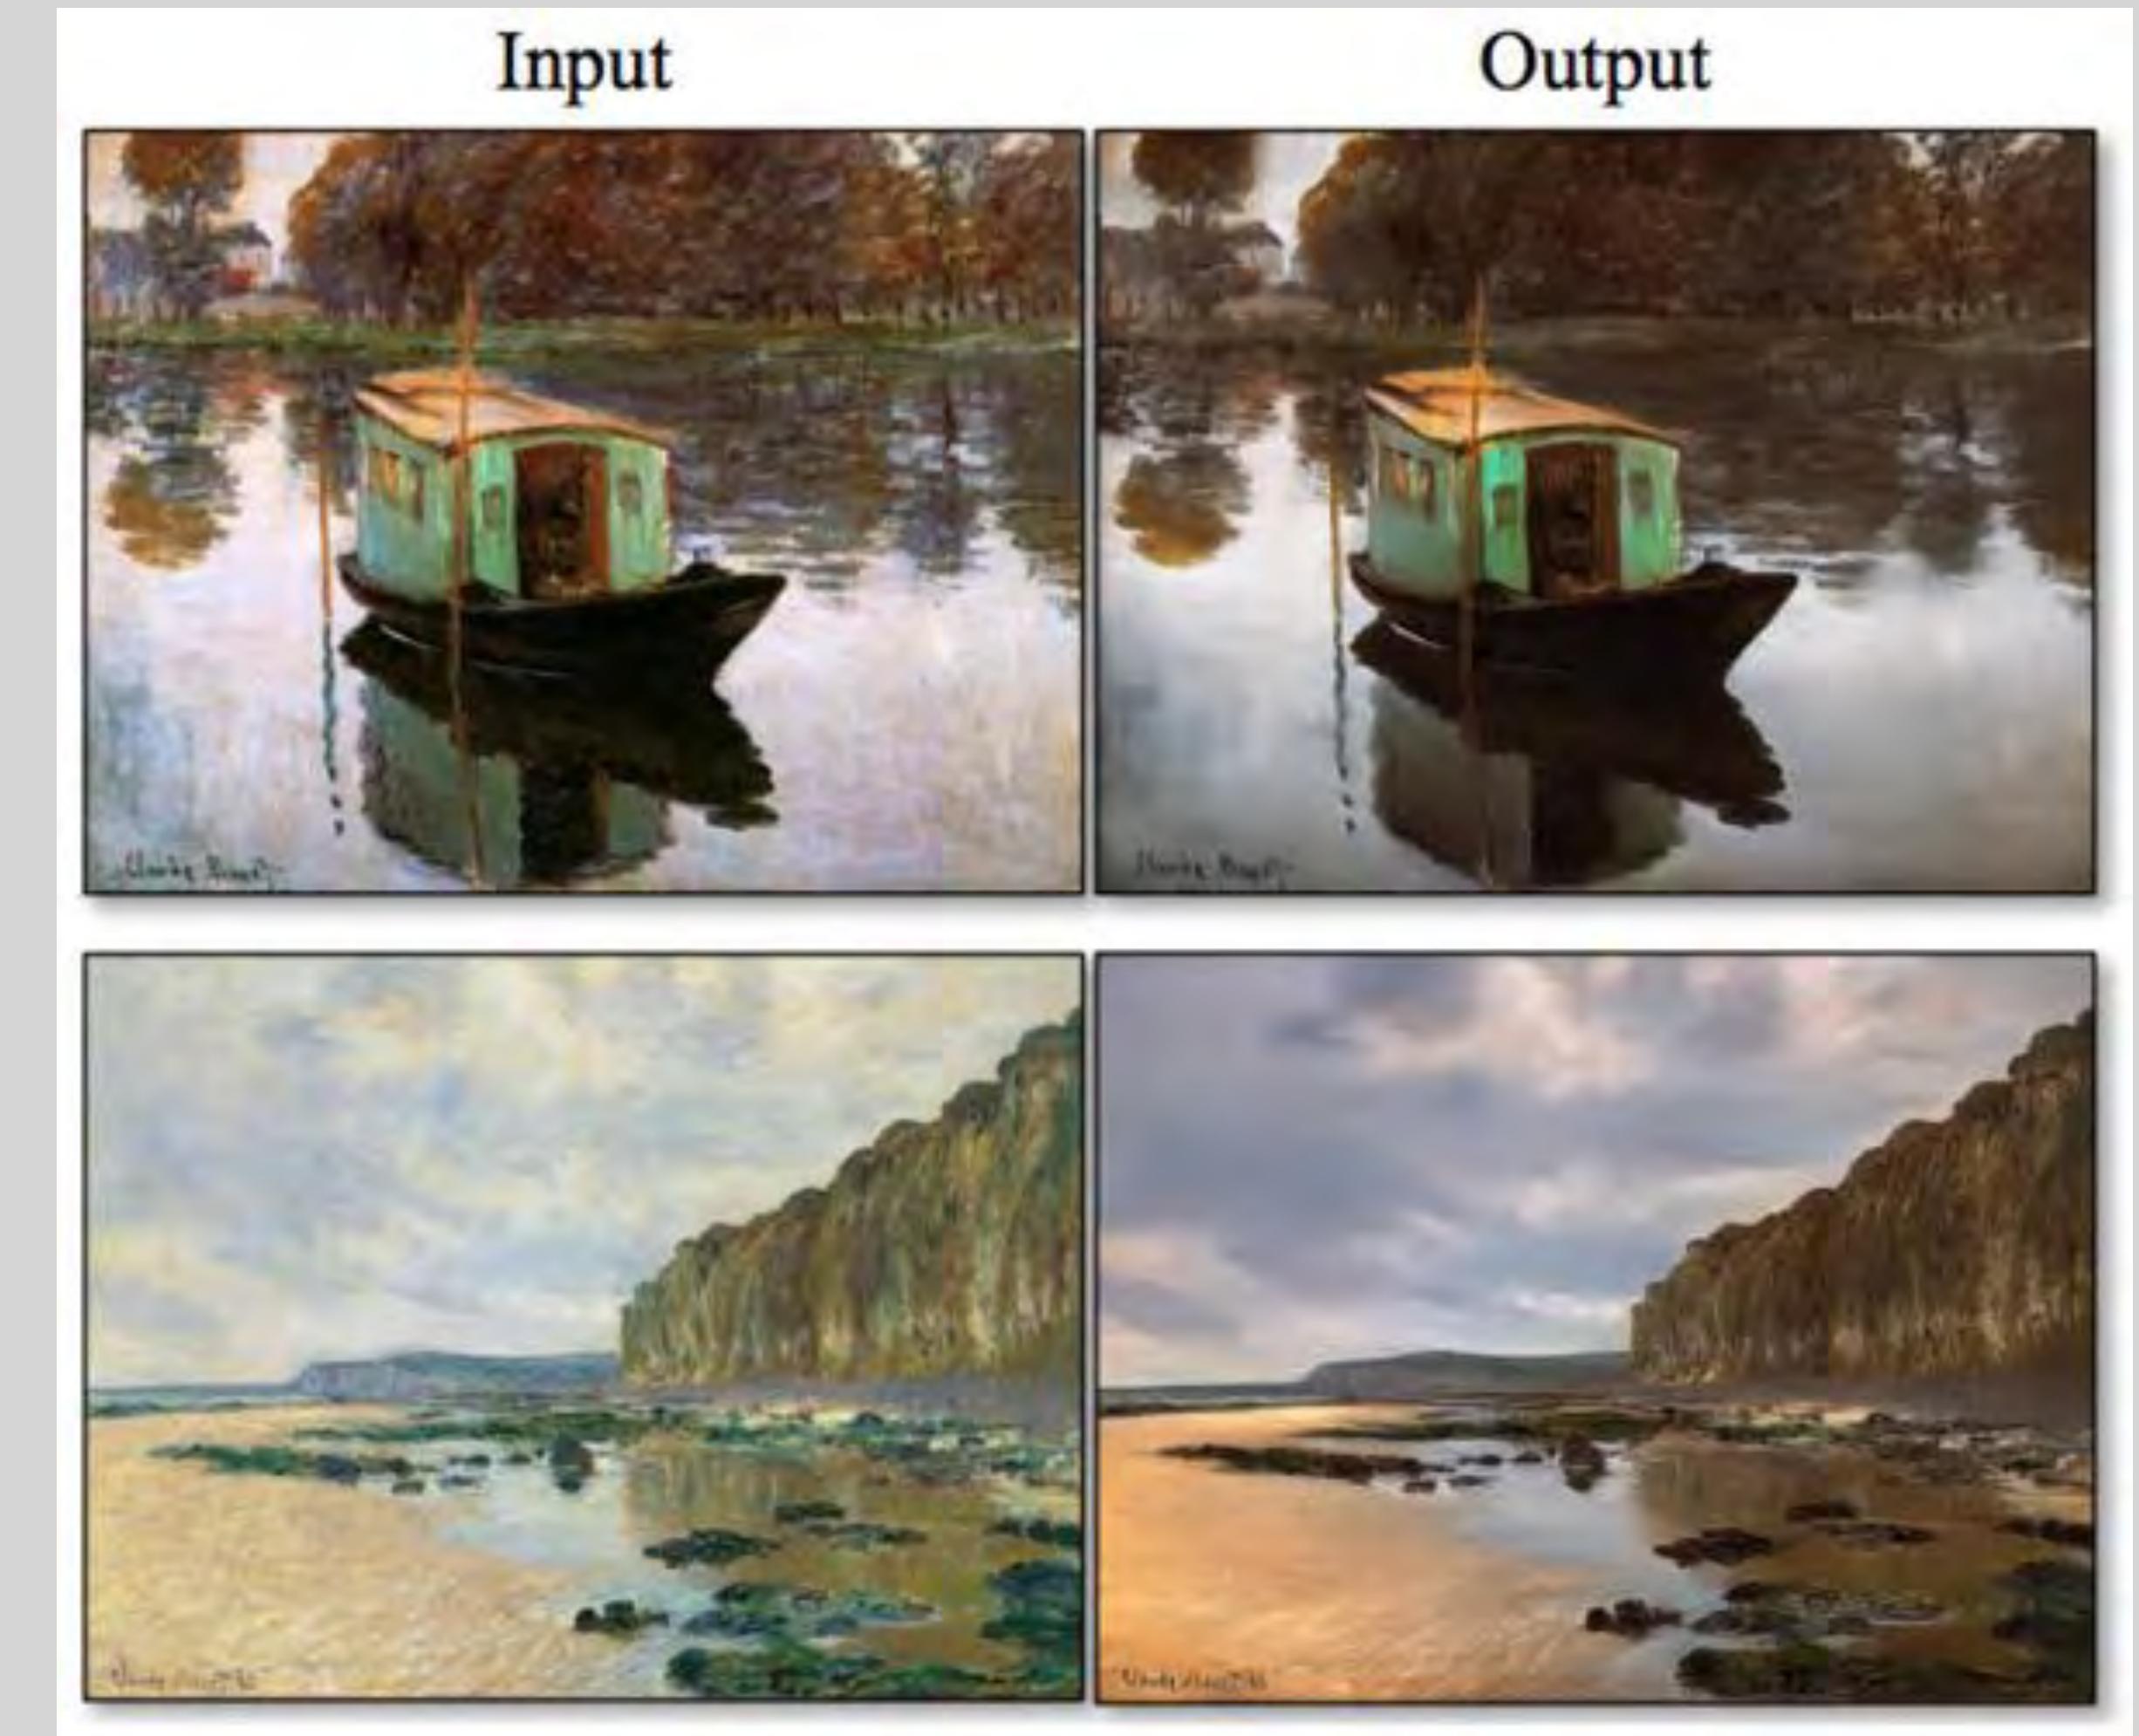


<https://affinelayer.com/pixsrv/index.html>

CycleGAN

There are many applications of CycleGAN, such as enhanced image resolution and style transfer. More specifically, it can be used for converting pictures to paintings (and vice versa), along with mapping one animal to another, as shown in the following example screenshots. The following are the examples cited in the CycleGAN paper on implementing CycleGAN:

Mapping Monet's paintings to a photographic style looks as follows:



Previous Work with GAN

A dense forest scene with a path winding through tall trees and fallen logs.

Video &
Motion Graphics
Through
Latent Walks



Different Projects Based our Students Interest.

Previous Projects



Logo-icon Generator



Font Generator

Icon-logo Project

Commission

Creation of logo for a
Japanese tea shop
(Branding design)

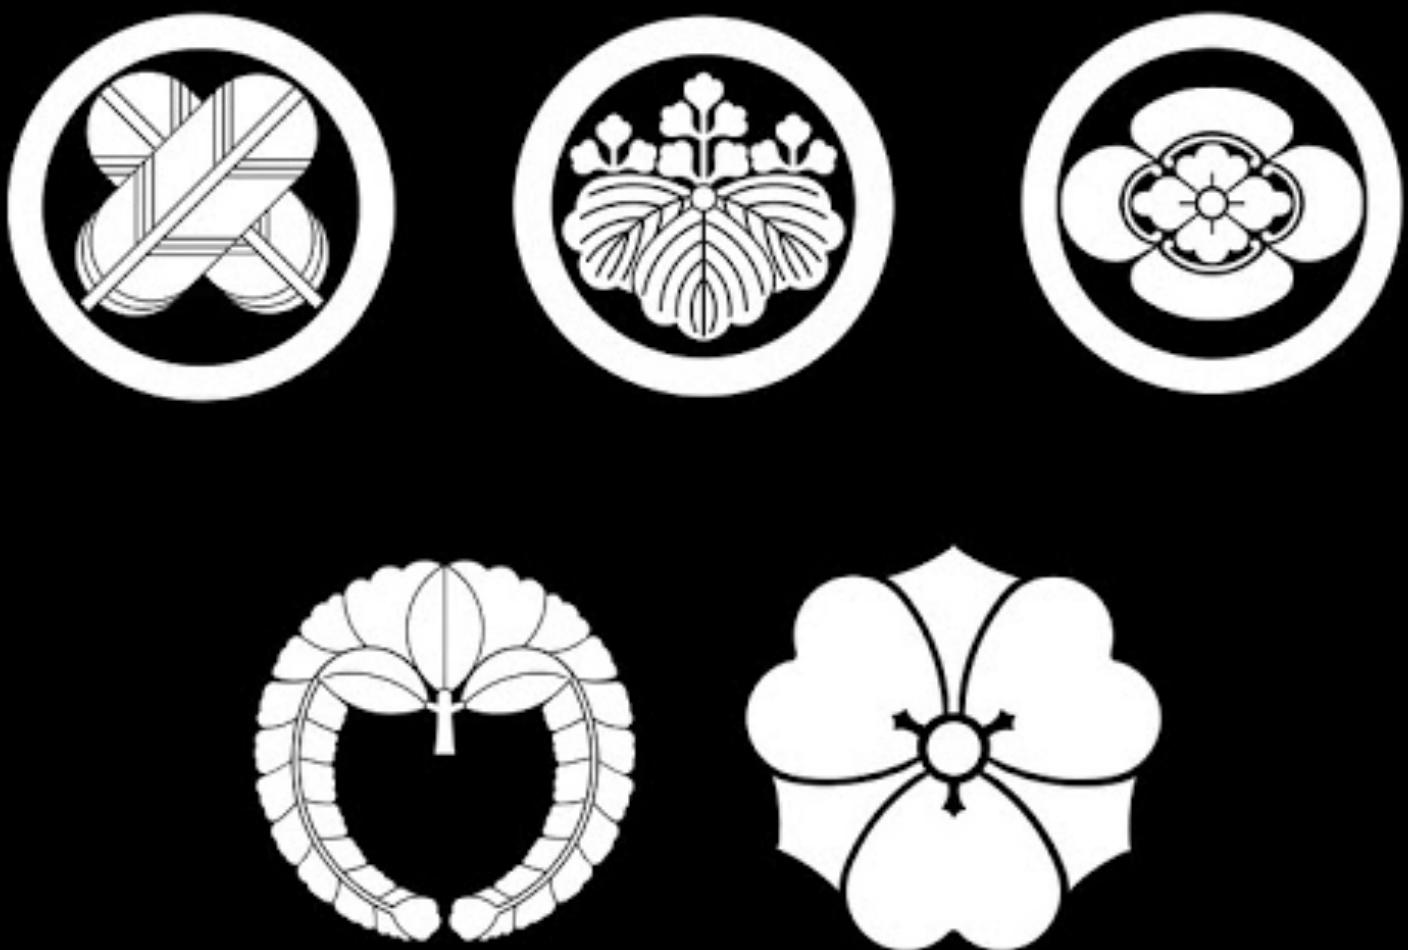
Objective

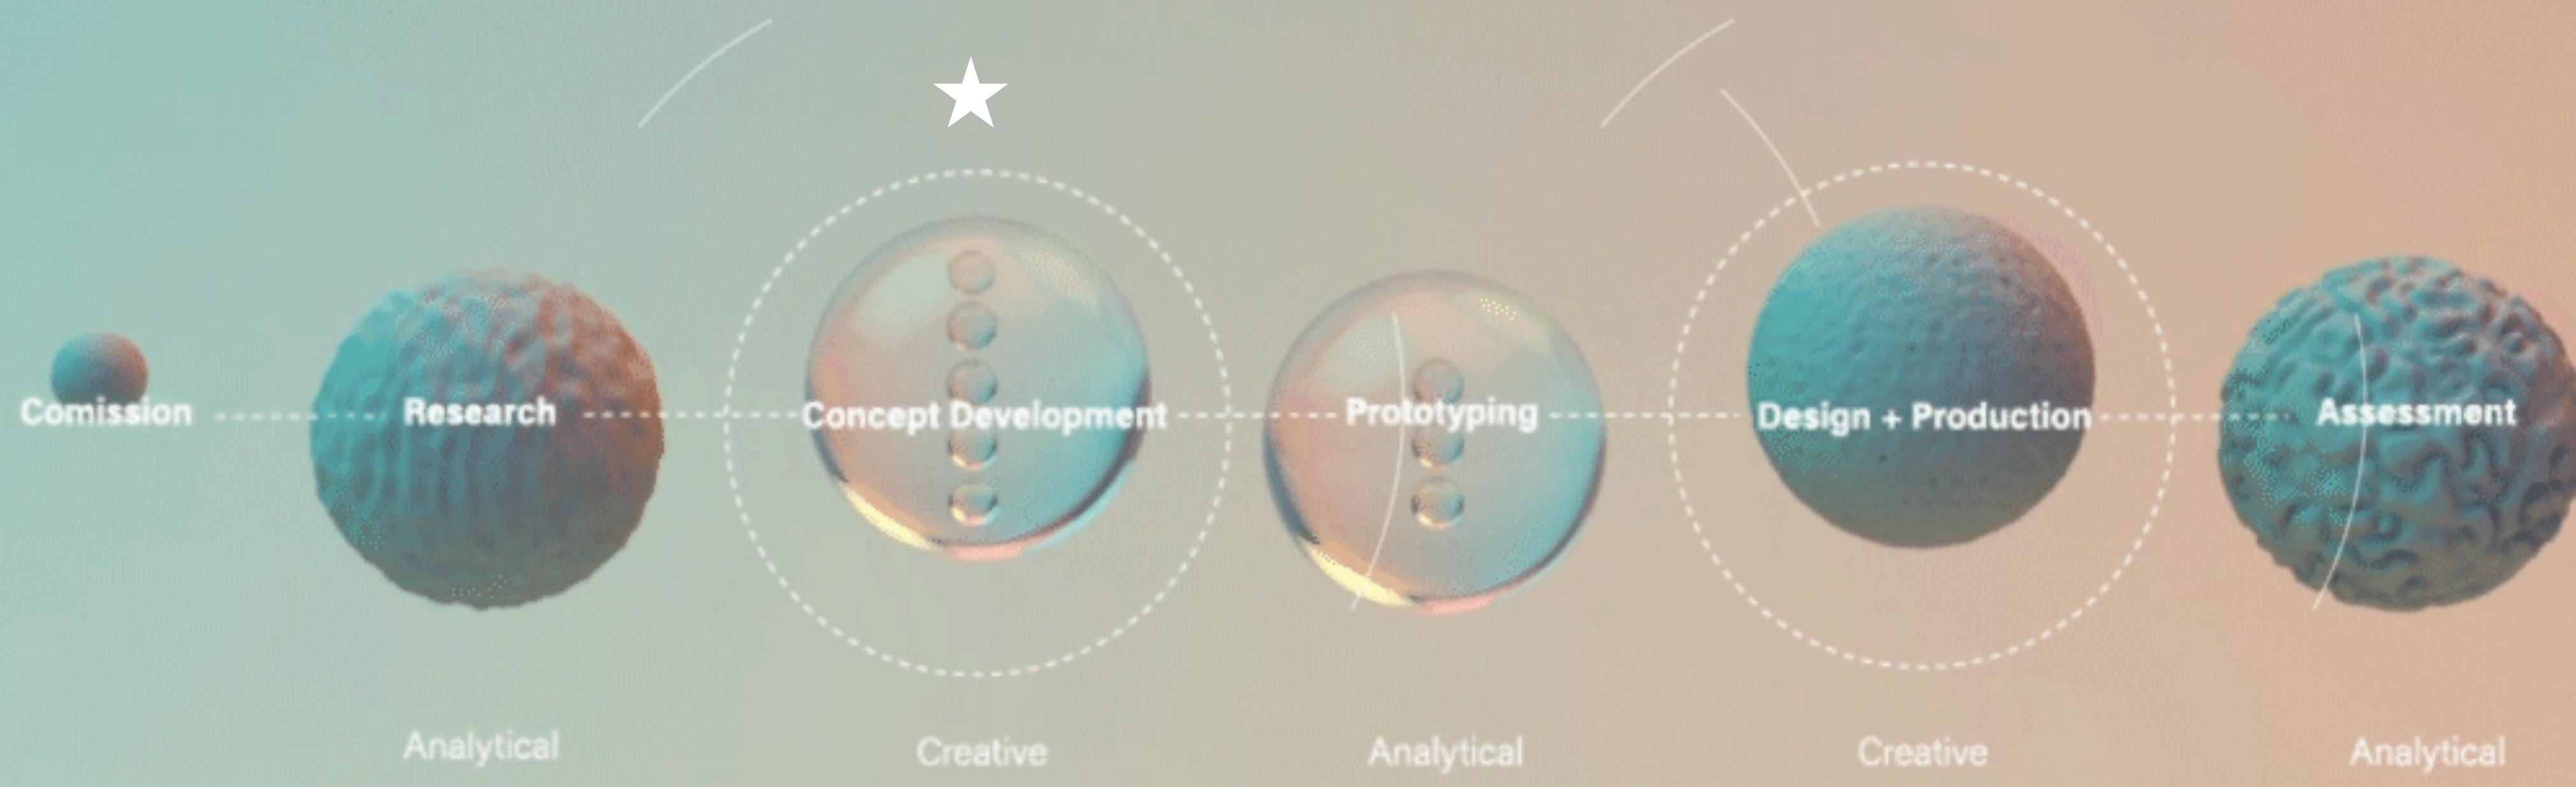
Creation of new Kamon
Using different elements related to
the Japanese tea shop

Kamon

Japanese family crest or emblems are graphic symbols associated with family clans or establishments and usually are used to decorate different objects like flags, armor, curtains, and personal items such as clothing.

The role of these family crests has changed over time. They are now used in a **commercial approach** as company logos.





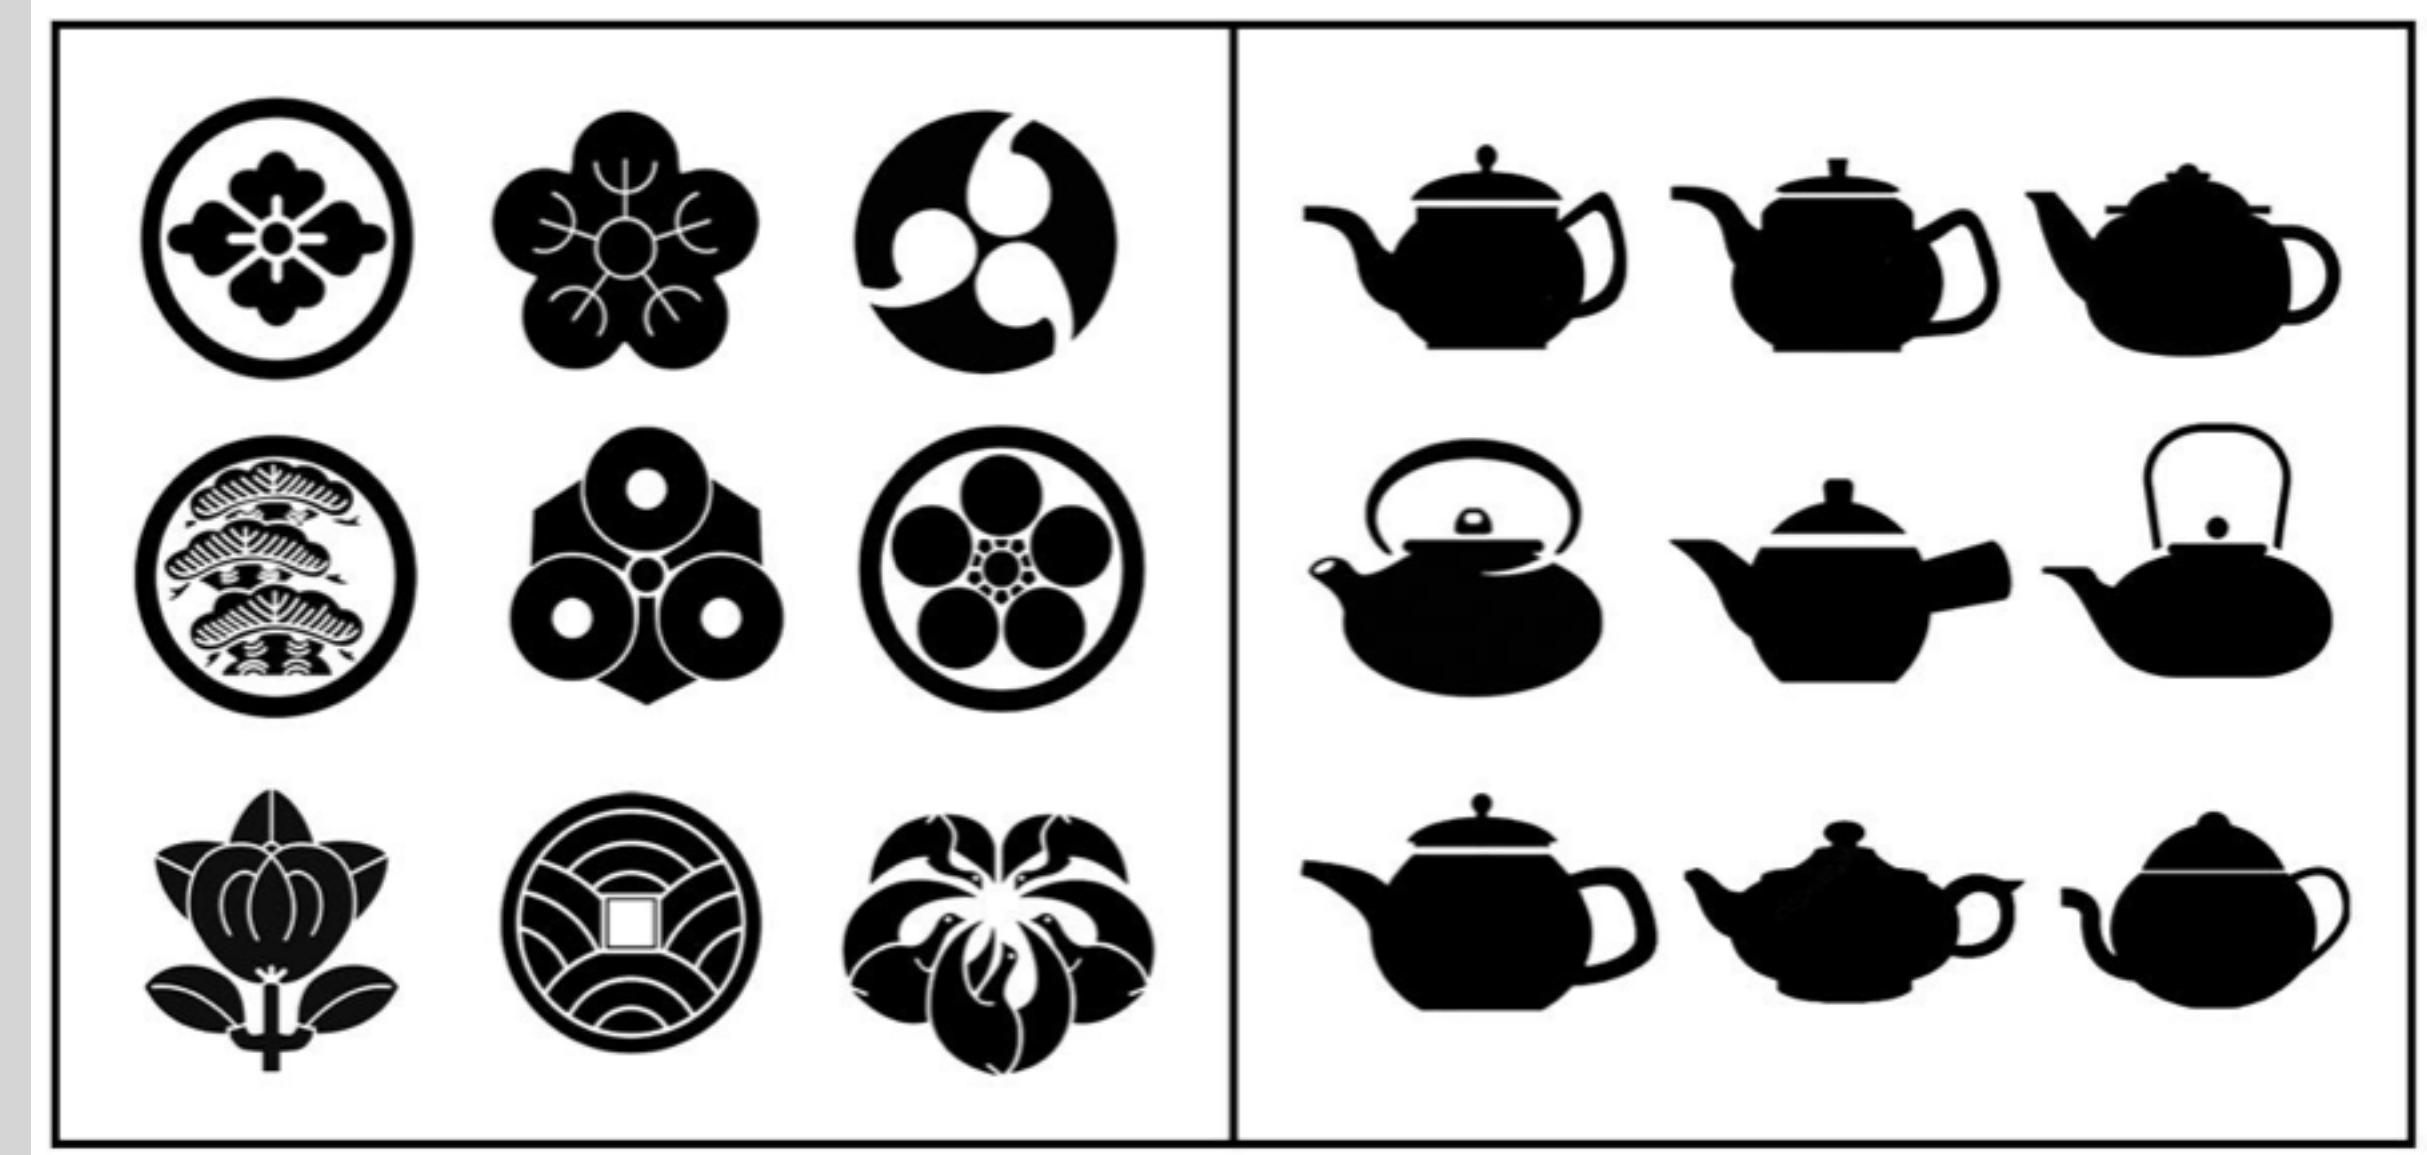
Ai as an Artificial co-worker

New Kamon

First test

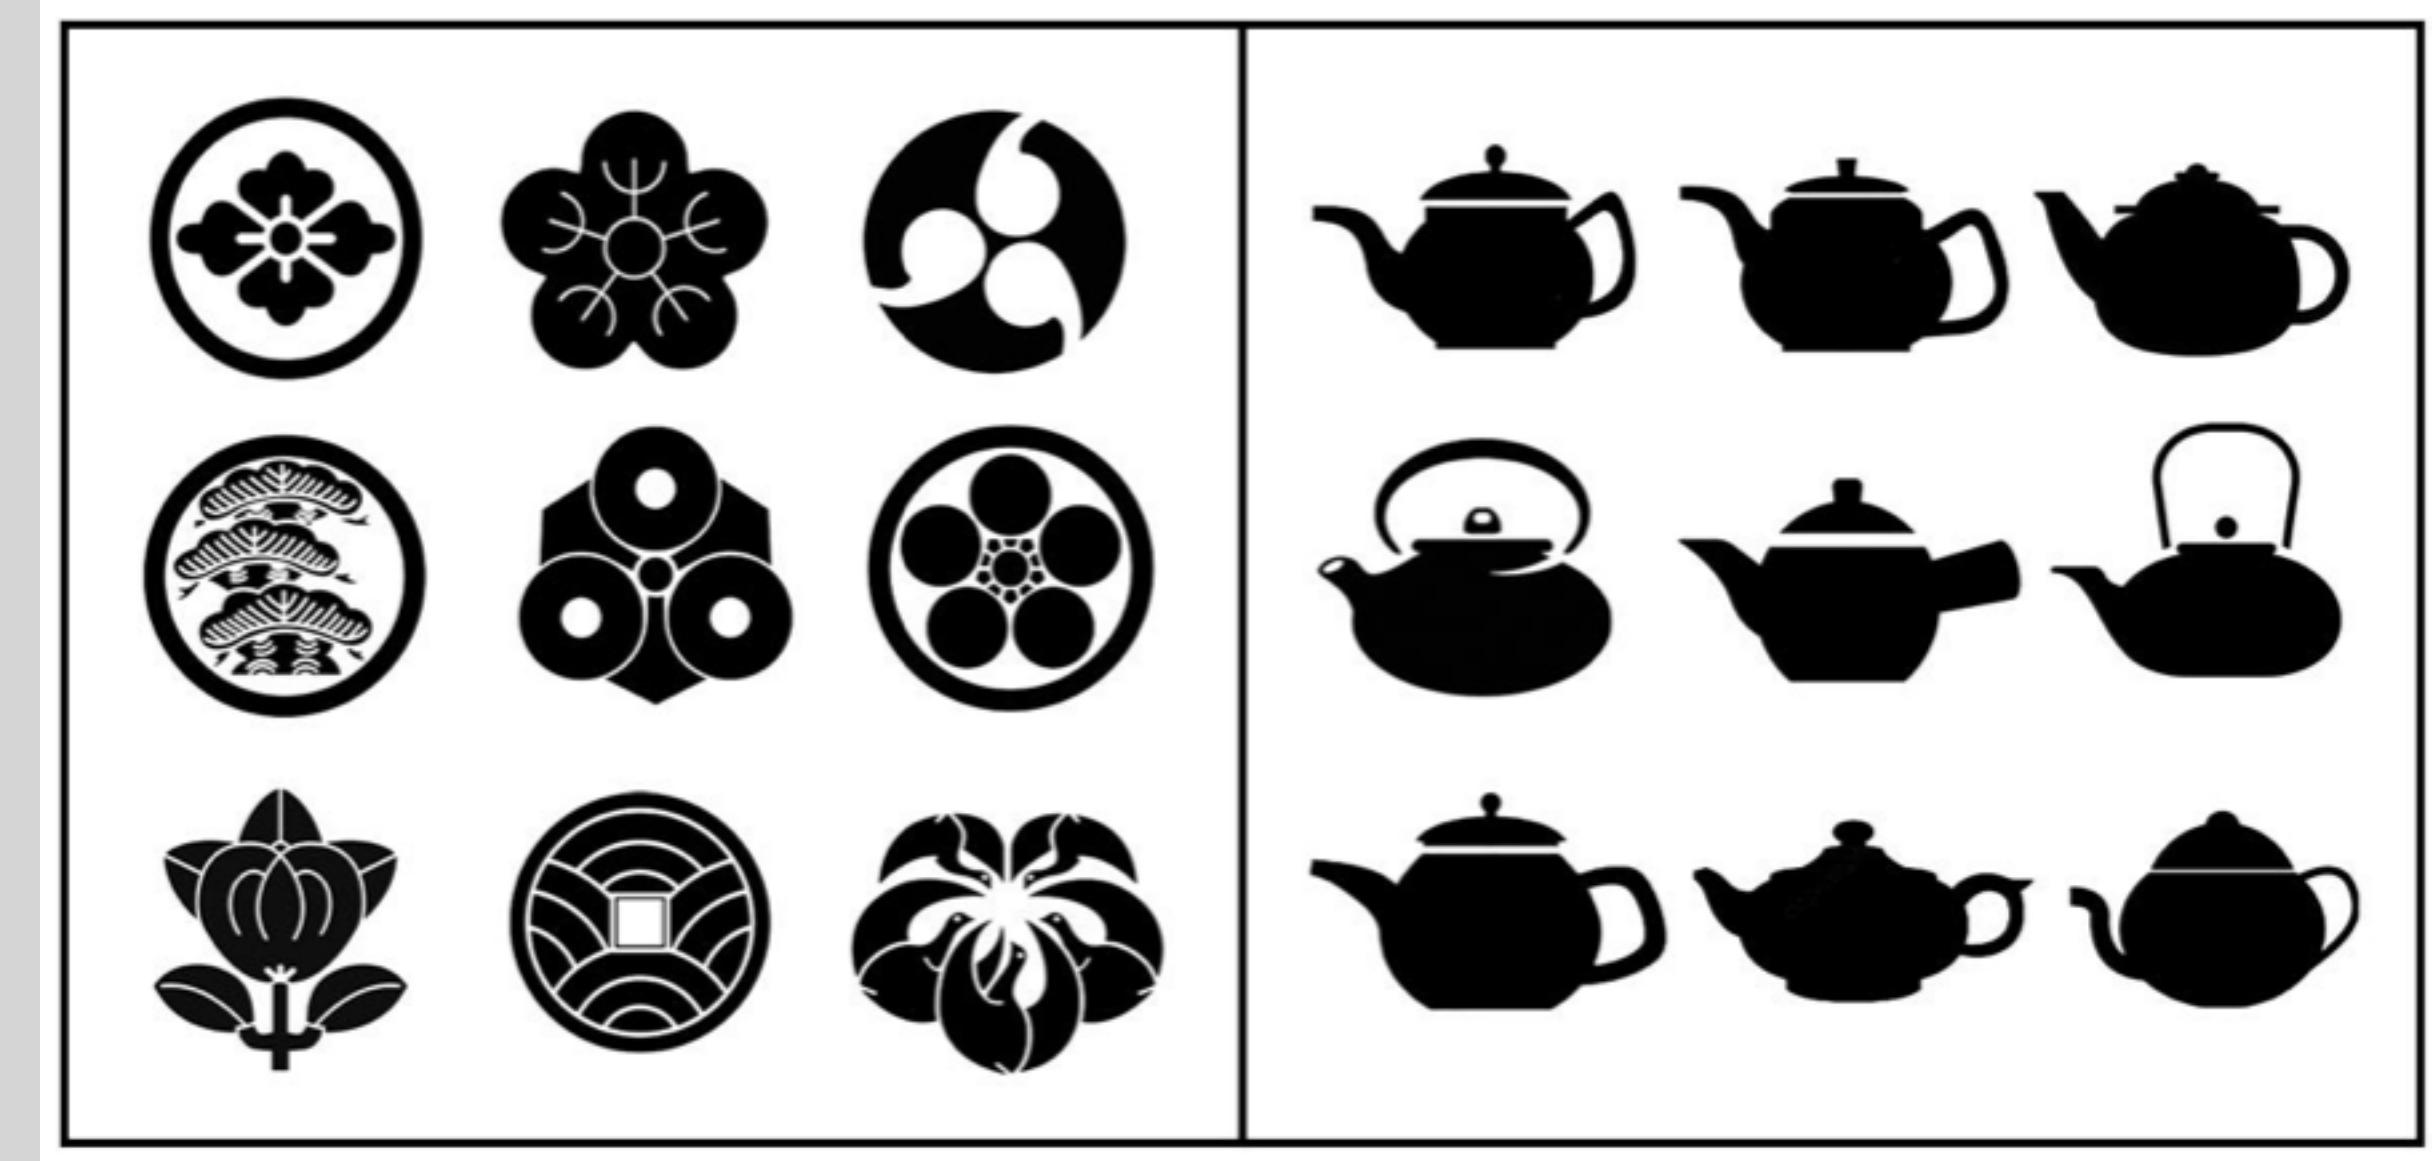
Mixing two
domains

Dataset was made of 200 standardized black
and white graphic symbols (256 x 256 pixels)
that belongs to two different domains: Japanese
Kamons and Teapots.



Dataset

We can establish
parameters and patterns
in the dataset and force the
output to follow and apply
these parameters and
patterns.



Dataset - 200 images



Results



Second test

Data Set - 13 Domains (650 Images)

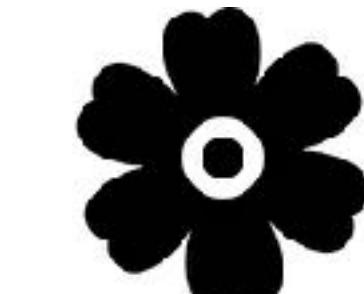
Cakes



Trees



Flowers



Kamon1



Kamon 2



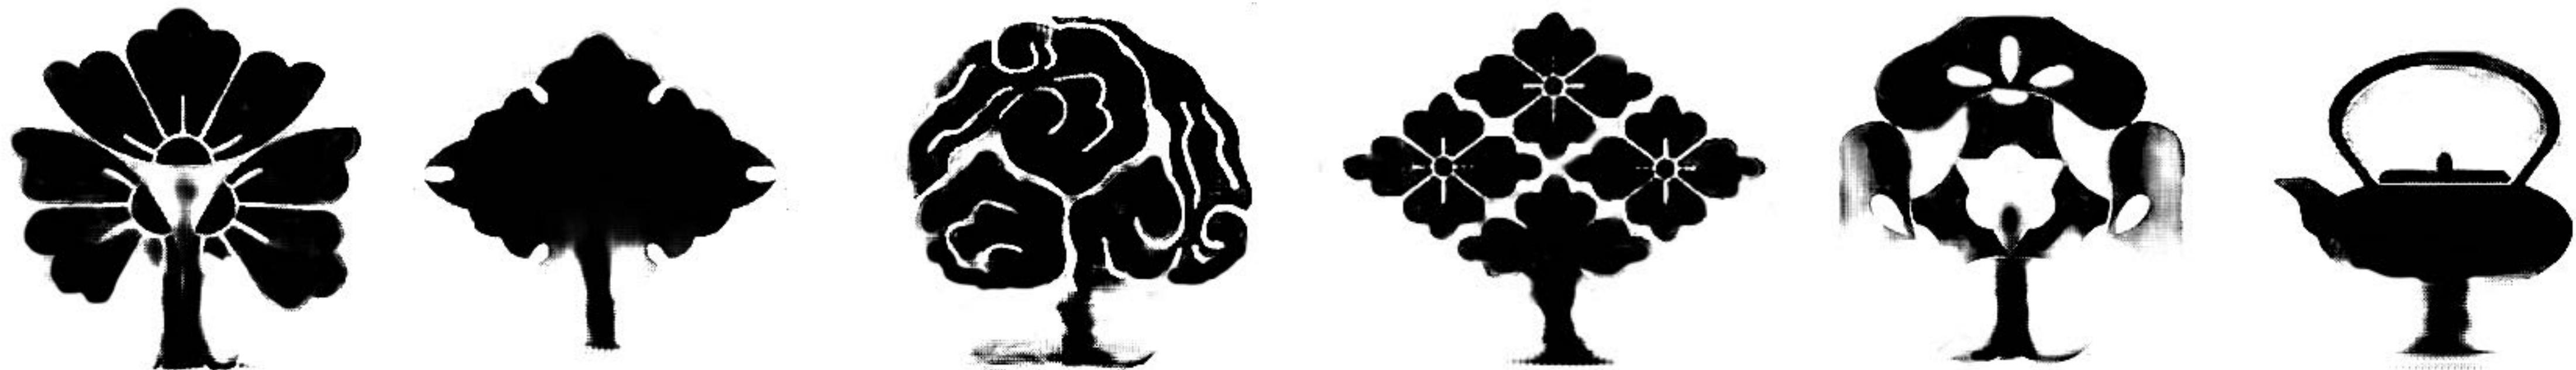
Kamon 3



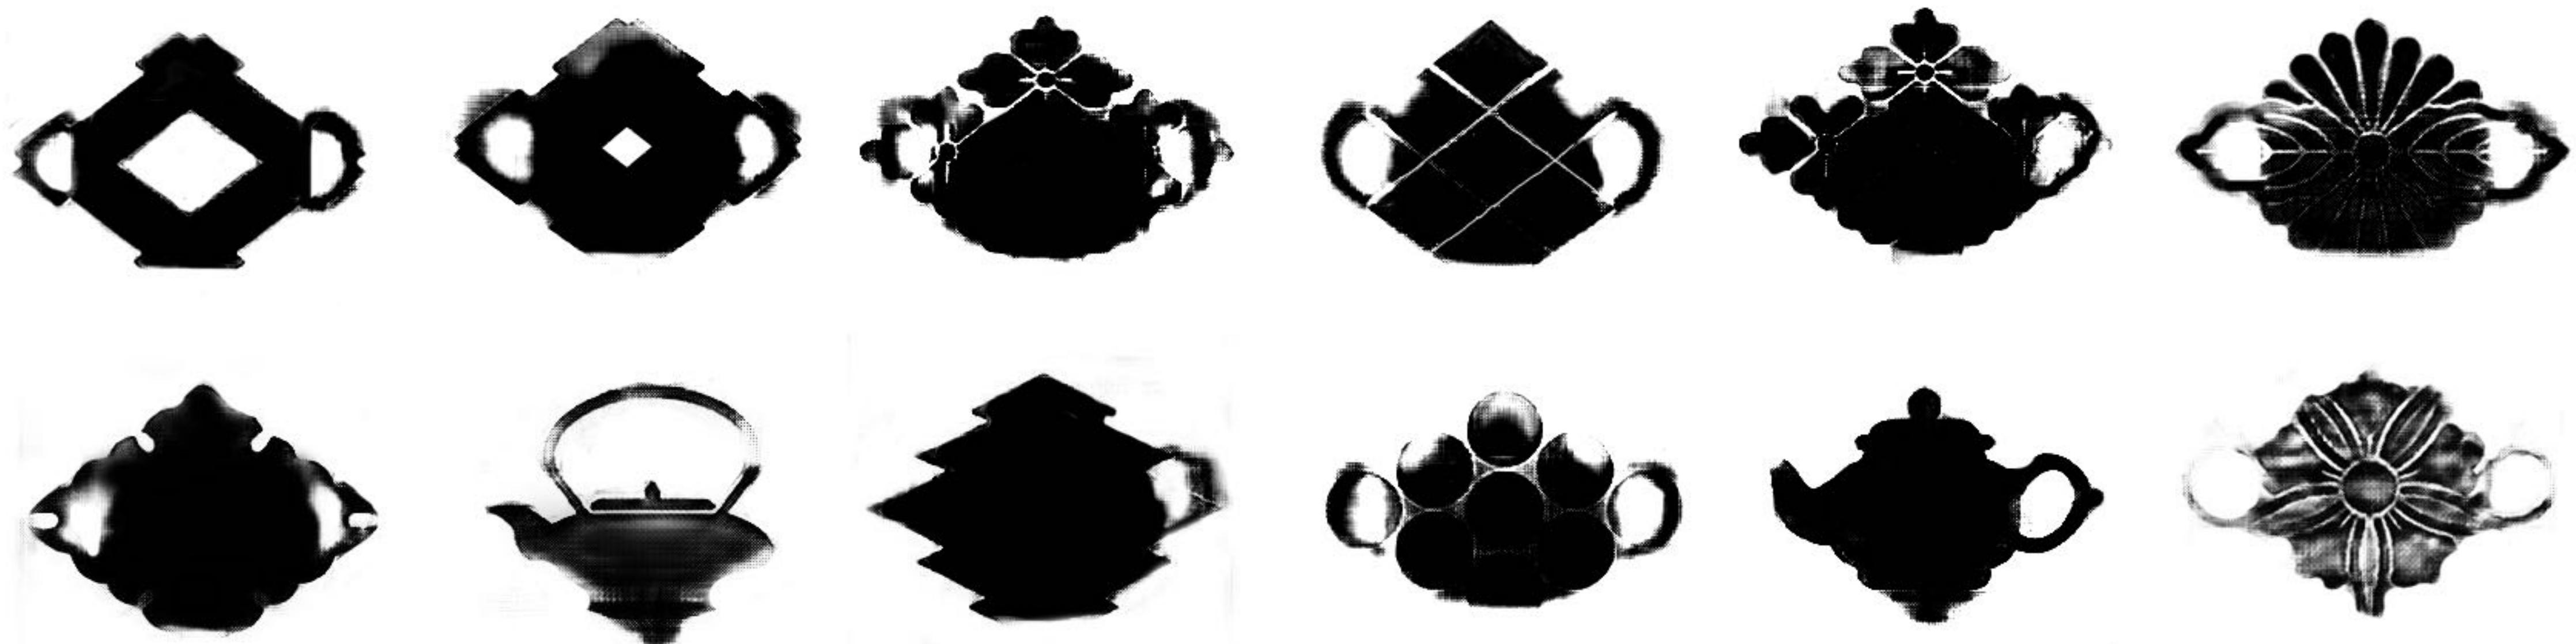
Kamon 4



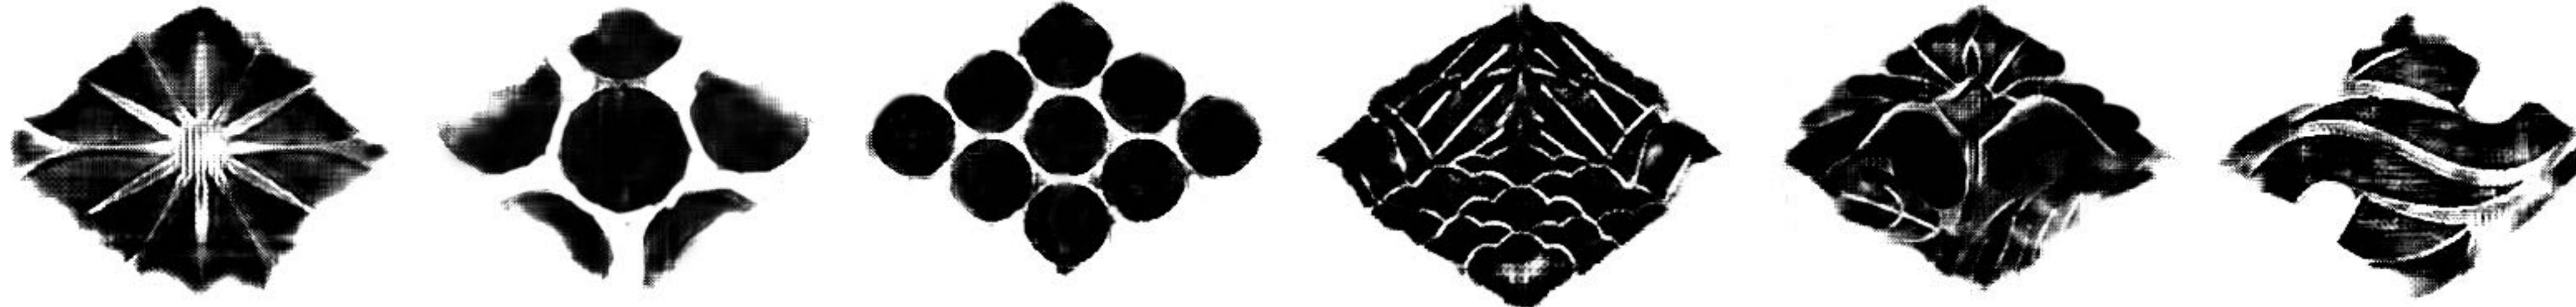
Results



Results



Results



Data set

We realized that by **combining different dimensions** in the dataset. (Different domains or categories) and forcing them to output a mix between this two domains you can generate new dimensions.

We can see this as the perfect approach to use GAN in a generative design approach.

Evaluation



Mr. Shoryu Hatoba / Monsho-Uwaeshi (artisan painter who paints family crests onto kimonos by hand).

What exactly constitutes a
“proper appearance and feel of a crest.”?

Hatoba:That....is extremely hard to explain.

This will get into where we draw the line between what is and can't be considered as being “crest-like.”

The difference **between a logo and a crest can be understood just by looking at the two**, but to explain this using words is pretty difficult.

Ask Ai vision

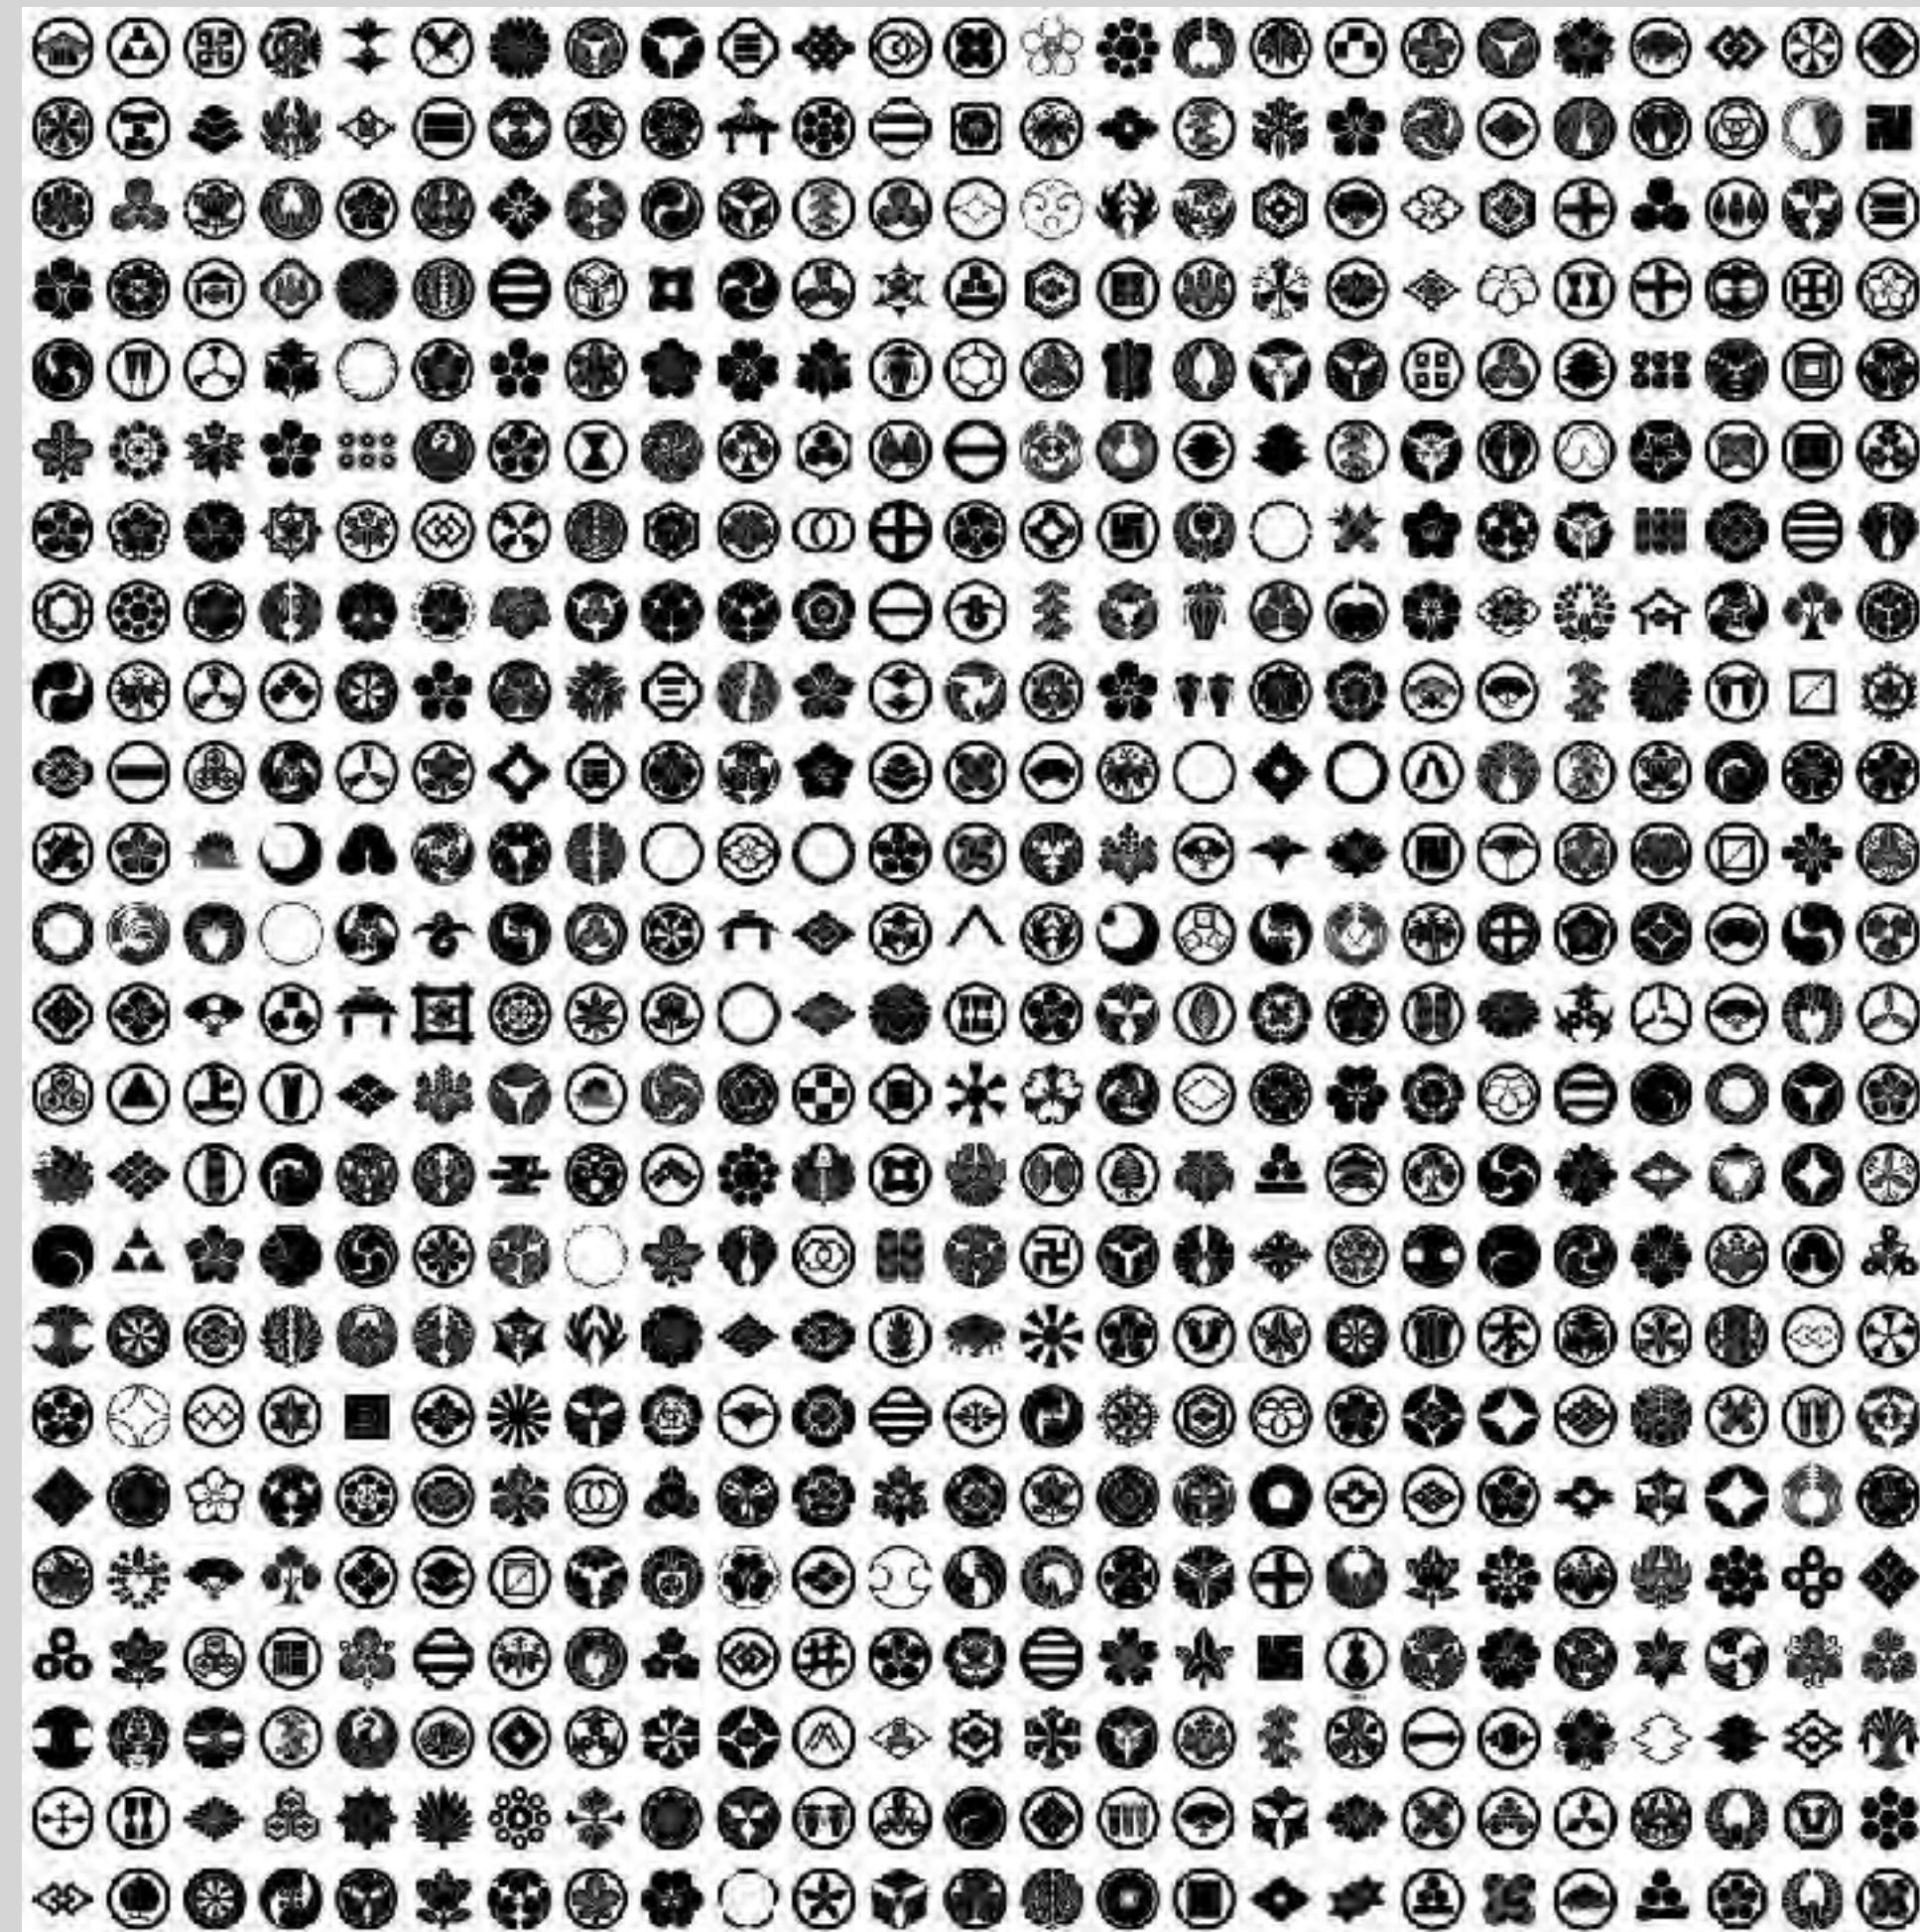
100 Kamon dataset samples:

30 labeled as **Emblems**

59 labeled as **Illustrations**

83 labeled as **Logos**

68 labeled as **Symbols**



Font Design Project

Goal

Create a Font that is legible
and usable.

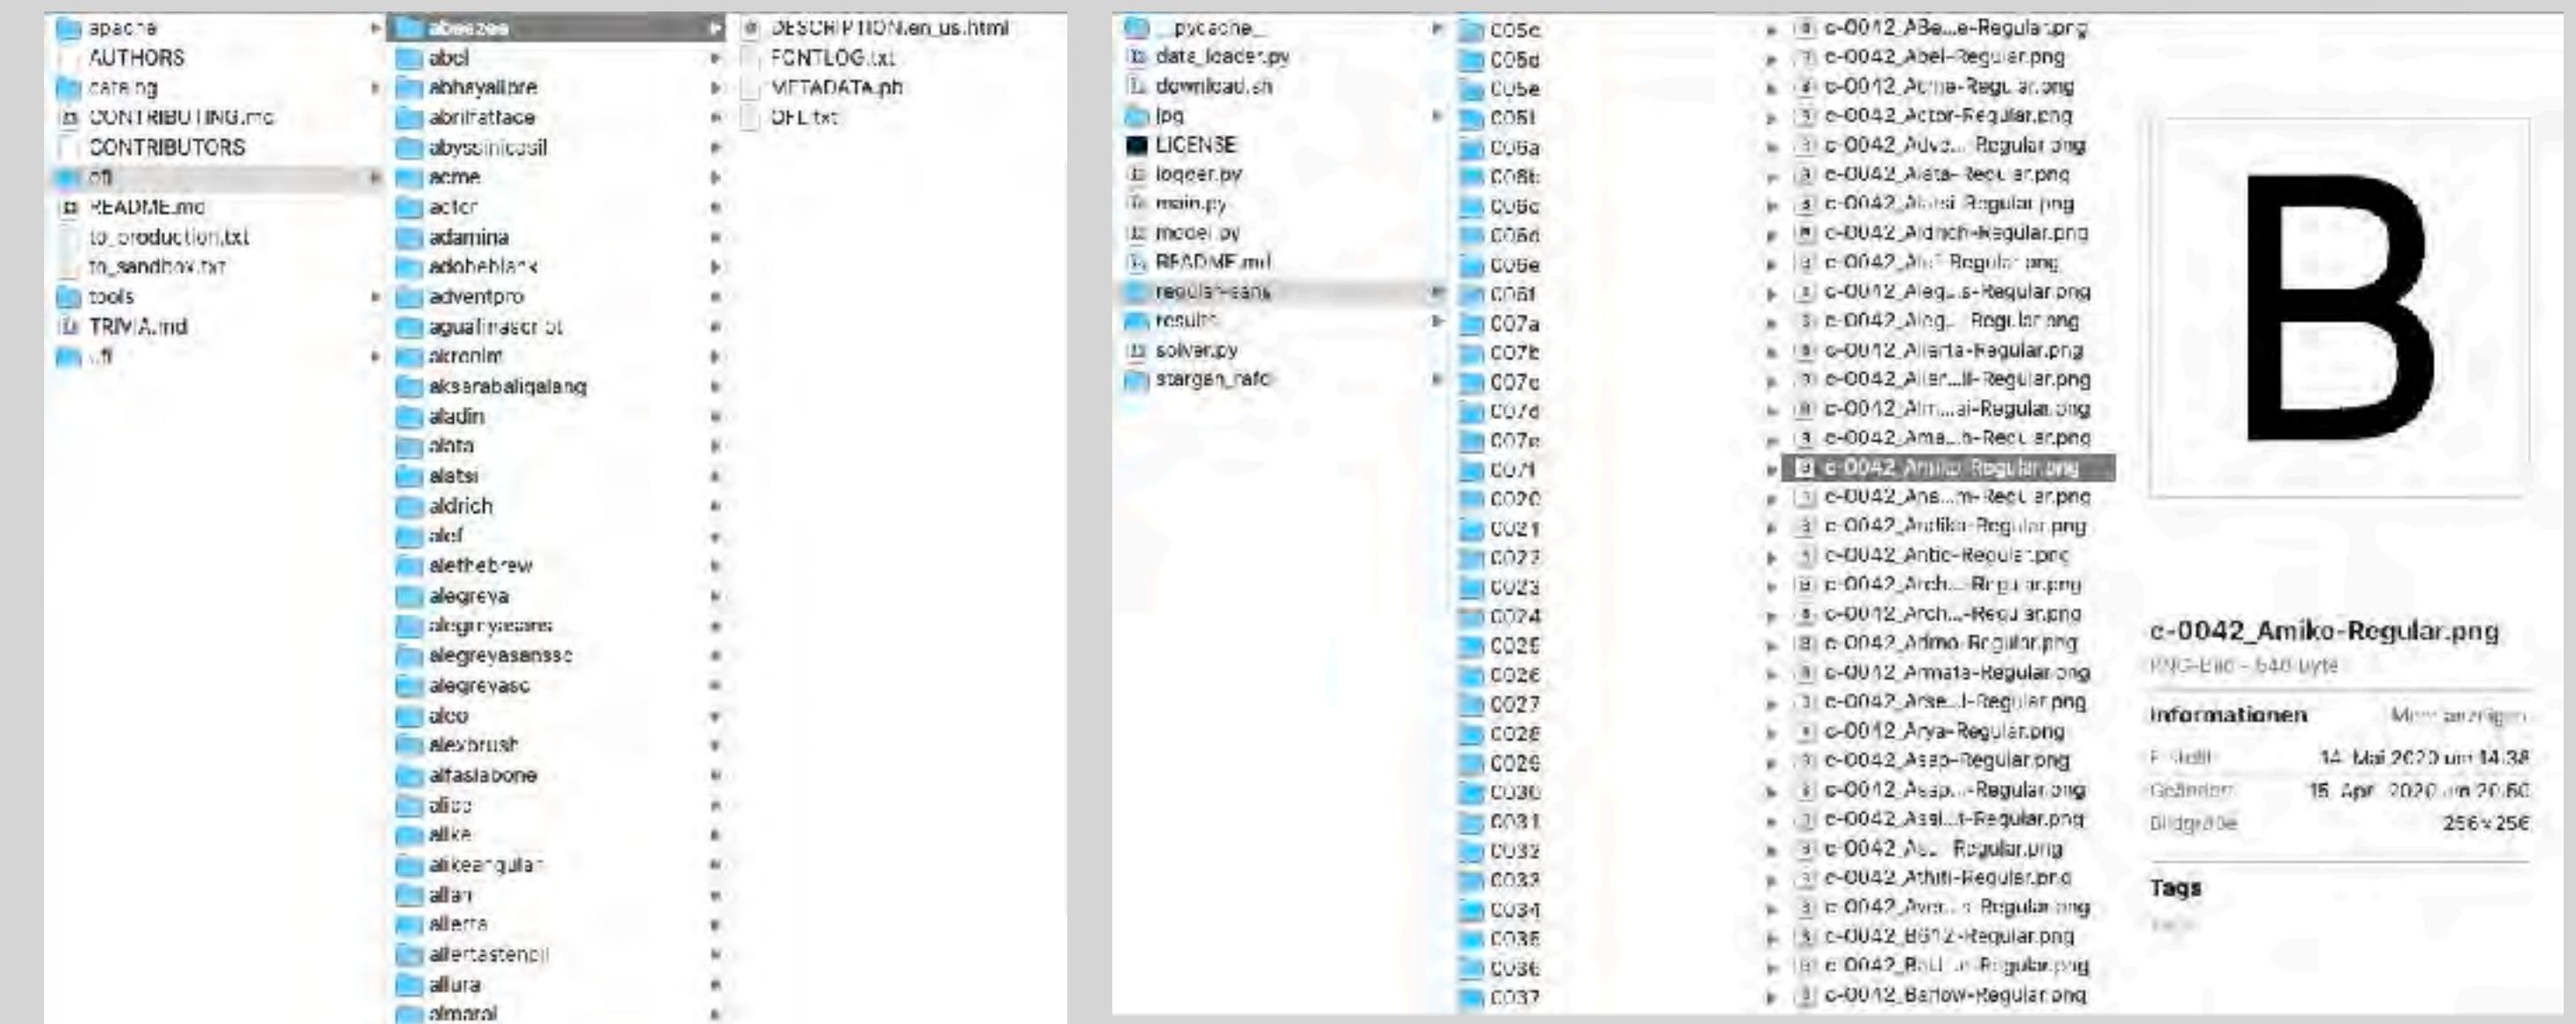
Objective

Develop a tool that creates a new
character set for every input.

Dataset – 22,000 Items

We got a bunch of Fonts (Google Master) and put them in a logical directory hierarchy.

Then we extracted all the glyphs and divided them by the UniCode.



Training

Y
K
2
A
E
9
G
W
M
R
P
S
I
N
-
Z
?

Training

Y
K
2
A
E
9
G
W
M
R
P
S
I
N
-
Z
?

Evaluating

! " # % ^ & ' () * + . - : /

0 1 2 3 4 5 6 7 8 9 ; ; < = > ?

Ⓐ A B C D E F G H I J K Ⓛ M N O

P Ⓛ R S T U V W X Y Z [\] ^ -

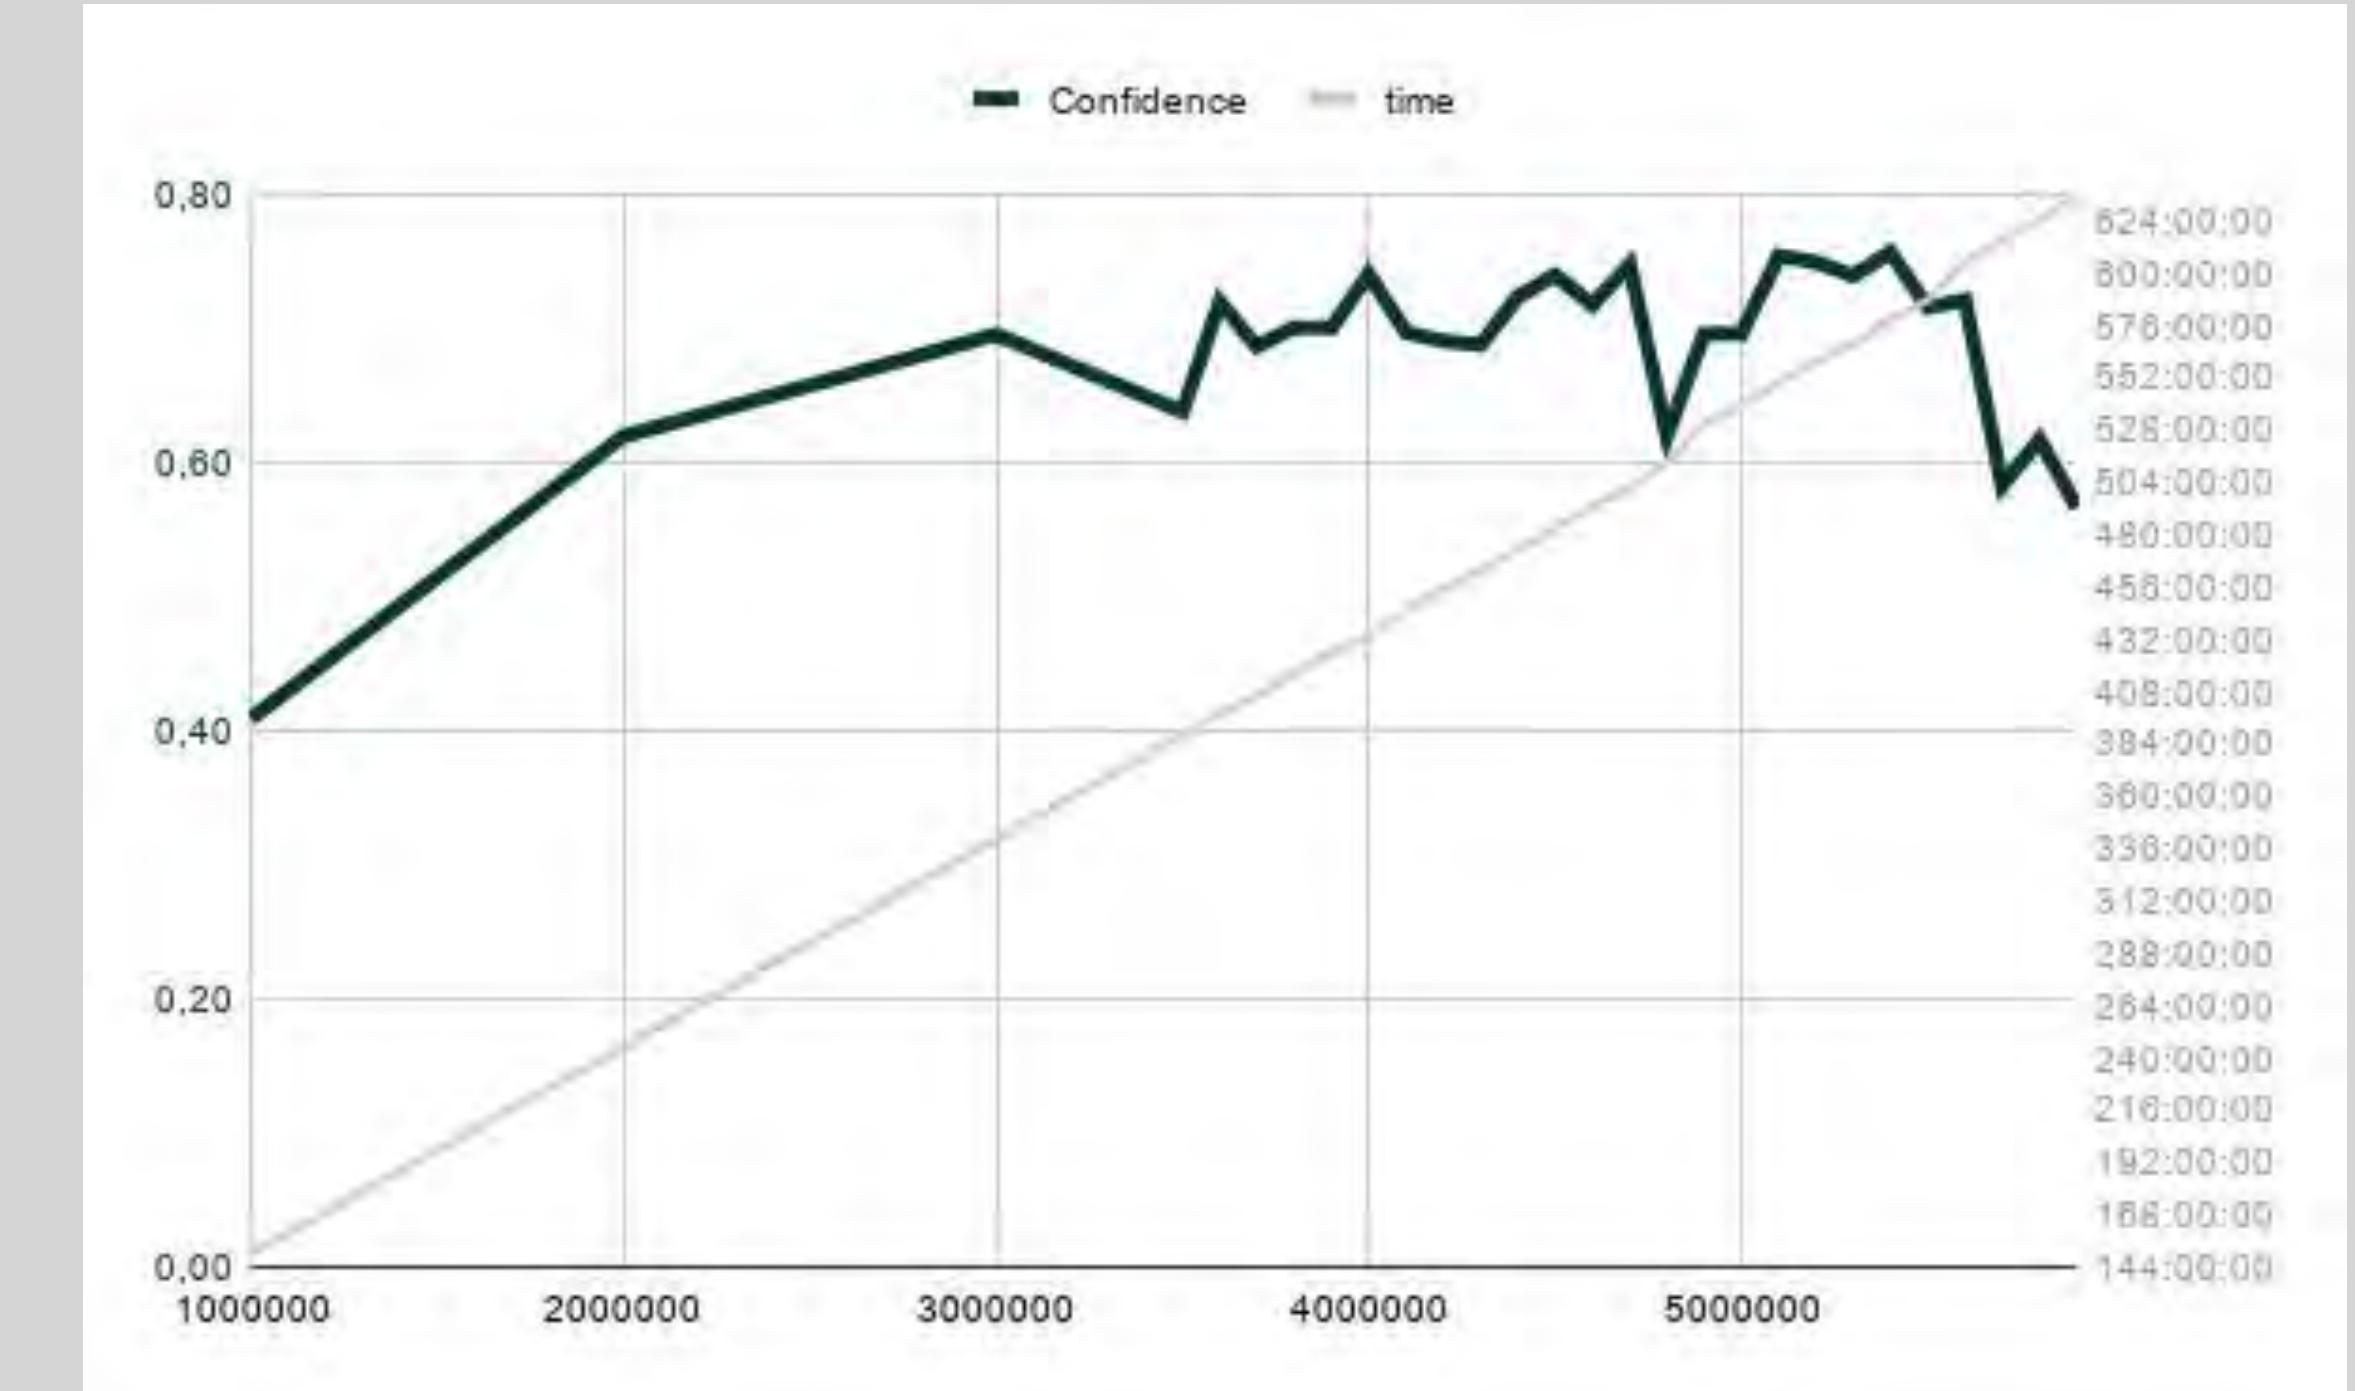
, a b c d e f g h i j k l m n o

p q r s t u v w x y z { | } ~

Ask Ai vision

Comparing the confidence of googles
character recognition.

Google Cloud's Vision API offers powerful pre-trained machine learning models through REST and RPC APIs. Assign labels to images and quickly classify them into millions of predefined categories.



Creating the actual Font file



0	!	"	#	\$	%	&	'	()	*	+	,	-	.	/
0	1	2	3	4	5	6	7	8	9	:	;	<	=	>	?
@	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
Ø	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
P	Q	R	S	T	U	V	W	X	Y	Z	[\	^	-	
Ø	R	S	T	U	V	W	X	Y	Z	[\	^	-		
~	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o
~	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o
p	q	r	s	t	u	v	w	x	y	z	{	}	-		
p	q	r	s	t	u	v	w	x	y	z	{	}	-		

Possibilities



Q&A

Questions & Answers

Cape Workshop 2021

	Thursday	Friday	Saturday	Sunday	Monday	Tuesday	Wednesday	Thursday	Friday
Time JP	February 18	February 19	February 20	February 21	February 22	February 23	February 24	February 25	February 26
10:00-12:00						Students Team work	Students Team work	Students Team work	Students Team work
13:00-14:00	Introduction to the Workshop and Instructions. Students Presentation 5 to 10 Mins Presentation per School (School Introduction and team members)	Generative Design using GAN Lecture Part. 2			Showing Results of Training sessions and Tracking	Presentation of first Ideas	Generative Design Lecture Part.3		
14:00-15:00	Presentation by Nagase Sensei & "Team formation"	Exercise using Generative Adversarial Networks	Training Day & Plant Tracking	Training Day & Plant Tracking	Conceptualization & preparation of first pitch presentation	Feedback Session		Preparation for Final Presentation	
15:00-16:00	Introduction to Generative Design Lecture Part 1.	Closing Exercise 1			Team Work	Advise Session	Design & Prototyping Team Work	Preparation for Final Presentation	Final Presentation
16:00-17:00									Feedback & Closure

* During the sessions marked in Yellow, Professors are Invited to Attend

Create Github Account

Next Session: GAN Lecture 2 Exercise by teams



See you tomorrow!

Contact us!

Hisa Nimi

hisanimi@chiba-u.jp

Design Researcher
Chiba University

Juan Carlos Chacón

juancarlos@chiba-u.jp

Design Researcher
Chiba University